

Empirical Performance Investigation of a Büchi Complementation Construction

Daniel Weibel

August 2015

Master's Thesis

Supervised by:
Prof. Dr. Ulrich Ultes-Nitsche



UNIVERSITÉ DE FRIBOURG
UNIVERSITÄT FREIBURG

Abstract

Acknowledgements

First of all, I would like to thank my supervisors Prof. Dr. Ulrich Ultes-Nitsche and Joel Allred for their support and patience during the entire time I was working on this thesis. A very special thanks goes to Dr. Ming-Hsien Tsai from National Taiwan University. He is one of the main authors of the GOAL tool and steadily provided me with all the information I needed to fully integrate the plugin, developed as part of this thesis, in the GOAL tool. Without his help, the plugin would not have been possible in its present form. I am especially thankful to Assoc. Prof. Katarzyna Wac from the University of Geneva for granting me a lot of freedom from my employment in her research group to work on my thesis. Without this extra time, it would have been very difficult to complete the thesis within the deadlines of the university. I also would like to thank Christian Göttel for sharing information about his own thesis about a very similar subject, from which I adopted several ideas. Last but not least, I would like to thank Michael Rolli and Nico Färber from the Grid Admin Team of the University of Bern for their help regarding technical questions about the computer cluster on which the experiments of this thesis have been executed.

Contents

1	Introduction	3
1.1	Context	4
1.1.1	Büchi Automata and Büchi Complementation	4
1.1.2	Application of Büchi Complementation	5
1.1.3	Importance of Büchi Complementation	7
1.2	Motivation	8
1.2.1	Theoretical Investigation of Worst-Case Performance	8
1.2.2	Empirical Investigation of Actual Performance	9
1.3	Aim and Scope	10
1.4	Overview	11
2	Background	12
2.1	Büchi Automata and Other Omega-Automata	13
2.1.1	Büchi Automata	13
2.1.2	Other Omega-Automata	15
2.2	Run Analysis of Non-Deterministic Automata	16
2.2.1	Run DAGs	17
2.2.2	Run Trees	17
2.2.3	Split Trees	18
2.2.4	Reduced Split Trees	19
2.3	Review of Büchi Complementation Constructions	20
2.3.1	Ramsey-Based Approach	21
2.3.2	Determinisation-Based Approach	21
2.3.3	Rank-Based Approach	23
2.3.4	Slice-Based Approach	24
3	The Fribourg Construction	26
3.1	The Construction	27
3.1.1	Basics	27
3.1.2	Upper-Part Construction	29
3.1.3	Lower-Part Construction	31
3.2	Optimisations	36
3.2.1	R2C: Omission of States Whose Rightmost Component is 2-Coloured	36
3.2.2	M1: Merging of Adjacent Components	37
3.2.3	M2: Single 2-Coloured Component	37
3.3	General Remarks	39
4	Empirical Performance Investigation of the Fribourg Construction	40
4.1	Implementation	41
4.1.1	The GOAL Tool	41
4.1.2	Implementation of the Construction	43
4.1.3	Verification of the Implementation	45
4.2	Test Data	46
4.2.1	GOAL Test Set	46
4.2.2	Michel Test Set	49

4.3	Experimental Setup	50
4.3.1	Internal Tests	51
4.3.2	External Tests	52
4.3.3	Time and Memory Limits	53
4.3.4	Execution Environment	54
5	Results and Discussion	56
5.1	Results of the Internal Tests	57
5.1.1	GOAL Test Set	57
5.1.2	Michel Test Set	65
5.2	Results of the External Tests	68
5.2.1	GOAL Test Set	68
5.2.2	Michel Test Set	71
5.3	Discussion	73
5.3.1	Summary of the Results	73
5.3.2	Gained Insights	74
5.3.3	Limitations of the Study	76
6	Conclusions	78
A	GOAL and Plugin Installation	79
B	Matrices Corresponding to Perspective Plots	80
C	Execution Times	82

Chapter 1

Introduction

Contents

1.1	Context	4
1.1.1	Büchi Automata and Büchi Complementation	4
1.1.2	Application of Büchi Complementation	5
1.1.3	Importance of Büchi Complementation	7
1.2	Motivation	8
1.2.1	Theoretical Investigation of Worst-Case Performance	8
1.2.2	Empirical Investigation of Actual Performance	9
1.3	Aim and Scope	10
1.4	Overview	11

AT THE BEGINNING of the 1960s, a Swiss logician named Julius Richard Büchi was searching a procedure for deciding the satisfiability of formulas of the monadic second order logic with one successor (S1S). In his quest, Büchi observed that an S1S formula can be represented by a certain type of finite state automaton running on infinite words, such that this automaton accepts a word representing an interpretation of the formula, if and only if the interpretation satisfies the formula. The proof of this equivalence between S1S formulas and this type of automata is known as *Büchi's Theorem*. It led Büchi to his desired decision procedure: to test whether an S1S formula φ is satisfiable, translate it to an equivalent automaton A , and test whether A is empty (that is, accepts no words at all). If A is empty, then φ is unsatisfiable, if A is non-empty, then φ is satisfiable. [8]

This thesis is about the type of automata that Büchi invented in order to solve the satisfiability problem of S1S. These automata are called *Büchi automata*. Büchi automata have important applications in logic-related domains. One of them is the automata-theoretic approach to model checking, a branch of formal verification. It is based on the specification of a property to verify as a Büchi automaton. The test whether the system satisfies the property require the *complementation* of this Büchi automaton. Complementing a Büchi automaton A means to construct another Büchi automaton B that accepts a word if and only if it is not accepted by A . An algorithm that takes as input a Büchi automaton and outputs the complement of this automaton is called a *Büchi complementation construction*.

Büchi complementation constructions are an active research topic since the introduction of Büchi automata more than 50 years ago. The reason for this persistence is that these constructions exhibit a very high *state complexity* (also called state growth). This means that, given an input Büchi automaton, its complement may turn out to be extremely large. The complexity is so high that it inhibits the practical application of Büchi complementation [62, 63]. For this reason, there is an ongoing quest to find more efficient Büchi complementation constructions, and to generally better understand the problem of Büchi complementation.

The aim of this thesis is to contribute to a better understanding of the Büchi complementation problem. Our chosen approach is an empirical one, an

In this introductory chapter, we first present the context in which this thesis is situated (Section 1.1). It includes Büchi automata and the problem of Büchi complementation in general, and the applications and significance of Büchi complementation. In Section ??, we motivate the empirical approach for increasing our understanding of Büchi complementation. In Section ??, we clarify the aim and scope of this thesis. Finally, in Section ??, we present the structure of the rest of this thesis.

1.1 Context

In this section, we situate Büchi complementation in its context. To this end, we first briefly present Büchi automata and Büchi complementation in general (Section 1.1.1). Next, we present an application of Büchi complementation in formal verification (Section ??). Finally, in Section 1.1.3, we highlight the importance of Büchi complementation for this application in formal verification.

1.1.1 Büchi Automata and Büchi Complementation

Büchi automata are finite state automata that process words of infinite length, so called ω -words. If Σ is the alphabet of a Büchi automaton, then the set of all the possible ω -words that can be generated from this alphabet is denoted by Σ^ω . A word $\alpha \in \Sigma^\omega$ is accepted by a Büchi automaton if it results in at least one run that contain infinitely many occurrences of at least one accepting state. A run of a Büchi automaton on a word is an infinite sequence of states. Deterministic Büchi automata have at most one run for each word in Σ^ω , whereas non-deterministic Büchi automata may have multiple runs for a word.

The complement of a Büchi automaton A is another Büchi automaton¹ denoted by \overline{A} . Both, A and \overline{A} , share the same alphabet Σ . Regarding a word $\alpha \in \Sigma^\omega$, the relation between an automaton and its complement is as follows:

¹Büchi automata are closed under complementation. This has been proved by Büchi [8], who, to this end, described the first Büchi complementation construction in history.

$$\alpha \text{ accepted by } A \iff \alpha \text{ not accepted by } \overline{A}$$

That is, all the words of Σ^ω that are *accepted* by an automaton are *rejected* by its complement, and all the words of Σ^ω that are *rejected* by an automaton are *accepted* by its complement. In other words, there is no single word of Σ^ω that is accepted or rejected by *both* of an automaton and its complement.

The difficulty of constructing the complement for a given Büchi automaton depends on whether this automaton is deterministic or non-deterministic. The complementation of *deterministic* Büchi automata is “easy” and can be done in polynomial time and linear space [26]. The complementation of *non-deterministic* Büchi automata, however, is very complex, and therefore subject of ongoing research. Note that whenever we talk about “Büchi complementation” in the following, we always mean specifically the complementation of *non-deterministic* Büchi automata.

The main problem with the complexity of Büchi complementation is the so-called state complexity (also called state growth, state blow-up or state explosion). This is the number of states of the output automaton in relation to the number of states of the input automaton. In particular, this means that the complement of a given Büchi automaton may be extremely large. This hinders the practical application of Büchi complementation, because the complementation operation may require very high time and computing resources. In the following we describe an important application of Büchi complementation that is negatively affected by the high state complexity of Büchi complementation.

1.1.2 Application of Büchi Complementation

In this section, we describe a specific application of Büchi complementation in formal verification. To this end, we proceed in two steps. First, we describe a general application of Büchi complementation, namely testing language containment of ω -regular languages. Next, we describe how language containment is used in a specific formal verification approach, which is called language containment approach to automata-theoretic model checking. Thus, these two descriptions are strung together, ultimately show how Büchi complementation is used in this model checking approach.

Note that for a comprehensive treatment of the interesting domains of model checking and other formal verification approaches, we refer to the works by Huth [19], Ben-Ari [5], and Baier [4].

Language Containment of ω -Regular Languages

Büchi complementation is used for testing language containment of ω -regular languages. The ω -regular languages are the class of formal languages that is equivalent to non-deterministic Büchi automata². At this point, we briefly describe the language containment in general, before in turn describing an application of the language containment problem in the next subsection.

Given two ω -regular languages L_1 and L_2 over alphabet Σ^ω the language containment problem consists in testing whether $L_1 \subseteq L_2$. This is true if all the words of L_1 are also in L_2 , and false if L_1 contains at least one word that is not in L_2 . The way this problem is algorithmically solved is by testing $L_1 \cap \overline{L_2} = \emptyset$. Here, $\overline{L_2}$ denotes the complement language of L_2 , which means $\overline{L_2}$ contains all the words of Σ^ω that are *not* in L_2 . The steps for testing $L_1 \cap \overline{L_2} = \emptyset$ are the following:

- Create the complement language $\overline{L_2}$ of L_2
- Create the intersection language $L_{1,\overline{2}}$ of L_1 and $\overline{L_2}$
- Test whether $L_{1,\overline{2}}$ is empty (that is, contains no words at all)

Thus, the language containment problem is reduced to three operations on languages, *complementation*, *intersection*, and *emptiness testing*. The common way to work with formal languages is not to handle the languages themselves, but more compact structures that represent them, such as automata. In the case of ω -regular languages, these are non-deterministic Büchi automata.

²Note that deterministic Büchi automata have a lower expressivity than non-deterministic Büchi automata, and are equivalent to only a subset of the ω -regular languages.

For solving $L_1 \subseteq L_2$, one thus works with two Büchi automata A_1 and A_2 that represent L_1 and L_2 , respectively. The problem then becomes $L(A_1) \subseteq L(A_2)$, and equivalently, $L(A_1) \cap \overline{L(A_2)} = \emptyset$. This is automata-theoretically solved as $\text{empty}(A_1 \cap \overline{A_2})$, which includes the three following steps:

- Construct the complement automaton $\overline{A_2}$ of A_2
- Construct the intersection automaton $A_{1,\overline{2}}$ of A_1 and A_2
- Test whether $A_{1,\overline{2}}$ is empty (that is, accepts no words at all)

If the final emptiness test on automaton $A_{1,\overline{2}}$ is true, then $L_1 \subseteq L_2$ is true, and if the emptiness test is false, then $L_1 \subseteq L_2$ is false. In this way, the language containment problem of ω -regular languages is reduced to three operations of *complementation*, *intersection*, and *emptiness testing* of non-deterministic Büchi automata. Thus, Büchi complementation is an integral part of language containment of ω -regular languages.

Language Containment Approach to Automata-Theoretic Model Checking

Above, we have seen that Büchi complementation is used for testing language containment of ω -regular languages. Now, we will see what language containment of ω -regular languages can be used for. To this end, we describe the language containment approach to automata-theoretic model checking

The language containment approach to automata-theoretic model checking is an approach to automata-theoretic model checking, which is an approach to model checking, which in turn is an approach to formal verification [62]. Figure 1.1 shows the branch of the family of formal verification techniques that contains the language containment approach to automata-theoretic model checking.

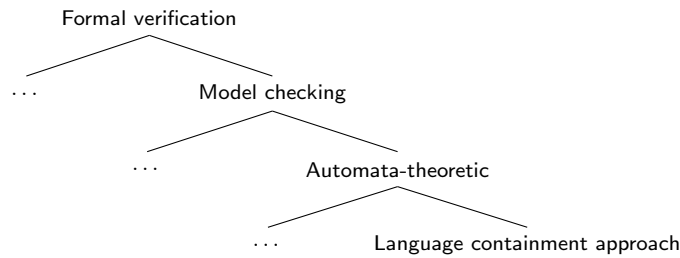


Figure 1.1: Branch of the family of formal verification techniques that contains the language containment approach to automata-theoretic model checking.

Formal verification is the use of mathematical techniques for proving the correctness of a system (software or hardware) with respect to a specified property [62]. A typical example is to verify that a program is deadlock-free (in which case the property would be “deadlock-freeness”). In general, formal verification techniques consist of the following three parts [19]:

1. A framework for modelling the system to verify
2. A framework for specifying the property to be verified
3. A verification method for testing whether the system satisfies the property

For the language containment approach to automata-theoretic model checking, the frameworks for points 1 and 2 are Büchi automata representing ω -regular languages. The verification method (point 3) is testing language containment of the system automaton’s language in the property automaton’s language. In some more detail, this works as follows. [61][62]

The system s to verify is modelled as a Büchi automaton, say S . This Büchi automaton represents an ω -regular language $L(S)$, and each word of $L(S)$ represents in turn a possible *computation trace* of the system. A computation trace is an infinite³ sequence of “situations” that the system is in at any point in

³The infinity of computation traces suggests that this type of formal verification (and model checking in general) is used for systems that are not expected to terminate and may run indefinitely. This type of systems is called *reactive* systems. They contrast with systems that are expected to terminate and produce a result. For this latter type of systems other formal verification techniques than model checking are used. See for example [19] and [5] for works that cover the formal verification of both types of systems.

time. Such a situation consists of a finite amount of information of, for example, the values of variables, registers, or buffers. The observation that such a trace can be represented as a word of an ω -regular languages comes from the fact that it can be represented intuitively as a linear Kripke structure (which in turn is an interpretation for a temporal logic formula that can also be used to represent computations), which in turn can be represented by a word of a language whose alphabet is ranges over the powerset of the atomic propositions of the Kripke structure. A work that explains these intimate relations between computation, temporal logic, formal languages, and automata in more detail is [61]. In simple words, the language $L(S)$, represented by the system automaton S , represents everything that the system *can* do.

Similarly, a property p to be verified is represented as a Büchi automaton, say P , which represents the ω -regular language $L(P)$, whose words represent computation traces. These computation traces are all the computations of a system like s that satisfy the property p . If for example p is “deadlock-freeness”, then the words of $L(P)$ represent all the possible computation traces that do *not* contain a deadlock. In this way, the language $L(P)$ represents everything that the system is *allowed* to do, with respect to a certain property.

The verification step is finally done by testing language containment of $L(S)$ in $L(P)$, that is, $L(S) \subseteq L(P)$. If this is true, then everything that the system *can* do is *allowed* to do, and the system satisfies the property p . If the language containment test is false, then the system *can* do a computation that is *not allowed* to do, and the system does not satisfy the property p .

Summarising, the language containment approach to automata-theoretic model checking requires language containment of ω -regular languages, which, as we have seen, requires Büchi complementation. In the following, we will highlight the particular importance of Büchi complementation for this type of formal verification.

1.1.3 Importance of Büchi Complementation

The high complexity of Büchi complementation makes the language containment approach to automata-theoretic model checking inapplicable in practice [50]. According to [62] and [63], there are so far no model checking tools that include the complementation of a property Büchi automaton. This is unfortunate, because the other operations for language containment test, intersection and emptiness testing of Büchi automata, have highly efficient solutions [61, 62]. Büchi complementation is thus a bottleneck. Existing model checking applications are forced to circumvent the need for Büchi complementation. In the following, we explain how this can be done, and why the approach via the complementation of a non-deterministic Büchi automaton is preferable.

One way to circumvent the complementation of non-deterministic Büchi automata is to specify the property as a deterministic Büchi automaton [50, 62]. As we have mentioned, the complementation of deterministic Büchi automata has an efficient solution. However, the disadvantage of this approach is that the property automaton may become exponentially larger than an equivalent non-deterministic automaton [50]. Furthermore, it is often more complicated and less intuitive to represent a property as a deterministic automaton [50].

Another to elude the need for Büchi complementation is to use a different model checking approach altogether. This leads us by the way back to the basic model checking question. In basic model checking, the property to be verified is represented as a formula φ of a temporal logic (typically LTL). The system to verify is represented as a Kripke structure K . Kripke structures are interpretations of temporal formulas. The verification step consists in checking whether K satisfies φ , or in other words, whether K is a model of φ (written as $K \models \varphi$). This check of modelhood is the reason why *model checking* is called this way [61].

The different model checking approaches differ in how they do the modelhood test. In the automata-theoretic approach, modelhood is tested in terms of automata. This is indeed possible *without* the need for Büchi complementation [62], as we explain in the following. The first step is to translate the Kripke structure K to a Büchi automaton A_K . Regarding the formula φ , the crucial point is that it is not directly translated to a Büchi automaton, but that it is first negated to $\neg\varphi$, and only then translated to a Büchi automaton $A_{\neg\varphi}$. This allows to do the language containment test $L(K) \subseteq L(\varphi)$ (which is equivalent to the modelhood test $K \models \varphi$), as $\text{empty}(A_K \cap A_{\neg\varphi})$, that is, without complementing a Büchi automaton.

This is because $A_{\neg\varphi}$ already represents the *negation* of the property, and thus it is not necessary to complement this automaton. In this way, the negation of the property, which is required for the language containment test, is pushed off from Büchi automata to the temporal logic formula, in which case it is easy. This approach is used by the SPIN model checker [17]. The disadvantage is that the typically used temporal logic, LTL, is less expressive than Büchi automata. Hence, specifying a property as an LTL formula restricts the set properties that can be used. It has even been stated that the expressivity of LTL is often insufficient for industrial applications [62].

The disadvantages of these alternative approaches highlight the importance of Büchi complementation for model checking. It shows that having more efficient Büchi complementation constructions, and to generally better understand the Büchi complementation problem, in order to handle it more adequately, would be of great practical value. In the next section, we motivate our approach to contribute to a better understanding of Büchi complementation.

1.2 Motivation

In the previous section we have seen that Büchi complementation is complex, and that it would be of great practical value to better understand it. In this section, we highlight the need for looking at the complexity of Büchi complementation in a way that has not received much attention in the past. This approach is based on empirical rather than theoretical investigations. In particular, it boils down to the question, how to assess the performance of a Büchi complementation construction? (Note that with *performance* we mean the state complexity that this construction exhibits. In the following, we will use the terms performance and state complexity interchangeably.) Studying the behaviour of complementation constructions is a way to better understand Büchi complementation itself.

In Section 1.2.1, we present the traditional way of assessing the performance of complementation constructions, which is based on theoretical rather than empirical grounds. Then, in Section 1.2.2, we present the empirical approach, which this thesis is following, and motivate its need. This includes a review of the research that has been done to date in this direction.

1.2.1 Theoretical Investigation of Worst-Case Performance

The traditional way to assess the performance of Büchi complementation constructions is to identify their *worst-case* state complexity. This is the state complexity a construction exhibits for a theoretical worst-case automaton. In other words, it means the maximum number of states a construction can generate.

For example, the original complementation construction by Büchi [8] has a worst-case state complexity of $2^{2^{O(n)}}$. At this point, two remarks. First, a state complexity is always given as a function of the number of states of the input automaton. Throughout this thesis, we will always denote this parameter by n . Second, the state complexity often includes the big- O notation, thus it is not an exact function. We will sometimes assume concrete functions by simply stripping big- O notation. This worst-case state complexity of Büchi's construction means that, given an input automaton with n states, the complement cannot have more than most $2^{2^{O(n)}}$ states. For example, if the input automaton has 8 states, then the upper bound for the number of states of the complement is 1.16×10^{77} (if assuming the function 2^{2^n}).

Different constructions exhibit different worst-case state complexities, and one of the main objectives of construction developers is to reduce this number. For example, the more recent construction that Schewe proposed in 2009 [45] has a worst-case state growth of only $(0.76n)^n + n^2$. For an input automaton with 8 states, this means that the complement has at most 119.5 million states. Clearly this is a tremendously smaller number than the maximum number of states of Büchi's construction.

A related track of research is concerned with finding general lower bounds for these upper bounds of the number of states that constructions can generate. This can be done by proving that there are certain automata whose complements have *at least* a certain number of states. Consequently, no complementation construction can have a worst-case state complexity that is below this number. A first such lower bound of $n!$ has been proposed in 1988 by Michel [30]. He proved that there exists a family of automata whose complement have at least $n!$ states. These automata are known as Michel automata. Later, in 2007,

Yan [68] proved that there are automata whose minimum complement size is even larger than $n!$, namely $(0.76n)^n$ (Michel’s $n!$ corresponds to approximately $(0.36n)^n$ [68]). This means that $(0.76n)^n$ replaced $n!$ as a new lower bound for the worst-case state complexity of Büchi complementation. This lower bound is valid until today.

Generally, construction developers aim at brining the worst-case state complexity of their construction close to this lower bound. For example, Schewe’s construction is very close to the lower bound of $(0.76n)^n$ [45]. If the worst-case state complexity of a construction matches this lower bound, it is said to be *optimal*. Note that the worst-case state complexity of a specific construction constitutes an upper bound for the general worst-case state complexity of Büchi complementation. This is because the currently valid lower bound cannot rise higher than the worst-case state complexity of a known construction. Therefore, new constructions with lower worst-case state complexities can be seen as narrowing the gap between the upper and lower bound of the exact state complexity bound of Büchi complementation.

1.2.2 Empirical Investigation of Actual Performance

Worst-case state complexities are interesting from a theoretical point of view. However, they are not necessarily good guides to the actual performance of a construction [54]. For example, if we have a concrete automaton of, say, 15 states, and we complement it with Schewe’s construction, the fact that the worst-case state complexity is $(0.76n)^n n^2$ reveals that the complement has in the worst case 1.6 quintillion (1.6×10^{18}) states. However, in a practical case, we do most probably not anticipate the worst case, but we are rather interested in the performance on a more likely “average case”.

In a practical application of Büch complementation, we could be interested in questions about complementation constructions as follows. What is a reasonable complement size to expect for the given automaton with n states? Are there generally easier and harder automata? What are the factors that make an automaton especially easy or hard? How does the performance of different constructions on the same automata vary? Are there constructions which are better suited for a certain type of automata than other constructions?

Questions like this cannot be answered by theoretical investigations of the worst-case state complexity of the constructions. However, they can be attempted to answer by empirical performance investigations of these constructions. With empirical performance investigations, we mean the testing of a construction with concrete automata, and the analysis of the results that the construction produces. Naturally, such investigations are based on an *implementation* of the investigated constructions and on *test data*, that is, the set of automata with which the constructions are tested.

There have been relatively few empirical investigations [54], compared to the number of theoretical works. What follows is a short review some of these empirical works that have been done in the past. This illustrates the empirical approach, and presents the line of research in which this thesis is situated.

Note that in the following review we mention many complementation constructions. All of these constructions will be presented in our review of Büchi complementation constructions in Chapter 2. Furthermore, we indicate whether a construction is Ramsey-based, determinisation-based, rank-based, or slice-based. These are the main complementation approaches that we also explain in Chapter 2.

1995 Tasiran et al. [50] create an efficient implementation of Safra’s construction [43] (determinisation-based) and used it for automata-theoretic model checking tasks with the HSIS verification tool [3]. They state that they could easily complement property automata with some hundreds of states, however, they do not provide a statistical evaluation of the results.

2003 Gurumurthy et al. [16] implement Friedgut, Kupferman, and Vardi’s construction [25] (rank-based) along with various optimisations that they propose as a part of the tool Wring [48]. They complement 1000 small automata, generated by translation from LTL formulas, and evaluate execution time, and number of states and transitions of the complement for the different versions of the construction.

2006 Althoff et al. [2] implement Safra’s [43] and Muller and Schupp’s [35] determinisation constructions⁴ in a tool called OmegaDet, applied them on the Michel automata with 2 to 6 states, and compared

⁴These determinisation constructions transform a non-deterministic Büchi automaton to a deterministic Rabin automaton (see Section ??), however, they are used as the base for determinisation-based complementation constructions.

the number of states of the determinised output automata.

- 2008** Tsay et al. [57] carry out a first comparative experiment with the publicly available⁵ GOAL tool [56][57][58][55]. They include the constructions by Safra [43] (determinisation-based), Piterman [38] (determinisation-based), Thomas [53] (WAPA⁶), and Kupferman and Vardi[25] (rank-based or WAA⁷). These constructions are pre-implemented in GOAL. As the test data, they use 300 Büchi automata, translated from LTL formulas, with an average size of 5.4 states. They evaluate and compare execution times, as well as number of states and transitions of the complements.
- 2009** Kamarkar and Chakraborty [21] propose an improvement of Schewe’s construction [45] (rank-based) and implement it, as well as Schewe’s original construction, on top of the model checker NuSMV [10][9]. They run the constructions on 12 test automata and compare the sizes of the complements. Furthermore, they run the same tests with the constructions by Kupferman, and Vardi [25] (rank-based or WAA) and Piterman [38] (determinisation-based) within GOAL, and compare the results to the ones of their implementation of Schewe’s construction.
- 2010** Tsai et al. [54] (paper entitled “State of Büchi Complementation”) carry out another experiment with GOAL. They compare the constructions by Piterman [38] (determinisation-based), Schewe [45] (rank-based), and Vardi and Wilke [65] (slice-based), with various optimisations that they propose in the same paper. As the test data, they use 11,000 randomly generated automata with 15 states and an alphabet size of 2. The test set is organised into 110 automata classes that consist of the combinations of 11 transition densities and 10 acceptance densities. This test set is repeatedly used in subsequent work (including in this thesis), and we will refer to it as the GOAL test set (because it has been generated with the GOAL tool). Tsai et al. provide sophisticated evaluation of the states of the complements for all the tested constructions and construction versions.
- 2010** Breuers [6] proposes an improvement for the construction by Sistla, Vardi, and Wolper [47] (Ramsey-based), and creates an implementation of it. He generates his own test data (inspired by the work of Tsai et al. [54]) consisting of *easy*, *medium*, and *hard* automata, based on different transition density and acceptance density values. He evaluates the complement sizes produced by the construction for automata of sizes 5, 10, and 15 of all these difficulty categories.
- 2012** Breuers et al. [7] wrap the implementation of their improvement of Sistla, Vardi, and Wolper’s construction [47] in the publicly available tool Alekto⁸, and run it on the GOAL test set. They compare the generated complement sizes, as well as the number of aborted complementation tasks (due to exceeding resource requirements) to the corresponding result for different constructions on the same test set by Tsai et al. [54].
- 2013** Göttel [14] creates a C implementation of the Fribourg construction [1], including the R2C optimisation (see Chapter 3), and executes it on the GOAL test set, as well as on the Michel automata with 3 to 6 states. He analyses the resulting complement sizes and execution times separately for each of the 110 classes that the GOAL test set consists of. The Fribourg construction⁹ is a slice-based complementation construction that is being developed at the university of Fribourg, and which lies at the heart of this thesis. The entire Chapter 3 of this thesis is dedicated to explaining the Fribourg construction.

1.3 Aim and Scope

The aim of this thesis is to contribute to a better understanding of the behaviour of Büchi complementation constructions. The way we want to achieve this aim is by investigating the behaviour of one specific construction.

⁵<http://goal.im.ntu.edu.tw/wiki/doku.php>

⁶Via **W**eak **A**lternating **P**arity **A**utomaton

⁷Via **W**eak **A**lternating **A**utomaton

⁸<http://www.automata.rwth-aachen.de/research/Alekto/>

⁹The authors of the constructions use the name *subset-tuple construction* (see [1]), however, in this thesis, we will use the name *Fribourg construction*.

The aim of this thesis is an in-depth empirical performance investigation of the Fribourg construction. As mentioned, the Fribourg construction is a Büchi complementation construction that is being developed at the University of Fribourg [1]. By empirically investigating the behaviour of this specific construction, we want to follow up the track of empirical research that we have outlined in the last section.

This thesis is certainly not sufficient to describe the performance of the Fribourg construction in its entirety, or in a way that is adequate to be relied on in industrial applications. Neither this thesis can answer general questions about the observed behaviour of Büchi complementation. Rather, we see this piece of work as a mosaic stone that we add to the very complex and multi-faceted picture of the complexity of Büchi complementation.

The empirical performance investigation will include testing of different versions of the construction, and comparison with other complementation constructions...

Aim: empirical performance investigation of a specific Büchi complementation construction, comparison with other constructions

Scope: two test sets, relatively small automata, no real world or “typical” examples,

1.4 Overview

Chapter 2

Background

Contents

2.1	Büchi Automata and Other Omega-Automata	13
2.1.1	Büchi Automata	13
2.1.2	Other Omega-Automata	15
2.2	Run Analysis of Non-Deterministic Automata	16
2.2.1	Run DAGs	17
2.2.2	Run Trees	17
2.2.3	Split Trees	18
2.2.4	Reduced Split Trees	19
2.3	Review of Büchi Complementation Constructions	20
2.3.1	Ramsey-Based Approach	21
2.3.2	Determinisation-Based Approach	21
2.3.3	Rank-Based Approach	23
2.3.4	Slice-Based Approach	24

IN THIS CHAPTER we treat several topics that serve as a background for the rest of the thesis. In particular, the goal of this chapter is to set the stage for our description of the Fribourg construction in Chapter 3, and the setup of our empirical performance investigation of the Fribourg construction in Chapter 4, as well as its results in Chapter 5.

In Section 2.1 we summarise some aspects about Büchi automata, as well as about other types of automata on infinite words, that are relevant to the purpose of Büchi complementation. In Section 2.2, we describe the principal techniques for run analysis of non-deterministic automata. This topic has a particular relation to Büchi complementation, and one of the run analysis techniques, reduced split trees, is the core of the Fribourg construction. In Section 2.3, we provide a review of proposed Büchi complementation constructions from the introduction of Büchi automata in 1962 until today. We organise the presentation of these constructions along the four main Büchi complementation approaches Ramsey-based, determinisation-based, rank-based, and slice-based. Some of the constructions that we describe in this section will be used in the performance investigation of the Fribourg construction in Chapter 4.

2.1 Büchi Automata and Other Omega-Automata

Büchi automata are a type of ω -automata. These are finite state automata that run on infinite words (so-called ω -words). Externally, ω -automata look the same as the traditional finite state automata on finite words. It is possible to interpret any such automaton on finite words as an ω -automaton, and vice versa.

The difference between ω -automata and automata on finite words is their acceptance condition. An automaton on finite words accepts a word, if after finishing reading it, the automaton is in an accepting state. For ω -automata, this acceptance condition is not possible, because an ω -automaton never finishes reading a word (because the word “never ends”). Instead, the acceptance condition of ω -automata is defined on the set of the so-called *infinitely recurring states*. We are going to describe this concept in Subsection 2.1.1 below.

In Subsection 2.1.1 of this section, we first treat Büchi automata, and in Subsection 2.1.2 the principal other types of ω -automata, in particular, Muller, Rabin, Streett, and parity automata. In the latter subsection, we also introduce a shorthand notation for different types of ω -automata that we will use throughout the thesis.

Note that in the entire section we omit overly formal notation and proofs of any kind, because it is not necessary for the aim of this thesis. More comprehensive and formally rigorous treatments of ω -automata can be found, for example, in the works by Thomas [51, 52], or Wilke [66].

2.1.1 Büchi Automata

Below we summarise the aspects of Büchi automata that are most significant for our purposes, including the acceptance condition of Büchi automata, the expressivity of non-deterministic and deterministic Büchi automata, and basics about the complementation of non-deterministic and deterministic Büchi automata. In the course of this, we always stress the difference between deterministic and non-deterministic Büchi automata, as this is one of the main sources for the intricacy of Büchi complementation [36].

Definition and Acceptance Condition

A non-deterministic Büchi automaton A is defined by the 5-tuple $A = (Q, \Sigma, \delta, q_0, F)$ with the following components:

- Q : a finite set of states
- Σ : a finite alphabet
- δ : a transition function, $\delta : Q \times \Sigma \rightarrow 2^Q$
- q_0 : an initial state, $q_0 \in Q$

- F : a set of accepting states, $F \subseteq Q$

Note that this is the same definition as for non-deterministic finite state automata on finite words [18]. The difference between Büchi automata and automata on finite words is only the acceptance condition of Büchi automata that we describe below. For deterministic Büchi automata, the definition is similar to the above one, but with a different transition function δ that returns none or a single state, instead of a set of states

An important concept in automata theory is the notion of a *run* of an automaton on a given word. A run ρ of automaton A on word α is a sequence of states $q \in Q$ that A visits in the process of reading α . For finite words, the length of a run is finite as well, however, for ω -words, the length of a run may be infinite. Note that deterministic automata have at most one run for a given word, whereas non-deterministic automata may have multiple possible runs for the same word.

The Büchi acceptance condition decides whether a run ρ is accepting or non-accepting. This in turn determines the acceptance or non-acceptance of a word: a word α is accepted by an automaton A , if and only if it has at least one accepting run in A . The decision whether a run ρ is accepting or non-accepting is based on the set of *infinitely recurring states* of ρ that we denote by $\text{inf}(\rho)$. This set contains all the states that occur infinitely often in ρ . In particular, the Büchi acceptance condition is as follows:

$$\text{Run } \rho \text{ is accepting} \iff \text{inf}(\rho) \text{ contains at least one accepting state}$$

That is, a run is accepting, if and only if the run contains at least one accepting state infinitely often. Formally, this can be written as $\text{inf}(\rho) \cap F \neq \emptyset$.

An intuitive way for describing the Büchi acceptance condition has been given by Vardi [61]. If we imagine the automaton having a green light that blinks whenever the automaton visits an accepting state, then the run is accepting if we observe the green light blinking infinitely many times. The fact that there are only finitely many accepting states, but the light blinks infinitely many times, implies that at least one accepting state is being visited infinitely often.

Expressivity

A particularity of Büchi automata is that deterministic and non-deterministic automata are *not* expressively equivalent. In particular, the class of languages corresponding to the deterministic Büchi automata is a strict subset of the class of languages corresponding to the non-deterministic automata. This result has been proved by Büchi himself in his 1962 paper [8].

This contrasts, for example, with finite state automata on finite words. In this case, non-deterministic and deterministic automata are expressively equivalent, and consequently every non-deterministic automaton can be turned into an equivalent deterministic automaton. With Büchi automata, however, this is not possible, because there exist languages that can be expressed by a non-deterministic Büchi automaton, but not by a deterministic one. An example of such a language is $(0 + 1)^* 1^\omega$, that is, the language of all words of 0 and 1 ending with an infinite sequence of 1. A non-deterministic automaton representing this language, cannot be turned into an equivalent deterministic automaton [61, 42]. Because of this fact, we say that Büchi automata can *in general* not be determinised to other Büchi automata. This fact has implications on the complementation of non-deterministic Büchi automata, as we will see below.

The class of languages that is equivalent to the *non-deterministic* Büchi automata is the class of *ω -regular languages*. A formal description of the ω -regular languages can be found, for example, in [51, 52, 66]. Regarding deterministic Büchi automata, consequently the set of languages that is equivalent to them is a strict subset of the ω -regular languages.

Complementation

Non-deterministic Büchi automata are closed under complementation. This means that the complement of every non-deterministic Büchi automaton is another non-deterministic Büchi automaton. This result

has been proved by Büchi in his 1962 paper [8]¹. Deterministic Büchi automata, on the other hand, are not closed under complementation [51]. In particular, this means that the complement of a deterministic Büchi automaton is still a Büchi automaton, however, possibly a non-deterministic one.

As we already mentioned, the algorithmic difficulty and complexity of complementation is very different for deterministic and non-deterministic automata. For deterministic Büchi automata, there exists a simple procedure, introduced in 1987 by Kurshan [26], that can complement a deterministic Büchi automaton to a non-deterministic Büchi automaton in polynomial time and linear space.

For non-deterministic Büchi automata, however, there exists no easy solution. The main reason is that Büchi automata can in general not be determinised. If they could be determinised, then a solution would be to transform a non-deterministic Büchi automaton to an equivalent deterministic one, and complement the deterministic Büchi automaton with Kurshan's construction. This is by the way the approach that is used for the complementation of non-deterministic automata on finite words: determinise a non-deterministic automaton with the subset construction [39], and then trivially complement the deterministic automaton by making the accepting states non-accepting, and vice versa. Unfortunately, for Büchi automata such a simple procedure is not possible, and this can be seen as the main reason that Büchi complementation is such a hard problem [36].

2.1.2 Other Omega-Automata

After the introduction of Büchi automata in 1962, several other types of ω -automata have been introduced. These automata differ from Büchi automata only in their acceptance condition, that is, in the way they decide whether a run ρ is accepting or non-accepting. All these acceptance condition are however based on the set of infinitely recurring states $\text{inf}(\rho)$ of ρ . The following are the most important of these alternative ω -automata along with the year of their introduction.

- Muller automata (1963) [33]
- Rabin automata (1969) [40]
- Streett automata (1982) [49]
- Parity automata (1985) [32]

Some of these automata types are used in complementation constructions, especially in determinisation-based complementation constructions (see Section 2.3). Table 2.1 lists the Muller, Rabin, Streett, and parity acceptance conditions, along with the Büchi acceptance condition for comparison.

Type	Definitions	Condition
Muller	$U \subseteq 2^Q$	$\text{inf}(\rho) \in U$
Rabin	$\{(U_1, V_1), \dots, (U_r, V_r)\}, U_i, V_i \subseteq Q$	$\exists i : \text{inf}(\rho) \cap U_i = \emptyset \wedge \text{inf}(\rho) \cap V_i \neq \emptyset$
Streett	$\{(U_1, V_1), \dots, (U_r, V_r)\}, U_i, V_i \subseteq Q$	$\forall i : \text{inf}(\rho) \cap U_i \neq \emptyset \vee \text{inf}(\rho) \cap V_i = \emptyset$
Parity	$\pi : Q \rightarrow \{1, \dots, k\}, k \in \mathbb{N}$	$\min(\{\pi(q) \mid q \in \text{inf}(\rho)\}) \bmod 2 = 0$
Büchi		$\text{inf}(\rho) \cap F \neq \emptyset$

Table 2.1: Acceptance conditions of Muller, Rabin, Streett, parity, and Büchi automata. Note that Q denotes the set of states and F the set of accepting states of the corresponding automaton.

Below, we briefly explain the acceptance conditions of Muller, Rabin, Streett, and parity automata in words. More detailed descriptions of these conditions can be found, for example, in [28], or [66].

Muller acceptance condition Includes a set U of subsets of states. A run ρ is accepting if and only if $\text{inf}(\rho)$ equals one of the pre-defined subsets in U . The Muller acceptance condition is the most general one, and the Rabin, Streett, parity, and Büchi acceptance conditions can be expressed as Muller acceptance conditions [28, 66].

¹Actually, the proof of closure under complementation of non-deterministic Büchi automata was a necessary step for Büchi to prove the equivalence of Büchi automata and S1S formulas (Büchi's Theorem). In order to prove this closure under complementation, Büchi described the first Büchi complementation construction in history.

Rabin acceptance condition Includes a list of pairs (U, V) where U and V are subsets of states. A run ρ is accepting if and only if *there is* at least one pair for which U does not contain any states of $\text{inf}(\rho)$ and V contains at least one state of $\text{inf}(\rho)$. A pair (U, V) is called a Rabin pair.

Streett acceptance condition Includes a list of pairs (U, V) where U and V are subsets of states. A run ρ is accepting if and only if *for all* the pairs either U contains at least one state of $\text{inf}(\rho)$ or V does not contain any states of $\text{inf}(\rho)$. A pair (U, V) is called a Streett pair. Note that the Streett condition is the dual of the Rabin condition. This means that if we have two identical Rabin and Streett automata with an identical list of pairs, then the Streett automaton accepts the complement language of the Rabin automaton [23].

Parity acceptance condition Assigns a number to each of the states of the automaton. A run ρ is accepting if and only if the smallest-numbered element of $\text{inf}(\rho)$ has an even number. The numbers that are assigned to the states are sometimes called colours [29].

Regarding the expressivity of these automata, it turns out that they are all equivalent to non-deterministic Büchi automata, and thus to the ω -regular languages [66]. This holds for non-deterministic *and* deterministic automata of these types. That means that, unlike Büchi automata, deterministic and non-deterministic Muller, Rabin, Streett, and parity automata are expressively equivalent.

At this point we introduce a notation that we will occasionally use for denoting different types of ω -automata. This notation has been used by Piterman [37] and later by Tsai et al. [54]. It consists of three-letter acronyms of the form

$$\{N, D\} \times \{B, M, R, S, P\} \times W$$

The first letter specifies whether the automaton is non-deterministic (N) or deterministic (D). The second letter stands for the acceptance condition: B for Büchi, M for Muller, R for Rabin, S for Streett, and P for parity. The last letter specifies on which structure the automaton runs. In our case these are always words, thus the last letter is always W . For example, NBW means non-deterministic Büchi automaton, DBW means deterministic Büchi automaton, NMW means non-deterministic Muller automaton, DMW means deterministic Muller automaton, and so on.

2.2 Run Analysis of Non-Deterministic Automata

As mentioned, in a deterministic automaton, every input word has at most one run. In a non-deterministic automaton, however, every input word may have multiple runs. The analysis of the different runs of a non-deterministic automaton on a given input word is called *run analysis* and is central to the Büchi complementation problem.

The reason that run analysis is central to Büchi complementation is as follows. In a Büchi complementation task, we are given a non-deterministic Büchi automaton A , and we attempt to construct its complement B . For constructing B we have in fact to decide for every word α whether B must accept it or not. This decision is intrinsically tied to the entirety of the runs of A on α in the following way:

$$B \text{ accepts } \alpha \iff \text{All the runs of } A \text{ on } \alpha \text{ are non-accepting}$$

B must accept a word α , if and only if *all* the runs of A on α are non-accepting. If, for example, A has 10 runs on α , and 9 of them are non-accepting and one is accepting, then A still accepts the word α , and consequently, B must not accept it. Only if all of the 10 runs of A on α are non-accepting, A does not accept α , and consequently B must accept α . This means that for constructing the complement B , we need to consider *all* the runs of the input automaton A on specific words, and the way this can be done is by run analysis on A .

In this section, we present the principal run analysis techniques for non-deterministic Büchi automata [66]. We start with *run DAGs* (DAG stands for directed acyclic graphs) in Section 2.2.1. Then in Sections ??, 2.2.3, and 2.2.4, we present three techniques based on trees with increasing sophistication, namely *run trees*, *split trees*, and *reduced split trees*. The last of the tree techniques, reduced split trees, lies at the heart of the Fribourg construction that we describe in Chapter 3.

In the following subsections, we will give examples for the different run analysis techniques that are based on the non-deterministic Büchi automaton in Figure 2.1. Note that the alphabet of this automaton is $\Sigma = \{a\}$, and thus the only ω -word in Σ^ω is a^ω . However, the automaton has multiple (infinitely many) runs for this word.




Figure 2.1: Non-deterministic Büchi automaton A used as example automaton in this section.

2.2.1 Run DAGs

Run DAGs arrange all the runs of an automaton A on a word in a directed acyclic graph (DAG). This graph can be thought of as a matrix with rows and columns. The rows are called levels, and each column corresponds to a state of A . Consequently, the width (that is, number of columns) of a run DAG equals the number of states of A . Each level i (starting from 0) corresponds to the situation after reading the first i symbols of the word. Figure 2.2 shows the first five levels of the run DAG of the example automaton A in Figure 2.1 on the word a^ω .




Figure 2.2: First few levels of the run DAG for the runs of automaton A (Figure 2.1) on the word a^ω .



Note that throughout this thesis, we are using rectangles with rounded corners for states of automata, and rectangles with sharp corners for vertices of graphs nodes of trees. In graphs and trees, we will however indicate vertices or nodes that correspond to accepting automaton states with double-lined rectangles.

As can be seen in Figure 2.2, every path of a run DAG corresponds to a run of the automaton on the given word. A run DAG is a structure that is able to represent an infinite number of runs by keeping a finite width. The number of levels of a run DAG is infinite for ω -words. A formal description of run DAGs can be found, for example, in [11].

Run DAGs are the basic structure for the rank-based complementation constructions, that we review in Section 2.3.3. Regarding this application, the fact that run DAGs have finite width is important, because the levels of run DAGs are mapped to states of the output automaton in these complementation constructions.

2.2.2 Run Trees

Trees in general are, after DAGs, the second structure that is used for run analysis. Run trees are the most basic of these tree structures. A run tree is basically a direct unwinding of all the runs of an automaton on a word as a tree. Figure 2.3 shows the first five levels of the run tree for the automaton A in Figure 2.1 on the word a^ω .


 Figure 2.3: First five levels of the run tree for the runs of automaton A (Figure 2.1) on the word a^ω .

 Figure 2.4: First five levels of the split tree for the runs of automaton A (Figure 2.1) on the word a^ω .

As can be seen in Figure 2.3, there is a one-to-one mapping of paths in a run tree (from the root to the leaves) to runs of the corresponding automaton. This means that if the automaton has an infinite number of runs on a given word, then the corresponding run tree has an infinite maximum width. This makes run trees impractical to be used in Büchi complementation. The following tree techniques, split trees and reduced split trees, sacrifice a part of the information about individual runs, with the benefit of making the tree structure more compact.

2.2.3 Split Trees

A split tree is basically a run tree where the accepting and non-accepting children of every node are merged. Consequently, a node of a split tree does not represent a single state of the automaton, but a set of states. The reduction of the number of children to at most two, makes the split tree furthermore a binary tree. Figure 2.4 shows the first five levels of the split tree of the automaton A (Figure 2.1) on the word a^ω .

Split trees sacrifice a part of the information about individual runs. For example, in Figure 2.4, we see that there must be a run that starts in q_0 (root node) and is in q_0 after reading four symbols (leftmost node on fifth level). However, the split tree does not provide the exact sequence of states. All that it reveals is that the sequence is $q_0 q_1^1 q_2^2 q_2^3 q_0$, where each q_i^j is either q_0 or q_2 . The run tree in Figure 2.3, on the other hand, shows the actual sequence unambiguously, which is $q_0 q_0 q_0 q_0 q_0$.

However, the split tree still provides an important piece of information, namely that the sequence $q_0 q_1^1 q_2^2 q_2^3 q_0$, whatever it actually looks like, does not contain any accepting states, because q_i^j can only be q_0 or q_2 , which are both non-accepting. It turns out that, with regard to the application of run analysis to Büchi complementation, this information is sufficient.

Split trees in fact embody a modified subset construction that does not mix accepting and non-accepting states. The result of applying such a construction to a non-deterministic automaton is an equivalent non-deterministic automaton whose degree of non-determinism is reduced to two (which means that each state has at most two successors for a given alphabet symbol). Such a construction has been described in [60].

A formal description of split trees can be found, for example, in [65] or [11]. Split trees are clearly more compact than run trees. However, their width may still grow infinitely. In the next section, we present a

further reduction of the information contained in the tree, that bounds the maximum width of the tree to a finite number.

2.2.4 Reduced Split Trees

Reduced split trees are split trees that sacrifice even more information. However, the information that is retained is still sufficient for Büchi complementation. The rule is that on each level, going from right to left or from left to right, only the first occurrence of each state is kept, and the others are omitted. This bounds the maximum width of the tree to the number of states of the corresponding automaton.

The direction in which the first occurrence of a state is determined depends on whether right-to-left or left-to-right split trees are used. In *right-to-left split trees*, the accepting child is put to the right of the non-accepting child. The split tree in Figure 2.4 is thus a right-to-left split tree. In *left-to-right split trees*, on the other hand, the accepting child is put to the left of the non-accepting child. In the literature both versions are used. For example, Vardi and Wilke [65] describe the left-to-right version, whereas Fogarty et al. [11] describe the right-to-left version. In this thesis, we will use exclusively the right-to-left version.

In a right-to-left reduced split tree, the levels are processed from right to left. This means that only the rightmost occurrence of each state on a level is kept and the others are omitted. Figure 2.5 shows the first five levels of the reduced split tree of the automaton A (Figure 2.1) on the word a^ω .




Figure 2.5: First five levels of the reduced split tree for the runs of automaton A (Figure 2.1) on the word a^ω .

As can be seen in Figure 2.5, each level contains at most one occurrence of each of the states q_0 , q_1 , and q_2 . Comparing the reduced split tree in Figure 2.5 with the split tree in Figure 2.4 reveals which states have been omitted in the reduced split tree. The root level and level 1 are similar in both trees. On level 2, the state q_2 in the leftmost node is omitted. This is because there is already a q_2 farther to the right, namely in the rightmost node of level 2. In level 3, there are two omissions of q_2 for the same reason. One of them causes an entire node to disappear, because this node contained q_2 as its only state. Finally, on level 4, there are three omissions of q_2 .

In the following, we illustrate what this omission of states entails and why the resulting “truncated” run analysis is still usable for Büchi complementation. By omitting states, we obviously omit runs from the tree. This omission of runs is targeted at so-called *re-joining runs*. Re-joining runs are runs that after a certain number of steps end up in the same state again. For example, the automaton from Figure 2.1 has seven re-joining runs that after four steps end up in the state q_2 . This can be easily seen in the corresponding run tree in Figure 2.3. A reduced split tree keeps exactly one of a group of re-joining runs and removes the others. This can be seen in Figure 2.5 which contains only one run from the root to q_2 in four steps.

The run that is kept is named *rightmost run*. The name refers to the fact that this run is the first of a group of re-joining runs that makes a *right turn* in the corresponding path of the tree. A right turn means the transition from a parent to an accepting child, which are situated to the right of the non-accepting child, hence the name². The rightmost run of a group of re-joining runs has the following crucial property:

Any re-joining run is accepting \implies The rightmost re-joining run is accepting

²Note that everything we explain here is also valid for left-to-right split trees, but one has to exchange *right*, *rightmost*, and *right turn* with *left*, *leftmost*, and *left turn*, respectively.

That is, if a group of re-joining runs contains an accepting run, then the rightmost run is accepting too. On the other hand, if the rightmost run is non-accepting, then all the other re-joining runs are non-accepting as well. A proof of this relation can be found, for example, in [65]³.

Practically, this means that we can reduce the existential question about the presence of acceptance in a group of runs to the rightmost run. Remember that for Büchi complementation we are interested in exactly such an existential question: is there an accepting run of the automaton A on the word α ? If yes (it does not matter how many accepting runs there are), then A accepts α and the complement must not accept it. If no, then A does not accept α and the complement must accept it. By considering only the rightmost runs, we can still reliably answer this question.

Thus, reduced split trees include a tremendous reduction of the number of runs to be analysed, but still retain all the needed information for Büchi complementation. Most importantly, this reduction of runs bounds the width of the trees to a finite number, what makes them usable in Büchi complementation constructions, such as the Fribourg construction (see Chapter 3), and other slice-based constructions (see Section 2.3.4).

Note that an alternative name for rightmost run is *greedy run* [1]. This name refers to the fact that the rightmost run is the first of a group of re-joining runs that visits an accepting state, what makes it “greedy”. It is interesting to note that reduced split trees are related to Muller-Schupp trees that are used in the Büchi determinisation construction by Muller and Schupp [35] (cf. [65, 11]). As a last remark, formal descriptions of reduced split trees can be found, for example, in [65] and [66] (left-to-right version), or [11] (right-to-left version).

This concludes our presentation of run analysis techniques. In the next section we will review the most prominent Büchi complementation constructions that have been proposed over the past 50 years. Many of them are based on the run analysis techniques that we have described in the present section.

2.3 Review of Büchi Complementation Constructions

Since the introduction of Büchi automata in 1962, many different Büchi complementation constructions have been proposed. In this section, we will briefly review some of the most important of them. Many of the constructions that we describe in this section are implemented in the GOAL tool, that we use for the empirical performance investigation of the Fribourg construction (see Chapter 4). Three of them are actively used in our investigation as a comparison against the Fribourg construction, namely the constructions by Piterman [37, 38], Schewe [45], and Vardi and Wilke [65].

We organise our review of Büchi complementation constructions into the four approaches that have been proposed by Tsai et al. [54]: *Ramsey-based*, *determinisation-based*, *rank-based*, and *slice-based*. As a compact reference, we summarise below all the constructions that are included in the following review under their respective approach, and with their authors, year of introduction, and references.

Ramsey-based approach		
Büchi	1962	[8]
Sistla, Vardi, and Wolper	1987	[46, 47]
Determinisation-based approach		
Safra	1988	[43, 44]
Muller and Schupp	1995	[35]
Piterman	2007	[37, 38]
Rank-based approach		
Klarlund	1991	[22]
Kupferman and Vardi	1997/2001	[24, 25]
Thomas	1999	[53]
Friedgut, Kupferman, and Vardi	2006	[12, 13]
Schewe	2009	[45]

³Lemma 2.6

Slice-based approach

Vardi and Wilke	2007	[65]
Kähler and Wilke	2008	[20]

2.3.1 Ramsey-Based Approach

The Ramsey-based complementation approach is the oldest approach and includes the earliest Büchi complementation constructions, including the construction described by Büchi himself [8]. Ramsey-based constructions are based on a branch of combinatorics called Ramsey theory [15], hence the name Ramsey-based approach⁴. They proceed by directly constructing the complement automaton out of the input automaton by applying combinatorial arguments.

Büchi (1962)

The first Büchi complementation construction in history, described by Büchi in his 1962 paper [8], is based on a combinatorial argument introduced by Ramsey in 1930 [41]. The construction is rather complicated [64], and has a doubly-exponential worst-case state growth of $2^{2^{O(n)}}$ [63]. This blow-up is enormous, for example, given an automaton with 9 states, the maximum size of the complement is 1.3×10^{154} states, tremendously more as there are atoms in the observable universe⁵.

However, as mentioned in our discussion about the use of worst-case state growth as a performance measure for complementation constructions in Section 1.2, this number applies only to the worst case, which is a very extraordinary special case. Furthermore, we believe that Büchi's aim was not to create an *efficient* construction, but rather to prove that such a construction exists in the first place. Because for Büchi this served as a proof that Büchi automata are closed under complementation, which was in turn required to prove Büchi's Theorem [65].

Sistla, Vardi, and Wolper (1987)

In 1987, Sistla, Vardi, and Wolper proposed an improvement of Büchi's construction [47]⁶. This construction was the first to include only an exponential, rather than doubly-exponential, worst-case state growth. The state growth function of the construction is $O(16^{n^2})$ [47]. Thus, given an automaton with 9 states, the maximum size of the complement is 3.4×10^{97} states.

2.3.2 Determinisation-Based Approach

The determinisation-based approach achieves complementation by chaining several conversions between different types of ω -automata. The most important of these conversions is the transformation of the non-deterministic Büchi automaton to a deterministic ω -automaton of a different type. Thus, the first step is to determinise the input automaton, hence the name determinisation-based approach.

Note that in this context *determinisation* refers to the conversion of a Büchi automaton to a *different type* of ω -automata, such as Muller, Rabin, Streett, or parity (see 2.1.2). As mentioned in Section 2.1.1, the determinisation of a Büchi automaton to a deterministic Büchi automaton is in general not possible, because deterministic Büchi automata are less expressive than non-deterministic Büchi automata.

This approach harnesses the fact that the complementation of deterministic automata is generally easier than the complementation of non-deterministic automata. This means that once the non-deterministic Büchi automaton is turned into a deterministic ω -automaton, it can be easily complemented, and the result can then be converted back to a non-deterministic Büchi automaton.

⁴Ramsey was a British mathematician who lived at the beginning of the 20th century

⁵Assuming a number of 10^{80} atoms in the observable universe, according to <http://www.wolframalpha.com/input/?i=number+of+atoms+in+the+universe>.

⁶This paper was preliminarily published in 1985 [46].

Safra (1988)

The first determinisation-based complementation construction has been described in 1988 by Safra [43, 44]. Safra’s main work is actually a construction for converting a non-deterministic Büchi automaton (NBW) to a deterministic Rabin automaton (DRW)⁷. This determinisation construction is what today commonly is known as *Safra’s construction*. However, Safra also describes a series further conversions that, when chained together, result in a complementation construction. The complete series of conversions is as follows:

1. NBW \longrightarrow DRW (Safra’s construction)
2. DRW \longrightarrow $\overline{\text{DSW}}$ (Complementation)
3. $\overline{\text{DSW}}$ \longrightarrow $\overline{\text{DRW}}$
4. $\overline{\text{DRW}}$ \longrightarrow NBW

That means, the DRW resulting from Safra’s determinisation construction is complemented to a DSW, which is converted to a DRW, which in turn is converted to an NBW. The resulting NBW is the complement of the input NBW. Note that a horizontal indicates that the automaton accepts the complement language of the automata without a bar.

Safra’s construction is in fact a modified subset construction that guarantees the equivalence of the input and output automata⁸. The main difference of Safra’s construction to the classical subset construction is that the states of the output automaton are labelled by trees of subsets of states, rather than just subsets of states. These trees are called *Safra trees*. Detailed analyses and explanations of Safra’s construction can be found, next to Safra’s papers [43, 44], in [42], [23], [2], or [28].

The complementation step from the DRW to the complement DSW can be trivially done by interpreting the Rabin acceptance condition as a Streett acceptance condition [23]. This works, because, as mentioned in Section 2.1.2, the two acceptance conditions are the duals of each other. The final conversions from DSW to DRW, and DRW to NBW are described by Safra in [43, 44] or alternatively in [23].

The entire construction from an NBW to a complement NBW has a worst-case state complexity of $2^{O(n \log n)}$. With this result, Safra’s complementation construction matched the lower bound of $n!$ that has been introduced in the same year by Michel [30]. However, as pointed out by Vardi [63], the big- O notation hides a large gap between the two functions, and Safra’s complementation construction is thus not as “optimal” as it seems. Thus, the quest for finding Büchi complementation constructions with an even lower worst-case state complexity continued.

Muller and Schupp (1995)

In 1995, Muller and Schupp proposed an improvement of Safra’s determinisation construction [35]. This construction can be used at the place of Safra’s original construction in the conversation chain that we described above for Safra’s construction. Muller and Schupp’s construction uses *Muller-Schupp trees* instead of Safra trees. While the principles of the two constructions are very similar, it is said that Muller and Schupp’s construction is simpler and more intuitive than Safra’s construction [42]. However, the drawback is that in many concrete cases Muller and Schupp’s construction produces larger output automata than Safra’s construction [2]. The worst-case state complexity of Muller and Schupp’s construction is, similarly to Safra’s construction, $2^{O(n \log n)}$. A detailed comparison of the determinisation constructions by Muller and Schupp, and Safra can be found in [2].

Piterman (2007)

A further improvement of Safra’s determinisation construction has been proposed by Piterman from EPF Lausanne in 2007 [38]⁹. The main difference to Safra’s construction is that Piterman’s construction

⁷This is the notation for ω -automata types that we introduced in Section 2.1.2.

⁸Safra shows in his papers [43, 44] nicely that applying the classical subset construction to an NBW A results in a DBW B that may accept some words that are not accepted by A . This means that the subset construction is not *sound* with respect to Büchi automata.

⁹A preliminary version of this paper has appeared in 2006 [37].

converts the input Büchi automaton to a deterministic parity automaton (DPW), rather than a deterministic Rabin automaton (DRW). Furthermore, Piterman uses a more compact version of Safra trees, what results by trend in smaller output automata. Also in terms of worst-case complexity, Piterman's construction is more efficient than Safra's construction. The blow-up in Piterman's construction from the NBW to the DPW is $2n^n n!$, whereas in Safra's construction the blow-up from the NBW to the DRW is $12^n n^{2n}$ [37, 38].

Since Piterman's determinisation construction produces a DPW, it entails different conversions for complementation than Safra's construction. In particular, the procedure is as follows [54]:

1. NBW \longrightarrow DPW (Piterman's construction)
2. DPW $\longrightarrow \overline{\text{DPW}}$ (Complementation)
3. $\overline{\text{DPW}}$ $\longrightarrow \overline{\text{NBW}}$

The complementation of the DPW can be trivially done by increasing the number assigned to each state by one [54] (see Section 2.1.2 for an explanation of the parity acceptance condition). The complemented DPW can then be converted directly to an NBW [54].

2.3.3 Rank-Based Approach

The rank-based approach is based on run analysis with run DAGs (see Section 2.2.1). The basic “mechanics” of rank-based construction is similar to the subset construction, that is, the output automaton is constructed state by state, by adding to every state a successor state for every symbol of the alphabet. In rank-based constructions, these states are derived from levels of a run DAG. The idea is that the run analysis with run DAGs is interweaved with the construction, so that the states of the output automaton represent levels of “virtual” run DAGs.

As we have seen in Section 2.2, the purpose of run analysis is to find out whether *all* runs of an input automaton A on a word α are non-accepting. Because in this case, the complement B must accept α , and in all other cases it must not accept it. In rank-based constructions this is achieved by the means of natural numbers called *ranks* that are assigned to the vertices of a run DAG (hence the name rank-based approach). These ranks are assigned in such a way that vertices corresponding to *accepting* states only get *even* ranks. Furthermore the ranks along paths of the run DAG are monotonically decreasing. The effect of this is that every path gets eventually trapped in a rank. If for a given path this trapped rank is odd, it means that the corresponding run does not visit any accepting states anymore starting from a certain point, and thus is non-accepting. If all the paths of a run DAG get trapped in an odd rank, then all the runs of the automaton on the given word are non-accepting. This situation is called *odd-ranking* and indicates that the complement must accept the word.

This basic idea is the same for all rank-based construction and described in many places, for example, [11, 25, 13, 63, 21, 45]. The individual rank-based constructions differ mainly in details of how the ranking is done.

Klarlund (1991)

Klarlund's construction [22] rather foreshadowed the rank-based constructions than being one itself. It was one of the first constructions that was neither Ramsey-based nor determinisation-based [22, 25]. Its fundamental idea is very similar to the later rank-based constructions, but it does not use the term “rank” and is not explicitly based on run DAGs. However, Klarlund uses so-called *progress measures*, which have a very similar role as the ranks of later constructions. This construction served as the base for Kupferman and Vardi's rank-based construction in 1997 [24]. The worst-case state growth of Klarlund's construction is $2^{O(n \log n)}$.

Kupferman and Vardi (1997/2001)

This construction has been published for the first time in 1997 [24] and again in 2001 [25]. Both publications are entitled “Weak Alternating Automata Are Not That Weak”. Kupferman and Vardi basically

pick up Klarlund’s construction and describe it in a different way [25]. Actually, they provide two different descriptions for the same construction, one based on weak alternating automata (WAA)¹⁰, and one rank-based.

The first description includes several conversions between ω -automata types. It starts by interpreting the input NBW as a universal co-Büchi automaton (UCW) which accepts the complement language of the initial NBW. This UCW is converted to a WAA, which is finally converted to an NBW. The second description equals the rank-based approach that we described in the introduction to this section. Kupferman and Vardi state that their two versions and Klarlund’s construction produce identical results. However, they claim that their descriptions make the construction easier to understand and implement. [25].

Thomas (1999)

This construction by Thomas [53] is based on the concept of alternating automata, like the previous construction by Kupferman and Vardi (1997) that proceeds via a WAA [24]. The difference of Thomas’ construction is that it performs the complementation step in the framework of alternating automata. In particular, the input NBW is converted to a weak alternating parity automaton (WAPA), which is complemented to another WAPA, which finally is converted back to an NBW. Note that Thomas uses the concept of ranks that are assigned to the states of an alternating automata in a similar way as we described in the introduction to this section. Thus, we categorise this construction under the rank-based approach, even if it does not proceed in the typical way.

Friedgut, Kupferman, and Vardi (2006)

In 2006, Friedgut, Kupferman, and Vardi [13]¹¹ proposed an improvement of the rank-based construction by Kupferman and Vardi from 2001 [25]. The improvement consists in a better ranking function called *tight ranking* that reduces the maximum number of states in the output automaton¹² The worst-case state complexity of the construction is $(0.96n)^n$, which is an enormous reduction from the 2001 construction by Kupferman and Vardi with a result of $(6n)^n$. The title of Friedgut, Kupferman, and Vardi’s paper is “Büchi Complementation Made Tighter”. It refers to the fact that this construction provides a new upper bound for the precise state complexity of Büchi complementation that considerably *tightens* the gap between the best known upper bound and the then known lower bound by Michel [30] of $n!$.

Schewe (2009)

The work by Schewe [45] is a further optimisation to the construction by Kupferman, Friedgut and Vardi [13]. It consists of the use of *turn-wise* tests in the *cut point construction* [45]. This optimisation allows to reduce the worst-case state complexity of the construction to $(0.76n)^n$ modulo a polynomial factor in $O(n^2)$ [45], that is, around $(0.76n)^n O(n^n)$. This comes very near to the lower bound of $(0.76n)^n$ that has previously been established by Yan [67, 68]. Schewe’s paper is entitled “Büchi Complementation Made Tight”, which refers to the fact that finally the gap between the lower and upper bounds for Büchi complementation has been made really “tight” (after having been made only “tighter” by Friedgut, Kupferman, and Vardi in 2006 [13]).

2.3.4 Slice-Based Approach

The slice-based approach can be seen as the brother of the rank-based approach. The main difference is that the slice-based approach is based on *reduced split trees* (see Section ??), rather than run DAGs. As for the rank-based approach, the “mechanics” of the slice-based approach is similar to the subset construction, that is, the output automaton is constructed state by state. In a slice-based construction these states are derived from levels of reduced split trees (as opposed to rank-based constructions where

¹⁰Weak alternating automata have been introduced by Muller, Saoudi, and Schupp in 1986 [34].

¹¹This paper has been preliminarily appeared in 2004 [12].

¹²Alternative descriptions of the tight ranking function can be found, for example, in [11], [45], or [63].

the states are derived from levels of run DAGs). The levels of reduced split trees are called *slices* [65], hence the name slice-based approach.

Slice-based constructions use decorations, also called colours [1], that they assign to the vertices of the corresponding reduced split trees (and thus to the components of the states that are derived from the slices of the tree). These decorations serve the same role as the ranks in rank-based constructions, namely the detection whether all runs of the input automaton on a specific word are non-accepting, in which case the complement must accept the word. A comparative analysis of the slice-based and the rank-based approach has been done by Fogarty et al. [11]¹³. Below we briefly present the two principal slice-based constructions.

Vardi and Wilke (2007)

With this construction, Vardi and Wilke have established the slice-based approach [54]. The construction is divided into an *initial phase* and a *repetition phase*. In the initial phase, slices (thus, states) are constructed without decorations. At some point, the construction guesses to transfer to the repetition phase in which the vertices of the slices (thus, components of states) are decorated by one of the decorations *inf*, *new*, or *die*. These decorations determine whether the enclosing state is accepting or non-accepting. A loose upper bound for the worst-case state complexity of this construction is $(3n)^n$ [65]. Vardi and Wilke describe their construction using left-to-right reduced split trees.

Kähler and Wilke (2008)

This construction by Kähler and Wilke [20] is a generalisation of the previous slice-based construction by Vardi and Wilke [65]. The construction is modified such that it can treat complementation and disambiguation¹⁴ of Büchi automata in a uniform way [54]. The generalisation however comes at the price of a higher worst-case state complexity of $4(3n)^n$ [54]. Kähler and Wilke also use left-to-right trees for their construction.

With this review of the four main Büchi complementation approaches, we conclude the present chapter. In the next chapter we will explain in detail the Fribourg construction, which is another slice-based construction, and which is the subject of our empirical performance investigation that we present later in the thesis.

¹³In this work, the authors even propose a construction that unifies the two approaches.

¹⁴Disambiguation means the conversion of an automaton to an equivalent *unambiguous* automaton. An automaton is unambiguous if every accepted word has only one accepting run [31]. Note that this differs from determinisation, as a word may still have multiple runs, but only one accepting run.

Chapter 3

The Fribourg Construction

Contents

3.1	The Construction	27
3.1.1	Basics	27
3.1.2	Upper-Part Construction	29
3.1.3	Lower-Part Construction	31
3.2	Optimisations	36
3.2.1	R2C: Omission of States Whose Rightmost Component is 2-Coloured	36
3.2.2	M1: Merging of Adjacent Components	37
3.2.3	M2: Single 2-Coloured Component	37
3.3	General Remarks	39

IN THIS CHAPTER we present the Fribourg construction. The Fribourg construction is the Büchi complementation construction whose empirical investigation of performance makes up the core of this thesis (Chapters 4 and 5). The Fribourg construction belongs to the *slice-based* complementation approach that we reviewed in Section 2.3.4. This means that it is based on reduced split trees, which we explained in Section 2.2.4. Of the two versions of reduced split trees, left-to-right and right-to-left, our description of the Fribourg construction in this chapter uses right-to-left reduced split trees.

The Fribourg construction is being developed at the University of Fribourg by Joel Allred and Ulrich Ultes-Nitsche¹. A detailed description of the construction has been published internally in 2014 as a technical report entitled “Complementing Büchi Automata with a Subset-tuple Construction” [1].

The aim of this chapter is to give an intuitive and practically oriented description of the construction. We try to make its concrete application as clear as possible. Our focus lies therefore not on a formal description of the construction. For the formal background, as well as proofs of correctness, we refer to the work by the authors of the construction [1].

This chapter is structured as follows. In Section ??, we describe the basic Fribourg construction. The construction is divided into two sub-constructions, and we explain each of them separately. In Section 3.2, we present three optimisations that have been proposed for the Fribourg construction. These optimisations will also be subject to our empirical performance investigation in the next chapter. Finally, in Section 3.3, we give some general remarks about the Fribourg construction, especially putting it into relation with other complementation constructions.

3.1 The Construction

In this section we explain the basic Fribourg construction without any applied optimisations. The construction consists of two sub-constructions that we call upper-part construction and lower-part construction. In the following we first describe features that are common to these two sub-construction (Section 3.1.1). Then, in Section 3.1.2 and 3.1.3, we describe in detail the upper-part construction and the lower-part construction, respectively. We illustrate the application of these two sub-construction with an example which we work through step by step.

3.1.1 Basics

Below we first give a high-level view of the construction, and then present the state generation procedure of the construction. The state generation procedure determines a successor state for an existing state, and lies at the heart of the construction. In its basic form, the state generation procedure is common to the two sub-constructions that we present in the next two sections.

High-Level View

As mentioned, the Fribourg construction consists of two sub-constructions, the *upper-part construction* and the *lower-part construction*. These two sub-constructions are chained together, such that the output of the upper-part construction becomes the input of the lower-part construction. Figure 3.1 illustrates this in a functional view of the Fribourg construction.

The automaton A to be complemented is the input for the upper-part construction. The upper-part construction builds up a new automaton by starting with an initial state and adding state by state. The output of the upper part construction is an intermediate automaton U . This intermediate automaton U will be the *upper part* of the final complement automaton. The lower-part construction takes U as input and adds further states and transitions to it until finally outputting an automaton B which is the complement of the input automaton A . We say that the lower-part construction attaches the *lower part* of the final complement automaton to the upper part. This terminology comes from the fact that it is

¹As mentioned, the authors call their construction *subset-tuple construction*, in this thesis we will however use the name *Fribourg construction*.




Figure 3.1: Functional view of the Fribourg construction with an automaton A as input and its complement B as output.

convenient to draw the states of the lower part below the states of the upper part [1]. This also explains why the two sub-constructions are called upper-part construction and lower-part construction.

State Generation

The two sub-constructions can be seen as modified subset constructions. Like the subset construction, they work on a set of existing states and add to each state one successor for each symbol of the alphabet. The principle of how successor states are generated is the same for both sub-construction. The difference between the two is that the lower-part construction adds additional information to the states (the colours), so that the states generated by the lower-part construction are different from the states generated by the upper-part construction. Below we explain this basic principle of how new states are generated.

A state of the Fribourg consists of a tuple (that is, an ordered sequence) of subsets of states of the input automaton (note that we will from now on refer to the states of the input automaton as the *input-states*, and to the states of the output automaton as *output-states*). This contrasts with the subset construction in which the output-states consist of a single subset of input-states. The states of the Fribourg construction are thus more structured than the states of the subset construction. We refer to the subsets of a tuple of a state as the *components* of this state.

The sequence of components of each state of the Fribourg construction is imposed by a slice (that is, a level) of a reduced split tree. We will explain the process of going from an existing state over the slices of a reduced split tree to a successor state further below. For now, we look at how a slice defines the structure of a state. Figure 3.2 shows a reduced split tree (on the left), and for each slice (framed by a shaded box), the state that this slice would define (on the right).




Figure 3.2: Relation between slices of a reduced split tree (left) and states of the output automaton of the Fribourg construction (right). The sequence of nodes in a slice determines the sequence of components in a state.

As can be seen in Figure 3.2, each node of a slice gives rise to a component in the corresponding state. To this end, it does not matter whether the nodes in the slices are siblings or whether their familial relations are. It is important to note, however, that the order of the nodes in a slice must be preserved in the order of the components in a state. It is similarly possible to interpret a state as a slice of a reduced split tree (indicated by the double-ended arrows in Figure 3.2). In this case, it is not possible to deduce the ancestors of the resulting nodes, however, this is not necessary for the construction.

The procedure of determining the successor of a given existing state on a certain alphabet symbol is as follows:

- Interpret the existing state as a slice of a reduced split tree
- Create the subsequent slice for the appropriate alphabet symbol

- Interpret the new slice as a state

The state that results from the new slice is the successor of the existing state for the given alphabet symbol. If, for example, we want to define the successor on the symbol a (a -successor) of the state $(\{q_0\})$, then as the first step we would create a slice representation of it that looks like the first slice from the top in Figure 3.2. Then, we create based on the input automaton the subsequent slice for symbol a , that might look like the second slice from the top in Figure 3.2. Interpreting this slice as a state results in $(\{q_0\}2, \{q_1\})$, which thus is the a -successor of $\{q_0\}$.

Note that the translation of states to and from slices is only a mental aid and not a necessary ingredient. It is possible to deduce $(\{q_0\}2, \{q_1\})$ from $\{q_0\}$, in our previous example, directly without the detour over slices of a reduced split tree. However, by conceptually using reduced split trees, we can cover the central points of the state generation process by re-using our knowledge of the previously known and independent concept of reduced split trees. This includes the splitting of accepting and non-accepting input-states, the placement of the accepting subset to the right of the non-accepting subset, and the retention of only the first occurrence of an input-state from right to left.

3.1.2 Upper-Part Construction

In this subsection, we recapitulate the upper-part construction, and demonstrate its application with an example. The example is based on the input automaton in Figure 3.3.

Description

The upper-part construction consists basically only of the state generation procedure that we explained the last section. It starts with an initial state with a single component containing the initial state of the input automaton. This will be the initial state of the final complement automaton. Then, the state generation procedure is applied for each alphabet symbol for each unprocessed state.

If a determined successor state is identical to a state that already exists in the automaton, then only a transition to this state is added. If a determined successor state is empty, that is, contains no components, then no new state or transition at all is added to the automaton. In this case, the state being processed has no successor for the corresponding alphabet symbol, and we say that it is incomplete.

The state generation process is repeated until all states have been processed. At this stage, a deterministic intermediate automaton without any accepting states has been produced. This intermediate automaton will be the upper part of the final complement automaton.

A peculiarity is that the upper part must be complete. An automaton is complete, if every state is complete, which means having an outgoing transition for each symbol of the alphabet. If the upper-part automaton is not complete, then it must be made complete by adding a sink state. A sink state is a state that serves as the “missing successor” of all the incomplete states. Note that this sink state that is potentially added to the upper-part automaton must be accepting.




Figure 3.3: Non-deterministic Büchi automaton used for the examples in this section.

Example

We demonstrate the application of the upper part construction by applying it to the example automaton A in Figure 3.3. Note that this automaton has an alphabet size of one, and thus can run only on a single

ω -word, namely a^ω . Furthermore, the automaton does not accept this word, which means that it does not accept *any* word and is thus empty. We selected this special automaton to keep our example handy and small. However, the construction works in exactly the same way for non-empty automata with larger alphabets.

The steps of the upper-part construction on the automaton A are shown in Figure 3.4 (a) to (d). Below, we walk through these steps one by one.





Figure 3.4: Steps of the upper-part construction applied to the example input automaton in Figure 3.3.


The first step in Figure 3.4 (a) is to start with the initial state. As can be seen, this state $(\{q_0\})$ contains only one component with the initial state of the input automaton A . This state will be the initial state of the final complement automaton B .

The next step Figure 3.4 (b) is to add the a -successor of the initial state $(\{q_0\})$. This is done by the state generation procedure that we described in the last section. In particular, this means: (1) interpret the state $(\{q_0\})$ as a slice of a reduced split tree, (2) based on the input automaton A , determine the subsequent slice for symbol a , and (3) interpret this new slice as a state. This new state is the a -successor of $(\{q_0\})$. We can illustrate these steps in the following way:




The state $(\{q_0\})$ results in a slice with only one node containing $\{q_0\}$. For creating the next slice, we determine the set of states of A that can be reached from $\{q_0\}$ on symbol a , which is $\{q_0, q_1, q_2\}$. According to the rules of reduced split trees, this set is split into a non-accepting $\{q_0, q_2\}$, and an accepting $\{q_1\}$. The non-accepting set $\{q_0, q_1, q_2\}$ becomes the left child, and the accepting set $\{q_1\}$ becomes the right child of the current node. Translating the resulting new slice to a state results in $(\{q_0, q_2\}, \{q_1\})$. This state is thus set as the a -successor of $(\{q_0\})$ in the automaton under construction.

The next step in Figure 3.4 (c) is to determine the a -successor the newly created state $(\{q_0, q_2\}, \{q_1\})$. Applying the same procedure as above, results in the following:



The state $(\{q_0, q_2\}, \{q_1\})$ corresponds to a slice with two nodes. For creating the next slice, these nodes must be processed from right to left. Regarding the node with $\{q_1\}$, the set of states in A that is reachable from $\{q_1\}$ on a is $\{q_2\}$, which becomes thus the only child of this node. Regarding the next node with $\{q_0, q_2\}$, the state set reachable from $\{q_0, q_2\}$ is $\{q_0, q_1, q_2\}$. Now, since q_2 already exists in the slice under construction, it is removed from the set. The resulting set $\{q_0, q_1\}$ is split into a non-accepting $\{q_0\}$ and an accepting $\{q_1\}$, which become the left child and right child of this node, respectively. The resulting slice corresponds to the state $(\{q_0\}, \{q_1\}, \{q_2\})$, which thus becomes the a -successor of $(\{q_0, q_2\}, \{q_1\})$.

The last step in Figure 3.4 (d) consists in determining the a -successor of $(\{q_0\}, \{q_1\}, \{q_2\})$. The state generation procedure in this case looks as follows:



The state $(\{q_0\}, \{q_1\}, \{q_2\})$ results in a slice with three nodes. Again, we have to process these nodes from right to left. The node with $\{q_2\}$ has a non-accepting $\{q_2\}$ as its only child. Regarding the node with $\{q_1\}$, the reachable set in A is $\{q_2\}$. However, since q_2 already exists in the slice, it is omitted, and since this empties the set, no child of this node is added to the new slice. Regarding the last node with $\{q_0\}$, the set of states reachable from $\{q_0\}$ is $\{q_0, q_1, q_2\}$, but again, due to the omission of q_2 and the splitting into an accepting and a non-accepting set, this results in the accepting right child $\{q_1\}$ and the non-accepting left child $\{q_0\}$. The state corresponding to this slice is $(\{q_0\}, \{q_1\}, \{q_2\})$, which is identical to the state we are currently processing. Consequently, we only add a transition from $(\{q_0\}, \{q_1\}, \{q_2\})$ back to itself.

At this stage, all the states in the intermediate automaton have been processed. Since the resulting automaton is complete, there is no need to add a sink state to it. This means that the upper-part construction is finished.

Note that in this example, we had to determine only one successor for each state, because there is only one symbol in the alphabet of our example automaton. For input automata with more than one alphabet symbols, the procedure that we did for each state must be repeated for each symbol of the alphabet. This increases the number of required steps, however, the basic principles generating states stay exactly the same.

3.1.3 Lower-Part Construction

Below, we describe in detail the lower-part construction. The lower-part construction can be seen as an extended version of the upper-part construction. The difference is that the lower-part construction additionally assigns colours to all the components. These colour assignments are the main subject of our explanations below. After explaining the lower-part construction, we demonstrate its application by an example. To this end, we will use and continue the same example that we used for the upper-part construction.

Description

The lower-part construction takes the intermediate automaton resulting from the upper-part construction as input. It works on the states of the upper part by regarding them as the unprocessed states to be processed. This results in additional successors for the states of the upper part, and finally in the entire lower part growing out of the upper part. Note that in the final complement automaton, there will be transitions from the upper part to the lower part, but not back from the lower part to the upper part. This means that once a run of the output automaton enters the lower part, it cannot go back to the upper part again.


The basic state generation procedure of the lower-part construction is similar as in the upper-part construction. The difference between the two sub-constructions is that the lower-part construction adds an additional piece of information to each component of new states. This piece of information consists of one of three colours, that we denote by 0, 1, and 2. We shed light on the meaning of these colours further below. For now, we explain how the colour of a component is determined.

The colour of component c of a new state q that is the successor (for a given alphabet symbol) of an existing state p , is determined by three pieces of information:

1. Whether the component c is accepting or non-accepting
2. The colour of the predecessor component of c
3. Whether the predecessor state p contains any 2-coloured components

Regarding the first point, a component is accepting if it contains accepting input-states. Regarding the second point, each component c in a state q has exactly one predecessor component c_{pred} in the predecessor state p . This relation emerges from the parent-child relation in the reduced split tree representation of the two states. That is, from this perspective c_{pred} is the parent of c . Consequently, every component has exactly one predecessor component, but every component may have zero, one, or two successor

components. This relation of components can be best illustrated by an example. Consider the following two state p and q and their corresponding slice representation:



The edges of the slice representation of the states (on the right) can be mentally merged into the states (indicated as dotted lines). This makes the relation between the components of the two states apparent. Regarding the components of q , we can see that the predecessor component of $\{q_0\}$ is $\{q_0\}$, of $\{q_1\}$ is $\{q_0\}$, and of $\{q_2\}$ is $\{q_2\}$. On the other hand, regarding the components of p , we can see that $\{q_0\}$ has two successor components ($\{q_0\}$ and $\{q_1\}$), $\{q_1\}$ has no successor components, and $\{q_2\}$ has one successor component ($\{q_2\}$). For the case of $\{q_1\}$ which has no successor components, we say that the input-runs leading through this components “disappear”.

Having explained the terms accepting component and predecessor component, we can now turn to the actual rules that determine the colour of a given component. One remark is however still necessary. The colouring rules need a way to recognise components that belong to the upper part of the output automaton. A way to guarantee this is to preliminarily assign the colour -1 to all the components of the upper-part states. In the following, we will rely on this convention. Figure 3.5 shows the complete rules for determining the colour of a component c .

Colour of c_{pred}	c is non-accepting	c is accepting
-1	0	2
0	0	2
1	2	2

(a) Case A: the predecessor state has *no* 2-coloured components.

Colour of c_{pred}	c is non-accepting	c is accepting
0	0	1
1	1	1
2	2	2

(b) Case B: the predecessor state *has* 2-coloured components.

Figure 3.5: Rules for determining the colour of a component c (in bold). The identifier c_{pred} denotes the predecessor component of c .

As can be seen in Figure 3.5 the colour rules are divided into two cases. Case A applies if the predecessor state does not contain any components with the colour 2, and case B applies if the predecessor state contains one or more components with the colour 2. In each of the cases, the colour of a component c is unambiguously determined by the two remaining factors, colour of the predecessor component c_{pred} , and by the fact whether c is an accepting or not. We provide the rationale behind these rules below, when we explain the meaning of the colours.

Note that in case A, the possible colours of the predecessor component includes -1 , but not 2. The missing of 2 is simply because by definition case A applies only if the predecessor state has no 2-coloured components, thus the colour of the predecessor component cannot be 2. However, in this case, the colour of the predecessor component may be -1 , namely in the case that the predecessor state is an upper-part state. In case B, on the other hand, colour -1 is missing. This is because case B applies only if the predecessor state contain 2-coloured components, in which case it cannot be an upper-part state.

Having explained the rules for determining the colours of components, we now give some insight into the meaning of these colours. The purpose of the colours is to signalise the presence or absence of certain runs of the input automaton (input-runs) entailed by a corresponding run of the output automaton (output-run). We divide these input-runs into *dangerous* and *safe* runs. Dangerous input-runs have visited at

least one accepting input-state since entering the lower part. These runs are dangerous in the sense that they might become accepting if they keep visiting accepting input-states. Safe runs, on the other hand, are runs that have not yet visited an accepting input-state since entering the lower part. Input-runs that never become dangerous correspond to non-accepting runs in the input automaton.

The purpose of the distinction of safe and dangerous input-runs is to recognise whether *all* input-runs corresponding to an output-run on a specific word are safe. In this case, the output-run must accept the word. From a practical point of view, the colours of the components of a state determine whether the state must be added to the accepting set (according to a rule that we present at the end of this section). In particular, the meaning of the three colours 0, 1, and 2 is as follows:

Colour 2 Signalises the presence of dangerous input-runs, that is, runs that have visited at least one accepting input-state. Colour 2 appears when a group of input-runs, that have previously been safe, visits an accepting component (see Figure 3.5 (a) lines 1 and 2). However, once a component has colour 2, it passes on this colour to all its successor components, no matter if they are accepting or non-accepting (see Figure 3.5 (b) line 3). This means that a path of 2-coloured components can only be cut if at some point the input-runs going through this component disappear (that is, the components has no successors). In this case, we say for brevity that the 2-coloured components disappears.

Colour 1 Means basically the same as colour 2, namely the presence of dangerous runs. The idea behind colour 1 is a caveat that can arise if we would assign colour 2 to *every* component whose corresponding input-runs just visited for the first time an accepting state. In this case, it could happen that in a sequence of output-states, there are constantly newly appearing 2-coloured components, and at the same time constantly disappearing 2-coloured components. The disappearance of the 2-coloured components means that the dangerous runs disappear, which in fact makes the corresponding output-states safe. However, this disappearance is hidden by the constant appearance of new 2-coloured components. The reason for colour 1 is to prevent this from happening. If the predecessor state already contains 2-coloured components (case B in Figure 3.5), then all the components that according to case A in Figure 3.5 would deserve colour 2, get colour 1 instead (see Figure 3.5 (b) line 1 and 2). We say that these are dangerous runs that are *on hold*. However, as soon as all 2-coloured components disappear, and consequently case A applies again, then all the successors of the 1-coloured component get colour 2 (see Figure 3.5 (a) line 3). In this way, the dangerous components that were previously on hold, become the new active dangerous components.

Colour 0 Signalises that all input-runs are safe. This means that these input-runs have not yet visited an accepting input-state since entering the lower part. If these runs stay safe indefinitely, then they are non-accepting in the input automaton.

As mentioned, the colours of the components of a state determines if a state must be made accepting. The corresponding rule is as follows:

Each state of the lower part that contains no 2-coloured components is an accepting state.

That is, states that contain only 0-coloured and 1-coloured components must be made accepting. The rationale behind this rule is that if an output-run visits such a state infinitely often, then all the input-runs on the corresponding word are non-accepting. For states with only 0-coloured components, this is clear, as colour 0 means that there are no dangerous runs. For states with 1-coloured components, the case is slightly more complicated, because 1-coloured components actually signalise the presence of dangerous runs. However, note that the situation that a state contains 1-coloured components, but no 2-coloured components (case A in Figure 3.5), can only occur if previously all 2-coloured components disappeared. If this happens an infinite number of times (which is implied by the output-run visiting this state infinitely often), then an infinite number of 2-coloured components disappear. This means that all the corresponding dangerous input-runs die after reading a finite number of symbols, which means that none of them can be accepting. Consequently, a state with 1-coloured components also guarantees that all the corresponding input-runs are non-accepting.

All these rules and mechanisms become clearer by seeing them in action. Below we demonstrate the application of the lower-part construction by an example.

Example

We continue the complementation of the example automaton A in Figure ?? that we left off after the upper-part construction in the last section. In this section, we step by step add the lower part to the upper part, and in the end we obtain the complement automaton of A .

Figure 3.6 shows a selection of the steps of the lower-part construction applied to the upper part depicted in Figure 3.4 (d). Note that we use the following notation to indicate the colour of a component $\{q_i, \dots\}$:

- $\widehat{\{q_i, \dots\}}$: colour -1
- $\{q_i, \dots\}$: colour 0
- $\overline{\{q_i, \dots\}}$: colour 1
- $\overline{\overline{\{q_i, \dots\}}}$: colour 2





Figure 3.6: Selected steps of lower-part construction applied to the upper part depicted in Figure 3.4 (d).

The first step of the lower-part construction in Figure 3.6 (a) is to start with the upper part (as depicted in Figure 3.4 (d)), and assign colour -1 to all the components. As mentioned, we use this convention with colour -1 throughout this thesis.

The following steps of the lower-part construction consist in fact of the state generation function (as described in Section 3.1.1) plus the assignment of colours as described above. In Figure 3.6 (b) we add the a -successors for all the upper-part states. The state generation procedure results in the successor states $(\{q_0, q_2\}, \{q_1\})$, $(\{q_0\}, \{q_1\}, \{q_2\})$, and $(\{q_0\}, \{q_1\}, \{q_2\})$. To determine the colours of the components of these states, we consult the colour rules in Figure 3.5. The predecessor states of all the new states do not contain a 2-coloured component, thus case A in Figure 3.5 (a) applies. Furthermore, since the predecessor states are upper-part states, all the predecessor components have colour -1 . That means that for each component only the rule in line 1 of Figure 3.5 applies. According to this rule, the non-accepting components get colour 0 and the accepting components get colour 2 . This results in the successor states $(\{q_0, q_2\}, \overline{\{q_1\}})$ for $(\widehat{\{q_0\}})$, $(\{q_0\}, \overline{\{q_1\}}, \{q_2\})$ for $(\widehat{\{q_0, q_2\}}, \widehat{\{q_1\}})$, and $(\{q_0\}, \overline{\{q_1\}}, \{q_2\})$ for $(\widehat{\{q_0\}}, \widehat{\{q_1\}}, \widehat{\{q_2\}})$. Note that this makes the upper-part states non-deterministic. In particular, this is the way how the runs of the final output automaton can switch non-deterministically to the lower part.


In Figure 3.6 (c), we create the a -successor of $(\{q_0, q_2\}, \overline{\{q_1\}})$. Again, the new state is first determined by

the state generation procedure, what results in $(\{q_0\}, \{q_1\}, \{q_2\})$. Now, we have to determine the colour for each of the components $\{q_0\}$, $\{q_1\}$, and $\{q_2\}$. In this case, the predecessor state $(\{q_0, q_2\}, \overline{\{q_1\}})$ contains a 2-coloured component, thus case B of the colour rules in Figure 3.5 applies. Next, we need to know the colour of the predecessor of each component. To figure this out, we refer to the slice representation of the two states that we used for the state generation procedure. As explained earlier, we can mentally merge the edges between the nodes of the slices into the two states. This results in the following picture, where the component relations are indicated by dotted arrows:




Consequently, the predecessor component of $\{q_0\}$ and $\{q_1\}$ is $\{q_0, q_2\}$, and the predecessor component of $\{q_2\}$ is $\overline{\{q_1\}}$. Regarding $\{q_0\}$ and $\{q_1\}$, the colour of their predecessor component is 0, thus the rule on line 1 of Figure 3.5 (b) applies. Since $\{q_0\}$ is non-accepting, it gets colour 0, and since $\{q_1\}$ is accepting, it gets colour 1. Note how the assignment of colour 1 instead of colour 2 to $\{q_1\}$ sets it *on hold* because there exists already a 2-coloured component in the predecessor state. Regarding $\{q_2\}$, the colour of its predecessor component is 2, and thus, according to the rule on line 3 of Figure 3.5 (b) it gets colour 2. This results in the state $(\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}})$.

Still in Figure 3.6 (c) we create the a -successor of this new state $(\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}})$. By the state generation procedure, we obtain $(\{q_0\}, \{q_1\}, \{q_2\})$. Since the predecessor state $(\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}})$ has a 2-coloured component, case B of Figure 3.5 applies. Regarding the predecessors of the components in $(\{q_0\}, \{q_1\}, \{q_2\})$ we observe the following relations:




By applying the appropriate colour rules depending on the colour of the predecessor component and the acceptance or non-acceptance of the component itself, $\{q_0\}$ gets colour 0, $\{q_1\}$ gets colour 1, and $\{q_2\}$ gets colour 2. The resulting state is $(\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}})$, which is identical to its predecessor state. Consequently, we just add an a -loop to $(\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}})$.

In Figure 3.6 (d), we add the successor of the state $(\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}})$. After the state generation procedure that generates $(\{q_0\}, \{q_1\}, \{q_2\})$, again case B of the colour rules in Figure 3.5 applies. The predecessor component relations are as follows:



By applying the colour rules, $\{q_0\}$ gets colour 0, $\{q_1\}$ gets colour 1, and $\{q_2\}$ gets colour 0. This results in the state $(\{q_0\}, \overline{\{q_1\}}, \{q_2\})$. It is interesting to see that between these two states, the 2-coloured component $\overline{\{q_1\}}$ disappears, because it has no successors. This disappearance is reflected in the fact that $(\{q_0\}, \overline{\{q_1\}}, \{q_2\})$ contains a 1-coloured component but no 2-coloured components, which can only be the case if a 2-coloured component disappears (as we explained above).

Still in Figure 3.6 (d), we create the a -successor of this new state $(\{q_0\}, \overline{\{q_1\}}, \{q_2\})$. After the state generation procedure, which results in $(\{q_0\}, \{q_1\}, \{q_2\})$, we have to apply case A of the colour rules in Figure 3.5, because $(\{q_0\}, \overline{\{q_1\}}, \{q_2\})$ does not contain a 2-coloured component. The component relations between the two states are as follows:



By applying the appropriate colour rules, $\{q_0\}$ gets colour 0, $\{q_1\}$ gets colour 2, and $\{q_2\}$ gets colour 0,

which results in the state $(\{q_0\}, \overline{\{q_1\}}, \{q_2\})$. This state already exists in the automaton, and thus we only add an a -transition to it.

At this point, all the states in the automaton have been processed. The last thing that remains to do is to define the accepting set. The rule is that every state of the lower part that does not contain any 2-coloured components must be made accepting. In our case, this applies only to the state $(\{q_0\}, \overline{\{q_1\}}, \{q_2\})$, and thus, this is the only accepting state of the automaton.

The resulting automaton, as depicted in Figure 3.6 (d), is the final complement automaton B . It is easy to verify that this automaton is indeed the complement of the input automaton A in Figure 3.3. Both automata have an alphabet size of one, and thus work only on the word $\Sigma^\omega = \{\alpha^\omega\}$. Whereas A does not accept this word (A is empty), B accepts it (B is universal), and is thus the complement of A .

A loose upper bound for the worst-case state complexity of the entire construction (including the construction of both, the upper and the lower part) has been calculated to be in $O((1.59n)^n)$ [1]. This result is however subject ongoing research, and may be sharpened to smaller number in future work.

This concludes our explanation of the basic Fribourg construction. In the next section, we present three optimisations that can be applied to the Fribourg construction in order to increase its performance.

3.2 Optimisations

The authors of the Fribourg construction propose three optimisations for the construction [1]. In this section, we briefly describe each of them. Further details about all the optimisations can be found in [1].

For easy reference, we define the abbreviations R2C, M1, and M2 for these three optimisations. The optimisations can be added to the construction like add-ons, and they can be combined in different ways. This results in different *versions* of the Fribourg construction. Throughout this thesis, we specify a specific version of a construction by appending the included optimisations after the construction name. For example, Fribourg+R2C means the Fribourg construction including the R2C optimisation, and Fribourg+M1+R2C means the Fribourg construction including the M1 and the R2C optimisation.

3.2.1 R2C: Omission of States Whose Rightmost Component is 2-Coloured

The R2C optimisation can be summarised as follows:

If the input automaton is complete, then states of the lower part whose rightmost component has colour 2 can be omitted.

It is important to highlight that this optimisation can only be applied if the input automaton is complete. We hint at this restriction by the letter C in R2C. If the completeness of the input automaton is given, then the optimisation allows to remove all the states whose rightmost component is 2-coloured.

The reason that it is allowed to remove these states is as follows. If the input automaton is complete, then every input-state has at least one successor state on every alphabet symbol. This means that in the corresponding reduced split trees, that are used for the Fribourg construction, the rightmost nodes are guaranteed to have a successor node. Note that this guarantee applies only to the rightmost nodes, because the successors of the other nodes might be omitted in the right-to-left processing of the slices. Consequently, in terms of states, all rightmost components are guaranteed to have a successor component.

If the colour of the rightmost component is 2, then all successor components inherits this colour (see the rule on line 3 of Figure 3.5 (b)). Since rightmost colours *always* have a successor, *all* the successor states with a rightmost 2-coloured component will contain a 2-coloured component. And since states containing a 2-coloured component are non-accepting, these states form a non-accepting cycle in the automaton. Consequently, it is safe to remove this cycle without changing the language of the automaton.

In our example from the last section, we could thus remove the states $(\{q_0, q_2\}, \overline{\{q_1\}})$ and $(\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}})$ from the output automaton in Figure 3.6 (d). However, remember that this is only allowed if the input

automaton is complete, which in our case is true for our input automaton in Figure 3.3. The R2C optimisation is described in [1] (Section 4.F).

3.2.2 M1: Merging of Adjacent Components

The M1 optimisation allows the merging certain combinations of adjacent components depending on their colour. By merging, we mean the creation of the union of two adjacent sets. Components can be merged in three cases in the following way:

1. Two adjacent 1-coloured components can be merged to a single 1-coloured component
2. Two adjacent 2-coloured components can be merged to a single 2-coloured component
3. Two adjacent 2-coloured and 1-coloured components, where the 1-coloured component is to the right of the 2-coloured component, can be merged to a single 2-coloured component

These mergers can be done recursively. This means that the result of a merger is in turn subject to further mergers. For example, given the state $(\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}}, \overline{\{q_3\}}, \overline{\{q_4\}}, \overline{\{q_5\}})$, the following chain of mergers is possible:

$$\begin{aligned}
 (\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}}, \overline{\{q_3\}}, \overline{\{q_4\}}, \overline{\{q_5\}}) &\longrightarrow (\{q_0\}, \overline{\{q_1, q_2\}}, \overline{\{q_3, q_4\}}, \overline{\{q_5\}}) && \text{By rules 1 and 2} \\
 &\longrightarrow (\{q_0\}, \overline{\{q_1, q_2, q_3, q_4\}}, \overline{\{q_5\}}) && \text{By rule 3} \\
 &\longrightarrow (\{q_0\}, \overline{\{q_1, q_2, q_3, q_4, q_5\}}) && \text{By rule 2}
 \end{aligned}$$

This means that with the M1 optimisation, we add the state $(\{q_0\}, \overline{\{q_1, q_2, q_3, q_4, q_5\}})$ to the automaton rather than $(\{q_0\}, \overline{\{q_1\}}, \overline{\{q_2\}}, \overline{\{q_3\}}, \overline{\{q_4\}}, \overline{\{q_5\}})$. For the formal description and proof of correctness of the M1 optimisation, we refer to [1] (Section 4.E).

The application of these mergers obviously reduces the maximum number of states that the construction can produce. The worst-case state growth regarding the number of states of the lower part has been calculated in [1] to be $O((1.195n)^n)$. This must be added to a worst-case state growth of the upper part of $(0.53n)^n$. However, the resulting term is clearly dominated by $O((1.195n)^n)$. It will be a subject of our empirical performance investigation in Chapter 4 to investigate the impact of the M1 optimisation on *acutal* automata.

3.2.3 M2: Single 2-Coloured Component

The M2 optimisation is the most complicated of the three optimisations. Its effect on the output automaton can be summarised as follows:

States of the lower part contain at most one 2-coloured component.

Note that without the M2 optimisation, it is possible that states contain two or more 2-coloured components. This can happen if the predecessor state contains no 2-coloured components and the new state contains multiple accepting components (see colour rules on line 1 and 2 of Figure 3.5 (a)). With the M2 optimisation, in this case only one of the accepting components gets colour 2, and the others get colour 1. The purpose of the M2 optimisation is to further reduce the maximum number of states that the construction can produce.

It is important to note that the M2 optimisation is dependent on the M1 optimisation. That means that the M2 optimisation can only be applied if the M1 optimisation is also applied. We selected the abbreviation M2 for this optimisation to highlight that it is based on M1.

Below we first explain how a single 2-coloured component of a state is created. Next, we explain in detail the steps that the disappearance of this single 2-coloured component entails.

Creation of a 2-Coloured Component

As mentioned, when creating a new state, only one of the components, that without the M2 optimisation would become 2-coloured, gets colour 2. This component is defined to be the rightmost of the candidates for colour 2. A possible procedure to ensure this restriction is to determine the colours for the components from right to left, and assign colour 2 only if it has not been previously assigned in the state. If it has been previously assigned, then assign colour 1 instead of 2.

A special case applies if the component that gets assigned colour 2 has a 2-coloured predecessor component and a sibling to the left of it. With sibling we mean an adjacent component that has the same predecessor component. If this two criteria are fulfilled, then the sibling gets colour 2 as well. Note that without this special treatment, it would get colour 1 instead of 2. Thus, there are two 2-coloured components in the state. However, since the M2 optimisation also requires the M1 optimisation to be applied, these two 2-coloured components are merged to a single 2-coloured components. In this way, statement that all states have at most one 2-coloured components is still true.

Actions on the Disappearance of a 2-Coloured Component

The restriction of having only one 2-coloured component in a state entails special actions when this 2-coloured components disappears (that is, has no successor components). Imagine a state p with a single 2-coloured component, and its successor state q . If after the assignment of colours as described above, q has no 2-coloured component, then we have the situation that the 2-coloured component of p disappeared. The required action in this situation is to change the colour of one of the 1-coloured components of q (if there are any) to 2. If q does not contain any 1-coloured components, then nothing is done and the creation of the state is finished. If q contains exactly one 1-coloured component, then the colour of this component is changed to 2. If q contains multiple 1-coloured components, then the following procedure is applied.

- Go to the position in q where the successor of the disappeared 2-coloured component of p *would* be
- From this position, go to the left until arriving at a 0-coloured component
 - If the search arrives at the leftmost component, continue with the rightmost component
- From this 0-coloured component, go to the left until arriving at the first 1-coloured component
 - If the search arrives at the leftmost component, continue with the rightmost component
- Change the colour of this 1-coloured component to 2

That is, if there are multiple 1-coloured components in the state, then we select the first 1-coloured component that is to the left of the first 0-coloured component that is to the left of the position where the successor of the disappeared component would be, and change its colour to 2. In the course of this, we understand the relation “to the left of” as cyclically, such that the component to the left of the leftmost component is the rightmost component.

The effect of changing the colour of a 1-coloured component of a state q to 2 is that q , which previously contained no 2-coloured components, now contains a 2-coloured component. However, since a 2-coloured component disappeared, q must still be accepting, even though it contains a 2-coloured component. For this reason, we marks states such as q with a star (*). In general, each state that has been subject to a change of a 1-coloured component to a 2-coloured component is marked with a star. Finally, the rule for adding states to the accepting set is changed such that it includes all the states of the lower part that contain no 2-coloured component *and* all the states that contain a 2-coloured component but are marked with a star.

Figure 3.7 shows the result of applying the Fribourg construction with the M1 and M2 optimisations to the example automaton in Figure 3.3. The difference of this automaton to the complement resulting from the Fribourg construction without the M1 and M2 optimisation (see Figure 3.6 (d)) is the state at the bottom right of the automaton. As can be seen in the figures, without the optimisation the component $\{q_1\}$ of this state is 1-coloured. With the M1 and M2 optimisations, however, $\{q_1\}$ is 2-coloured. In both cases, the 2-coloured component of the predecessor state disappeared, however, with the M1 and M2 optimisations, the resulting 1-coloured $\{q_1\}$ is made 2-coloured. Additionally, the state is marked with a star, so that it is still part of the accepting set. Regarding the successors of this marked state, the

same procedure is repeated indefinitely: the 2-coloured component disappears, and one of the 1-coloured components is made 2-coloured. This is the reason why this state has a loop.

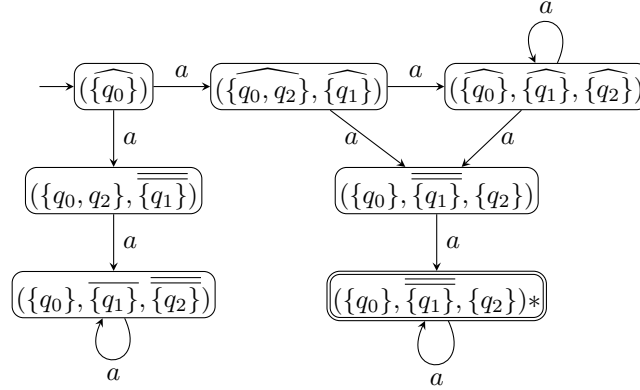


Figure 3.7: Result of applying the Fribourg construction with the M1 and M2 optimisation to example automaton in Figure 3.3.

The M2 optimisation further reduces the maximum number of states that the construction can produce. A loose upper bound of $O((0.86n)^n)$ for the worst-case state growth of the lower part of the automaton has been calculated in [1] (Section 4.H). However, calculations of actual values suggest that the real bound for the lower part might be as low as $O((0.76n)^n)$. The investigation of the worst-case state complexity of the M1 and M2 optimisation is subject to ongoing research by the authors of the Fribourg construction. Note that $O((0.86n)^n)$ or $O((0.76n)^n)$ is the maximum number of states of the lower part, and must be added to the maximum number of $(0.53n)^n$ upper-part states. However, the resulting term is dominated by the function for the lower part, that is, $O((0.86n)^n)$ or $O((0.76n)^n)$.

3.3 General Remarks

The Fribourg construction has many similarities to the slice-based construction by Vardi and Wilke[65] (see Section 2.3.4) that has been published in 2007. The upper part and lower part of the Fribourg construction correspond to the initial phase and repetition phase of Vardi and Wilke's construction. Furthermore, the colours 0, 1, and 2 of the Fribourg construction correspond to the decorations *inf*, *new*, and *die* of Vardi and Wilke's construction. However, the Fribourg construction has been developed independently, and is not based on Vardi and Wilke's construction.

Rather, the Fribourg construction is based on Kurshan's complementation construction for *deterministic* Büchi automata [26]. This construction also uses an upper part and a lower part which are both deterministic, except for non-deterministic transitions from the states of the upper part to states of the lower part. The Fribourg construction has emerged as an adaptation of Kurshan's construction for non-deterministic Büchi automata, with the slice-based run analysis approach in mind.

With these remarks, we conclude the present chapter about the Fribourg construction. At this point, we have provided all the background that is needed for our empirical performance investigation of the Fribourg construction, which we present in the next chapter.

Chapter 4

Empirical Performance Investigation of the Fribourg Construction

Contents

4.1	Implementation	41
4.1.1	The GOAL Tool	41
4.1.2	Implementation of the Construction	43
4.1.3	Verification of the Implementation	45
4.2	Test Data	46
4.2.1	GOAL Test Set	46
4.2.2	Michel Test Set	49
4.3	Experimental Setup	50
4.3.1	Internal Tests	51
4.3.2	External Tests	52
4.3.3	Time and Memory Limits	53
4.3.4	Execution Environment	54

IN THIS CHAPTER we come to the core of this thesis, namely the empirical performance investigation of the Fribourg construction. The aim of this investigation is to find out how the Fribourg construction performs on *actual* automata. The goal of this chapter is to describe the setup of our empirical investigation. The results of the study are presented in Chapter 5.

We approach the investigation of the Fribourg construction from two point of views. First, we want to test the construction with different combinations of its optimisations. Second, we want to compare the Fribourg construction with other complementation constructions. We refer to the first track of investigation as the *internal tests*, and to the second track as the *external tests*. Note that our main measure of performance is the number of states of the produced complement automata.

For an empirical performance investigation of a Büchi complementation construction, we need basically three things: an implementation of the construction, test data, and an experimental setup, including an execution environment. Regarding the *implementation*, we decided to create it as a part of the GOAL tool. GOAL is an ω -automata tool, which has already been used previously for empirical investigations of Büchi complementation constructions (see Section 1.2.2). Furthermore GOAL contains many pre-implemented complementation constructions, so that, for the external tests, we can run all the tested complementation constructions in a common framework.

Test data means basically a set of concrete automata that are given as input to a construction, in order to analyse the output of the construction. We use two different test sets for our investigation, that we refer to as the GOAL test set and the Michel test set (consisting of Michel automata [30]). The automata of both test sets have been used in previous investigations of Büchi complementation construction by other authors (see Section 1.2.2). We believe that the use of test data that has been previously used increases the validity of our results, as they can be compared to related work.

Our *experimental setup* includes a definition of the concrete tests that are run, that is, which construction, or version of a construction, is executed on which test data. It also states time and memory limits that are imposed on the experiments, and specifies the execution environment. Regarding the execution environment, since our study includes in large computations, we decided to use professional high-performance computing (HPC) facilities. Concretely, we execute all the experiments on a Linux-based HPC cluster named UBELIX of the University of Bern¹.

In this chapter, we discuss each of the mentioned points in detail. We present the implementation of the Fribourg construction in Section 4.1, the test data in Section 4.2, and the experimental setup in Section 4.3.

4.1 Implementation

As mentioned, we implemented the Fribourg construction as part of the GOAL tool. In this section, we first present the GOAL tool in a general way (Section 4.1.1). In Section 4.1.2, we describe how we implemented the Fribourg construction with its optimisations as a part of this tool. Clearly, an implementation is only useful if it is correct. Therefore, in Section 4.1.3, we describe how we verified the correctness of our implementation.

4.1.1 The GOAL Tool

GOAL stands for *Graphical Tool for Omega-Automata and Logics* and is being developed at the Department of Information Management at the National Taiwan University (NTU)². The tool has been described in various scientific publications [56, 57, 58, 55]. The GOAL tool is freely available on <http://goal.im.ntu.edu.tw>. It is a Java program, and thus runs on every platform having a Java Runtime Environment (JRE).

GOAL is a graphical and interactive tool for creating and manipulating ω -automata. It provides a large number of operations that can be applied to these automata. These operations range from input test-

¹<http://ubelix.unibe.ch>

²<http://exp.management.ntu.edu.tw/en/IM>

ing, conversions to other types of ω -automata, to union and intersection. Figure 4.1 shows a screenshot of GOAL’s graphical user interface with an open menu showing the breadth of operations that GOAL provides. Of course, complementation is also part of these operations, and GOAL includes implementations of many existing Büchi complementation constructions. We will discuss the complementation constructions that GOAL provides further below.

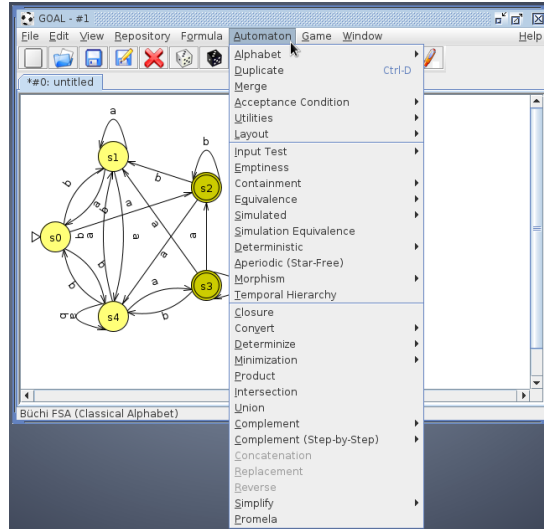


Figure 4.1: The graphical user interface of GOAL (version 2014–11–17). The open menu item gives an idea about the different types of manipulations that can be applied to ω -automata.

The versions of GOAL have names of the form YYYY–MM–DD that correspond to the release date of the version. The latest version at the time of this writing, and the version that we used in our experiments, is 2014–11–17.

Automata can be imported to and exported from GOAL in different formats. The default format is the GOAL File Format (GFF). Files of the GFF type have by convention the extension `.gff`.

As mentioned, GOAL has a graphical user interface, as shown in Figure 4.1. However, almost the entire functionality of GOAL is also available through a command line interface. For example, complementing an automaton can then be done with the command `gc complement -m safra -o out.gff in.gff` from the commandline. The effect of this command is to complement the automaton in the file `in.gff` with Safra’s complementation construction, and write the complement automaton to the file `out.gff`. The command `gc` is the entry point to GOAL’s command line interface. For our empirical performance investigation, we use the command line interface of GOAL.

Let us now have a closer look at the Büchi complementation constructions that are pre-implemented in GOAL. The 2014–11–17 version of GOAL provides a number of 10 complementation constructions. We list these constructions along with their authors and reference to the literature in Table 4.1.

As can be seen in Table ??, almost all complementation constructions that GOAL provides have been reviewed in Section 2.3. The only exception is ModifiedSafra, which is a minor modification to Safra’s construction that has been proposed by Althoff et al. [2]. Furthermore, GOAL has at least one construction of each of the four main complementation approaches. Ramsey belongs to the Ramsey-based approach, Safra, ModifiedSafra, Piterman, and MS belong to the determinisation-based approach, Rank, WAA, and WAPA to the rank-based approach, and Slice and Slice+P to the slice-based approach. Note that the identifiers that are defined in Table 4.1 for each construction are also used by GOAL, and we will use them throughout this thesis to refer to the corresponding constructions.

The construction Slice+P is actually the construction Slice with the P option. That is, from the perspective of GOAL, there is only the single construction Slice, however, if the P option is specified, then the slice-based construction by Vardi and Wilke [65] is used, and otherwise the slice-based construction by Kähler and Wilke [20] is used.

At this point, it is worth pointing at a related project of the same research group, called the *Büchi*

Identifier	Description	Authors (year)	Reference
Ramsey	Ramsey-based construction	Sistla et al. (1987)	[46, 47]
Safra	Safra’s construction	Safra (1988)	[43, 44]
ModifiedSafra	Slightly modified Safra’s construction	Althoff et al. (2006)	[2]
Piterman	Safra-Piterman construction	Piterman (2007)	[37, 38]
MS	Muller-Schupp construction	Muller, Schupp (1995)	[35]
Rank	Rank-based construction	Schewe (2009)	[45]
WAA	Via weak alternating automata	Kupferman, Vardi (1997/2001)	[24, 25]
WAPA	Via weak alternating parity automata	Thomas (1999)	[53]
Slice+P	Slice-based construction	Vardi, Wilke (2007)	[65]
Slice	Enhanced slice-based construction	Kähler, Wilke (2008)	[20]

Table 4.1: The Büchi complementation constructions provided by GOAL version 2014–11–17.

Store [59]. This is an online repository of categorised ω -automata that can be downloaded in different formats (including GFF). The Büchi Store is accessible over the web on <http://buchi.im.ntu.edu.tw/>. Furthermore, there is a binding in GOAL to the Büchi Store, so that automata can be directly downloaded into GOAL. For our study we did not use the Büchi Store, but it might be an interesting option for future work.

4.1.2 Implementation of the Construction

In this section, we describe our implementation of the Fribourg construction as a part of GOAL. In particular, our implementation has the form of a *plugin* for GOAL. This is possible thanks to the plugin interface that GOAL provides. Below, we first briefly explain this plugin architecture of GOAL, and then present our plugin containing the implementation of the Fribourg construction.

The GOAL Plugin Architecture

GOAL has been designed from the ground up to be modular and extensible. To this end, it has been created with the Java Plugin Framework (JPF)³. This framework allows to build applications whose functionality can be easily and seamlessly extended by writing plugins for it. These plugins can be installed in the main application without the need to recompile the whole application. Rather, the plugin is compiled separately and the resulting bytecode files are copied to the directory tree of the main application. It is not even necessary to know the source code of the main application in order to write a plugin. The interfaces of JPF itself, and the documentations of the relevant classes of the main application are all that plugin writers need to know.

In some more detail, JPF requires an application to define so called *extension points*. For any extension point, multiple *extensions* can be provided. These extensions contain the actual functionality of the application. A JPF application basically consists of extensions that are plugged into their corresponding extension points. A plugin is an arbitrary bundle of extensions and extension points. It is the basic unit of organisation in the Java Plugin Framework.

One of the extension points of GOAL is called **ComplementConstruction**. The extensions to this extension point contain the actual complementation constructions that GOAL provides. For adding a new complementation construction to GOAL, one has thus to create a new extension to **ComplementConstruction**. This extension can then be wrapped in a plugin, and the plugin can be compiled and installed in the main application, what makes it an integral part of it. This means that once the plugin is installed, the new construction is included in GOAL in the same way as all the other constructions.

It is thanks to this open architecture of GOAL that welcomes extension that we were able to seamlessly add the Fribourg construction to this tool. In particular, we created a GOAL-plugin that contains the

³<http://jpf.sourceforge.net/>

Option	Description
R2C	Apply R2C optimisation
M1	Apply M1 optimisation
M2	Apply M2 optimisation
C	Make input automaton complete
R	Remove unreachable and dead states from output automaton
RR	Remove unreachable and dead states from input automaton
MACC	Maximise accepting states of input automaton
B	Use the “bracket notation” for state labels

Table 4.2: The options for the Fribourg construction in our Fribourg construction plugin for GOAL.

implementation of the Fribourg construction. We present this plugin below.

The Fribourg Construction Plugin

The name of our plugin is `ch.unifr.goal.complement`⁴. The plugin is freely available on https://frico.s3.amazonaws.com/goal_plugins/ch.unifr.goal.complement.zip and can be installed very easily by any GOAL user. Note that we provide instructions for installing the plugin in Appendix A.

We aimed at fully integrating the Fribourg construction in GOAL. The current version of the plugin includes menu integration of the Fribourg construction, the pertinent saving of option defaults, and a step-by-step execution facility for the Fribourg construction. This last point is a great way for learning how the Fribourg construction works.

Our implementation of the Fribourg construction also includes the three optimisations R2C, M1, and M2 that we described in Section 3.2. These optimisations are freely selectable by the user and can be combined in different ways. In the GUI, the optimisations are presented to the user as selectable check boxes before the start of the construction. In the command line mode, they can be set as flags of the entered command.

In addition to these three optimisations, we added further options to our construction. Table 4.2 provides a listing of all of them. Note that each option has an identifier consisting of upper-case letters. We use these identifiers throughout the rest of this thesis to refer to the corresponding options.

We discuss each of the options in Table 4.2 below. The first three options are the three optimisations to the Fribourg construction that we described in Section 3.2. The R2C optimisation is implemented so that it applies only to input automata that are complete. That is, if R2C is activated for the complementation of an automaton that is not complete, then this option has no effect. The options M1 and M2 stand for the M1 and M2 optimisations. Since M2 is dependent on M1, it is not possible to select M2 without also selecting M1. This restriction is enforced in both the GUI and the command line interface.

The C option is one of the options that modifies the input automaton before the actual complementation starts. This option first checks if the input automaton is complete, and if this is not the case, makes it complete by adding a sink state. The purpose of the C option is to be used in conjunction with the R2C optimisation. Making an automaton complete before the start of the construction ensures that the R2C optimisation will be applied. The question arises whether, in terms of performance, it is worth to do this, because making an automaton complete increases its size by one, what increases the potential size of the output automaton. This question has been investigated in previous work on the Fribourg construction by Göttel [14]. In this thesis will re-investigate this point in an extended form (see Section 4.3).

The R option modifies the output automaton at the end of the construction. In particular, it removes all the so called unreachable and dead states from the complement. Unreachable states are states that cannot be reached from the initial state. Dead states are states from which it is not possible to reach an accepting state. These states can be removed from any automaton without changing the language

⁴By convention, JPF plugins are named after the corresponding Java package.

of the automaton. The pre-implemented complementation constructions Ramsey, Piterman, Rank, and Slice contain a similar R option.

An example usage case of the R option is to investigate the number of unreachable and dead state that a complementation construction produces. A way to do this is to complement the same automaton with and without the R option, and calculating the difference of the sizes of the two complements. Such investigations have been done with by Tsai et al. [54] in their comparative study of several Büchi complementation constructions in GOAL. We will do similar investigations, although it is not the main focus of our study.

The RR option is similar to the R option, with the difference that it removes the unreachable and dead states from the input automaton rather than from the output automaton. This option is not common among the other complementation constructions in GOAL, and we are not using it in our study.

The MACC option is again an option that modifies the input automaton before the start of the construction. It maximises the accepting set of the input automaton. This means that as many states as possible are made accepting so that the language of the automaton is not changed. The idea is that the larger number of accepting state simplifies the complementation task. This technique has been introduced and empirically investigated by Tsai et al. [54]. The pre-implemented complementation constructions Ramsey, Piterman, Rank, and Slice also contain the MACC option. In our study we will however not use the MACC option.

Finally, the B option only affects the formatting of the labels of the output-states. It has no effect on the construction itself. In particular, it activates the bracket notation, which is an informal alternative way to indicate the colours of components of states. It does so by using different types of brackets for the subsets, rather than wrapping the subset in a tuple with the colour number, like the default notation.

4.1.3 Verification of the Implementation

It is important to verify that our implementation of the Fribourg construction is correct. With correct we mean that the implementation produces correct complements. Note that this is not the same sense of correctness as the verification that our implementation represents in all details the specification of the construction as devised by its authors. However, it was our aim to do this, and thanks to the close collaboration with the authors of the construction, we are confident that this is the case.

In order to verify that our implementation produces correct complements, we performed so called complementation-equivalence tests in GOAL. This works as follows. Take an automaton A and complement it with one of the existing constructions in GOAL, resulting in the complement automaton B' . Then, complement the same automaton A with our implementation of the Fribourg construction, which yields the automaton B'' . Now, by the means of GOAL's equivalence operation test whether A' and A'' are equivalent. If yes, then our implementation produced a correct complement.

This approach makes a couple of assumptions. First, it relies on the correctness of the pre-implemented complementation constructions, and on the equivalence operation of GOAL (which is also based on complementation). Second, with this empirical approach it is not possible to conclusively verify the correctness of our implementation. Every passed complementation-equivalence test just further confirms the hypothesis that our implementation is correct, but it can never be proved. Despite these issues, a sufficient number of passed tests, still makes us sufficiently confident that our implementation is correct.

We tested the Fribourg construction in different versions with all the options that we described in the last section (except the B option as it does not influence the construction). The tested versions are the following.

- Fribourg
- Fribourg+R2C+C
- Fribourg+M1
- Fribourg+M1+M2
- Fribourg+R
- Fribourg+RR

- Fribourg+MACC

We tested each version with 1,000 randomly generated automata. The automata had a size of 4 and an alphabet size between 2 and 4. The pre-implemented that we used as the ground-truth is Piterman. The result was that all the tests were successful.

With four states, the automata used for the tests are rather small, and it would be interesting to do the tests with larger automata. However, our limited computing and time resources prevented us from doing so. The equivalence test of GOAL is implemented as reciprocal containment tests, which in turn is based on complementation. This means that the equivalence test includes the complementation of the complements of the test automata. By using larger test automata, their complements might be already so large that a further complementation is practically infeasible with our available resources. Nevertheless, we think that the number of passed tests with these small automata indicates the absence of implementation errors with a high probability.

4.2 Test Data

After an implementation, the test data is the second main ingredient of an empirical study. The test data consists of the sample automata that are given as input to the construction, in order to analyse the produced output. The test data is an important part of every empirical study and should meet several requirements: it should include enough test cases so that the results are significant, it should not be biased in favour of the tested algorithm, and it should cover test scenarios that are relevant for evaluating the performance of the tested algorithm. The ideal case is to use publicly available test data, that has been previously used for other studies. In this way, the test data can be critically inspected by everybody, and furthermore the results of the study are comparable to the results of other studies.

We decided to use two sets of test data that have both been used in other studies. The first is the GOAL test set, that has been created for and used in the study by Tsai et al. [54], as well as in other studies, such as [7] and [14]. It consists of 11,000 non-deterministic Büchi automata with a size of 15. The second set of automata, which we call Michel test set, consists of the smallest four Michel automata. Michel automata are the automata that Michel used to prove the lower bound of $n!$ for the worst-case state complexity of Büchi complementation [30]. The Michel automata have been used as the test data for other empirical studies, such as [2], or [14].

The two test sets that we use for our study are accessible via the following links, that are maintained by the author of this thesis:

- https://frico.s3.amazonaws.com/test_sets/goal.zip⁵
- https://frico.s3.amazonaws.com/test_sets/michel.zip

Note that the automata in the test sets are in the GOAL File Format. In the following, we first describe the GOAL test set in Section ??, and then the Michel test set in Section ??.

4.2.1 GOAL Test Set

The GOAL test set is the larger and more complex one of the two test sets. In this section, we first introduce the GOAL test set and describe its basic structure. In the second part, we present the results of an analysis that we did in order to reveal further properties of the GOAL test set, namely the number and distribution of complete, universal, and empty automata.

Structure

The GOAL test set has been created by Tsai et al. [54] for an empirical study evaluation the effects of several optimisations on existing Büchi complementation constructions. This study has been executed in

⁵The original link to the GOAL test set, maintained by the authors of GOAL, is http://goal.im.ntu.edu.tw/wiki/lib/exe/fetch.php?media=goal:c1aa2010_automata.tar.gz.

GOAL and the automata in the test set are available in the GOAL file format. This is why we call the data GOAL test set.

The GOAL test set consists of 11,000 automata of size 15^6 . All the automata in the test set have 15 states and an alphabet size of 2 (consisting of the symbols 0 and 1). The further properties of the automata, namely accepting states and transitions, are determined by the values of two parameters called *transition density* and *acceptance density*.

The transition density determines the number of transitions of an automaton. In particular, the transition density t is defined as follows. Let n be the number of states of automaton A , and t its transition density. Then A contains tn transitions for each symbol of the alphabet. In the case that tn is not an integer, it is rounded up to the next integer. That is, if one of our automata with 15 states and the alphabet 0, 1 has a transition density of 2, then it contains exactly 30 transitions for symbol 0 and 30 transitions for symbol 1.

The acceptance density a is defined as the ratio of accepting states to non-accepting states in the automaton. It is thus a number between 0 and 1. If automaton A has n states and an acceptance density of a , then it has an accepting states. In the case that an is not an integer, it is rounded up to the next integer. For example, an automaton of the test set with an acceptance density of 0.5 contains 8 accepting states.

The GOAL test set is structured into 110 classes with different transition density/acceptance density pairs. These 110 pairs result from the Cartesian product of 11 transition densities and 10 acceptance densities. The set of transition densities \mathcal{T} and acceptance densities \mathcal{A} are the following:

$$\begin{aligned}\mathcal{T} &= (1.0, 1.2, 1.4, 1.6, 1.8, 2.0, 2.2, 2.4, 2.6, 2.8, 3.0) \\ \mathcal{A} &= (0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0)\end{aligned}$$

Thus, there is one class whose automata have a transition density of 1.0 and an acceptance density of 0.1, another class with a transition density of 1.0 and an acceptance density of 0.2, and so on. Each of the 110 classes contains 100 automata. According to [54], these parameter values were chosen to provide a broad range of complementation problems ranging from easy to hard.

Completeness, Universality, and Emptiness

We tested all the automata in the GOAL test set for completeness, universality, and emptiness. It is interesting to know about these special cases of automata because they may have an effect on the complementation results. The ideal complement of an empty or universal automaton has a size of one. Thus, by comparing this against the number of actually produced states, the number of “superfluous” states that the construction generated becomes apparent. Regarding completeness, the R2C optimisation applies only to complete automata, thus it is interesting to see on which automata this optimisation has in fact an effect.


GOAL provides a command for testing emptiness. However, it does not provide commands for testing completeness and universality. Therefore, we implemented these commands on our own and bundled them as a separate GOAL plugin. This plugin is called `ch.unifr.goal.util` and is also available as described in Appendix A.

With these GOAL commands we tested and recorded for each of the 11,000 automata whether it is complete, univversal, or empty. We then determined the number of complete, universal, and empty automata in the entire test set, as well as in each of the 110 transition density/acceptance density classes. The number of complete, universal, and empty automata in the entire test set of 11,000 automata are as follows:


- 990 complete automata (9%)
- 6,796 universal automata (61.8%)
- 63 empty automata (0.6%)

⁶There exists a second version with 11,000 automata of size 20.

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	0	0	0	0	0	0	0	0	0	0
1.2	0	0	0	0	0	0	0	0	0	0
1.4	0	0	0	0	0	0	0	0	0	0
1.6	0	0	0	0	0	0	0	0	0	0
1.8	1	1	0	0	0	1	1	1	0	0
2.0	0	5	1	2	2	3	1	2	2	2
2.2	5	10	8	5	3	5	8	6	7	1
2.4	10	6	11	11	8	6	10	20	9	7
2.6	17	17	12	16	14	19	22	21	19	19
2.8	27	20	29	32	26	27	30	25	24	19
3.0	37	37	40	39	34	37	38	35	38	39

(a) Number of *complete* automata per class.

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	4	5	5	7	8	4	6	10	4	3
1.2	1	3	5	8	8	12	10	13	4	14
1.4	2	17	13	17	20	24	22	21	27	26
1.6	16	28	30	37	49	42	42	49	45	45
1.8	31	40	55	59	64	67	76	70	63	78
2.0	60	64	85	75	83	83	79	90	87	83
2.2	67	87	86	88	89	91	89	89	89	86
2.4	88	89	86	92	95	95	94	97	96	97
2.6	86	93	92	97	97	97	98	96	98	96
2.8	94	97	95	94	97	99	98	97	97	100
3.0	99	99	99	97	99	98	100	100	100	99

(b) Number of *universal* automata per class

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	17	7	4	5	2	4	3	1	1	0
1.2	4	2	1	1	0	1	0	0	0	0
1.4	2	1	0	0	0	0	0	0	1	2
1.6	0	0	0	0	0	0	1	0	0	0
1.8	1	0	0	0	1	0	0	0	0	0
2.0	0	0	0	0	0	0	0	0	0	0
2.2	0	0	0	0	0	0	0	0	0	0
2.4	0	0	0	0	0	0	0	0	0	0
2.6	0	0	0	0	0	0	0	0	0	0
2.8	0	0	0	0	0	0	0	0	0	0
3.0	0	0	0	0	0	0	0	0	0	0


(c) Number of *empty* automata per class

Figure 4.2: Number of complete, universal, and empty automata in each of the 110 transition density/acceptance density classes of the GOAL test set. Note that each class contains 100 automata.

The fact that only 9% of the automata are complete means that the R2C optimisation of the Fribourg construction affects only 9% of the test data. This is an interesting fact for the analysis of the effect the R2C optimisation has on the overall performance of the construction. A surprisingly high number of 61.8% of the automata are universal. A reason for this might be the small alphabet consisting of only

two symbols of the automata. With a certain number of transitions in the automata, there seems to be a high probability that the automata are universal. Conversely, the number of empty automata is very low. This can be seen as the reverse of the same effect that causes the number of universal automata to be high.

In Figure 4.2, we show the number of complete, universal, and empty automata for each of the 110 classes. By the way, in this figure we introduce two ways for representing per-class data that we will use throughout this thesis. On the left side of Figure 4.2 the per-class data is represented as matrices. These matrices always have 11 rows and 10 columns, and the rows always represent the transition densities and the columns represent the acceptance densities. On the right side of Figure 4.2, the same data is visualised as so called perspective plots. The corner of the perspective plots that is closest to the viewer corresponds to the upper-left corner of the corresponding matrices. Thus, looking at a perspective plots is like looking at the corresponding matrix from beyond the upper-left corner. The advantage of matrices is that they show all the data values explicitly. The advantage of perspective plots is that they show the patterns in the data more intuitively. When we present the results of our study in Chapter 5, we mainly use perspective plots, however, we provide all the corresponding matrices in Appendix B.

Regarding the complete automata per class in Figure 4.2 (a), we can see that their number increases with the transition density. Up to a transition density of 1.6 there are no complete automata at all, and then it starts to increase up to a number between 34 and 40 for the transition density of 3.0. Since each class contains exactly 100 automata, these numbers can similarly be viewed as percentages. That the number of complete automata increases with the transition density is because a higher number of transitions per alphabet symbol in the automaton increases the probability that each state has at least one outgoing transition for each alphabet symbol. For example, with a transition density of 1.0 and 15 states, the automaton contains exactly 15 transitions for each alphabet symbol. It is still possible that this automaton is complete, but the probability is very low, because there must be a one-to-one mapping of transitions and states. On the other hand, with a transition density of 3.0, there would be 45 transitions per alphabet symbol, and the probability that each state gets one of them is much higher.

The number of universal automata per class in Figure 4.2 (b) also increases with the transition density, although much stronger. While in the classes with a transition density of 1.0, there are between 3 and 10 universal automata, in the classes with a transition density of 3.0 there are between 97 and 100. As already mentioned, the small alphabet size of the GOAL test set automata and a sufficiently high number of transitions results in a high probability that an automaton accepts every possible word, and thus is universal. In Figure 4.2 (b) we can also see that low acceptance densities result by trend in slightly fewer universal automata. This is because with fewer accepting states there is less chance that a given word is accepted. As we identified the small alphabet size as a possible reason for the high number of universal automata, it would be interesting to investigate the number of universal automata in automata with a bigger alphabet size. We let this as an idea for future work.

Conversely to the high number of universal automata, the number of empty automata is very low. The totally 63 empty automata are mainly concentrated in the upper-left corner of the matrix in Figure 4.2 (c). That is, the automata with a low transition density and a low acceptance density have the highest probability to not accept any word, and thus being empty. The reasons for this is roughly the opposite reasons for the distribution of the universal automata.

4.2.2 Michel Test Set

The Michel test set is very different from the GOAL test set. It consists of only four automata, that however exhibit an exceptionally high state growth. Michel automata have been introduced in 1988 by Michel [30] in order to prove a lower bound for the worst-case state complexity of Büchi complementation of $(n - 2)!$ [52]. The Michel automata are characterised by the parameter m . They have an alphabet size of $m + 1$, and $m + 2$ states. Michel proved that the complements of these automata cannot have less than $m!$ states. Since the number of states of the input automata is $n = m + 2$, the minimum state growth of these automata is $(n - 2)!$.

In practice, however, the state growth that is observed for Michel automata is much higher. It is so high that for practical reasons we are restricted to include only the first four Michel automata with

$m = \{1, \dots, 4\}$ in our test set. For the Michel automata with $m \geq 5$, the required time and computing power for complementing them with our implementation would by far exceed our available resources. We present extrapolations toward larger Michel automata in Section 5.1.2. Despite the small number of tested automata, we still obtain very interesting results, as we will see in Chapter 5.

The four Michel automata in our test set are shown in Figure 4.3. We will call them Michel 1, Michel 2, Michel 3, and Michel 4, respectively. As mentioned, Michel automata have $m + 2$ states and an alphabet size of $m + 1$. Furthermore, they all have a single accepting state. Our Michel automata 1 to 4 have thus 3, 4, 5, and 6 states, and alphabet sizes of 2, 3, 4, and 5, respectively.




Figure 4.3: The Michel automata with $m = \{1, \dots, 4\}$, an alphabet size of $m + 1$, and $m + 2$ states.

It is interesting to use Michel automata to investigate the performance of Büchi complementation constructions, because they force the constructions to generate large number of states. By furthermore using several Michel automata with different sizes, it is possible to fit a state growth function on the resulting complement sizes. The state growth observed from the Michel automata is most likely not equal to the worst-case state growth (there are even worse automata), but it still provides a lower bound for the worst-case state growth. This can be interesting in case the worst-case state growth has not yet been precisely determined. This function will certainly not be equal to the theoretical state growth function, but it still provides a lower bound for the worst-case state growth.

4.3 Experimental Setup

In this section, we describe the experimental setup of our empirical performance investigation. This includes the specification of the concrete constructions (or versions of constructions) that are tested, constraints that are imposed on the runs, and the execution environment. As mentioned, the investigation is divided into the internal tests and external tests. In the internal tests we compare different versions of the Fribourg construction with each other. In the external tests, we compare one of the versions of the Fribourg construction with a set of different complementation constructions.

In Section 4.3.1, we present the tested versions for the internal tests. These versions differ for the GOAL and the Michel test set, and we present them separately. In Section ??, we present the tested constructions (including one of the versions of the Fribourg construction) for the external tests. In Section 4.3.4, we describe the computing environment in which the experiments were executed. Finally, in Section 4.3.3, we present the time and memory limits that were imposed on the individual executions.

4.3.1 Internal Tests

The versions of the Fribourg construction used for the internal tests consist of combinations of the three optimisations R2C, M1, and M2, and of the additional options C and R (see list of options for the Fribourg construction in Table 4.2). The sets of versions are different for the GOAL and the Michel test set. Below we first present versions that we use with the GOAL test set, and then the versions used with the Michel test set.

GOAL Test Set

For the internal tests on the GOAL test set, we use the following eight versions of the Fribourg construction:

1. Fribourg
2. Fribourg+R2C
3. Fribourg+R2C+C
4. Fribourg+M1
5. Fribourg+M1+R2C
6. Fribourg+M1+R2C+C
7. Fribourg+M1+M2
8. Fribourg+R

The first version, Fribourg, is the plain Fribourg construction without any optimisations or options. The next two versions, Fribourg+R2C and Fribourg+R2C+C, aim at investigating the R2C optimisation. In Fribourg+R2C, the R2C optimisation is applied only to complete input automata. As we have seen in Section 4.2.1 these are just 9% of the automata. Fribourg+R2C+C, on the other hand, makes the incomplete input automata complete so that the R2C optimisation is applied to all automata. The fact that we want to find out is whether it is worth to increase the size of an automaton by one (for adding the sink state) for the sake of being able to apply the R2C optimisation.

A very similar question has been investigated in previous work on the Fribourg construction by Göttel [14]. In terms of our above listing, he compares the performance of Fribourg with Fribourg+R2C+C. As the test data, he also uses the GOAL test set. His results are that the overall mean complement size resulting from Fribourg+R2C+C is higher than for Fribourg. However, by looking closely at Göttel's results we suppose that the median complement size (which is not reported in [14]) might be lower for Fribourg+R2C+C than for Fribourg. This would be an interesting relation, and therefore, we decided to re-investigate this question.

Fribourg+M1 and Fribourg+M1+M2 aim at investigating the M1 and M2 optimisations. As M2 can only be applied together with M1, there are only these two possible combinations. As we will see in Chapter 5, Fribourg+M1 shows a better performance on the GOAL test set than Fribourg+M1+M2. Therefore, we do not further investigate any versions based on Fribourg+M1+M2, however, we further investigate versions based on Fribourg+M1. This is why the next two versions are Fribourg+M1+R2C and Fribourg+M1+R2C+C. With these versions, we investigate the effect of the R2C optimisation in combination with the M1 optimisation. The last version, Fribourg+R, corresponds to the first version, Fribourg, but with the difference that in Fribourg+R all unreachable and dead states are removed from the output automaton. This allows to determine the number of unreachable and dead states that have been produced by Fribourg.

The versions Fribourg+M1+R2C and Fribourg+M1+R2C+C enhance the “better” of Fribourg+M1 and Fribourg+M1+M2 with R2C and R2C+C, respectively. As we will see in Chapter 5, the better one of these two versions terms of median complement sizes is Fribourg+M1. That is, the application of M2 results in a decline, rather than a gain, in performance compared to the application of M1 alone. It is important to highlight that such results are specific to the used test data, and not universally valid. With different test data Fribourg+M1+M2 might be better than Fribourg+M1. As we will see in the next section, this is for example the case for our Michel test.

The last version, Fribourg+R, corresponds to the application of the plain Fribourg construction without any optimisations or options, but with the difference that at the end of the construction all unreachable and dead states are removed from the complement automaton. Comparing the results of Fribourg+R with the results of Fribourg reveals how many unreachable and dead states the Fribourg construction produces. This idea is inspired by the study by Tsai et al. [54] in which the number of unreachable and dead states is one of the main performance metrics for the tested constructions..

Michel Test Set

For the internal tests on the Michel test set, we use the following six versions of the Fribourg construction:

1. Fribourg
2. Fribourg+R2C
3. Fribourg+M1
4. Fribourg+M1+M2
5. Fribourg+M1+M2+R2C
6. Fribourg+R

The rationale behind the selection of these versions is basically the same as for the GOAL test set. However, there are the following differences. First, the Michel automata are complete, thus there is no need to include the C option. Second, for the Michel test set, Fribourg+M1+M2 is more performant than Fribourg+M1. This is why Fribourg+M1+M2+R2C is based on Fribourg+M1+M2 rather than Fribourg+M1.

4.3.2 External Tests

In the external tests, we compare the results of the best version of the Fribourg construction on each test set, with the results of a set of other complementation constructions that are implemented in GOAL. As already mentioned, the best versions of the Fribourg constructions on the two test sets turn out to be the following:

- Fribourg+M1+R2C for the GOAL test set
- Fribourg+M1+M2+R2C for the Michel test set

Regarding which other complementation constructions to use, we principally could use all the constructions listed in Table 4.1. However, in preliminary tests we observed that most of these constructions are not performant enough to be reasonably used with the GOAL test set. Using them would require to allocate very high time and memory resources. Tsai et al. [54] made a similar experience with these constructions in GOAL when they observed that Ramsey could not complete any of the 11,000 complementation tasks in the GOAL test set within the time limit of 10 minutes and a memory limit of 1 GB.

There are three constructions that are performant enough to be used with the GOAL test set for our purposes, namely Piterman, Slice, and Rank. For this reason, we only include these three constructions in the external tests⁷. Interestingly, these constructions represent three of the four main complementation approaches, namely the determinisation-based (Piterman), the rank-based (Rank), and the slice-based (Slice) approach. In detail we use the following constructions for the external tests:

- Piterman+EQ+RO
- Slice+P+RO+MADJ+EG
- Rank+TR+RO

Note that regarding the options, we include those options for each construction which are set as default in their implementation, except the MACC (maximise accepting set of input automaton) and R (remove

⁷Actually, Safra would also be performant enough, but since we can use Piterman, which is an improvement of Safra, we chose to not include Safra.

unreachable and dead states of output automaton) options, because they are not part of the constructions, but rather pre-processing and post-processing steps⁸. Some of the used options are described by Tsai et al. in [54], in which case we indicate it below.

Piterman+EQ+RO is the determinisation-based construction by Piterman [37, 38] with the EQ and RO options of its implementation in GOAL. These options are described in GOAL as follows:

- EQ: merge equivalent states during the conversion of the complement parity automaton to the complement Büchi automaton (described in [54])
- RO: reduce transitions in the conversion from the complement parity automaton to the complement Büchi automaton

Slice+P+RO+MADJ+EG is the construction by Vardi and Wilke [65] with the RO, MADJ, and EG options (note that the P option causes the usage the construction by Vardi and Wilke [65], as opposed to the construction by Kähler and Wilke [20], which is used without the P option). The options have the following meaning:

- RO: reduce the out-degree of states
- MADJ: merge adjacent 0-components or *-components (described in [54])
- EG: use enhanced guessing (described in [54])

Rank+TR+RO is the rank-based construction by Schewe [45] with the TR and RO options. The meaning of these options is as follows:

- TR: use tight rankings
- RO: reduce the out-degree of states

Note that in Chapter 5, we refer to these three construction as Piterman, Slice, and Rank, for short. However, what we actually mean are the three precise versions as described above.

4.3.3 Time and Memory Limits

Similarly to Tsai et al. [54], we set a time and a memory limit for each complementation task of the GOAL test set. The time limit is 600 seconds (10 minutes) CPU time, and the memory limit is 1 GB for the Java heap memory. These limits roughly coincide with the limits in [54], which are also 10 minutes and 1 GB, however, it is not specified whether the 10 minutes are CPU time or real time, and whether 1 GB applies to the process' entire memory, or to a portion of it, like the Java heap. Complementation tasks that exceed the time or memory limit are aborted, and consequently do not deliver any results.

Note that these limits apply only to the GOAL test set. For the Michel test set, we do not set any limits, because in this case there are only four automata, and we want each of them to be completed.

The reasons for these limits are restricted computing and time resources. Clearly, the ideal case would be to let every complementation task run to completion, no matter how long it takes and how much memory it uses. However, because of the high complementation complexity that Büchi automata may have, some few extreme cases may cause our available resources to be exceeded⁹. The study is for example ultimately limited by the physically available memory on the nodes, and the maximum running time of a job permitted by the cluster management software. With these limits we can thus cut off such extreme cases and keep the required resources for the study in affordable bounds.

We implemented the time limit by the means of the `ulimit` Bash builtin, which allows to set a maximum running times for processes. Processes that reach this time limit are aborted by the operating system.

The memory limit, as mentioned, defines the maximum size of the Java heap. The heap is the main memory area of the Java process. It is where all the objects that are created by the Java program are stored. Concretely, this means that the states of the complements that are generated by the tested constructions are stored on the heap. It is possible to set the maximum Java heap size with the `Xmx`

⁸For Piterman, we also exclude the SIM option, which is set by default. This is because it simplifies the intermediate complement parity automaton, which can be seen as a post-processing step.

⁹As we will see in Chapter 5, the distribution of the complexity of the tested Büchi automata is extremely right-skewed, that is, most are easy, and very few are hard.

option to the Java Virtual Machine¹⁰. We also set the initial size of the Java heap to 1 GB by the means of the `Xms` option, so that the heap does not need to be enlarged what could distort the measured execution times.

The presence of time and memory limits, and thus aborted complementation tasks, require the use of so-called *effective samples*. The concept of effective samples is also used by Tsai et al. . These limits are based on the limits that have been used by Tsai et al. [54]. The effective samples are those automata which have been successfully completed by *all* constructions that are compared to each other. Imagine, for example, two constructions C_1 and C_2 that are used to complement a test set of 1,000 automata. C_1 successfully complements all the automata, whereas C_2 gets aborted for 100 of these automata. If we would statistically evaluate and compare the 1,000 results of C_1 and the 990 results of C_2 , then C_2 would be advantaged, because the results of the 100 apparently hardest automata are not taken into account, whereas they are included in the evaluation of C_1 . This is why in this case only the 990 effective samples that have been successfully completed by both constructions must be evaluated. In Chapter 5, we will therefore only evaluate the effective samples of the compared constructions.

4.3.4 Execution Environment

Due to the computational intensity, we executed all the complementation tasks, that we outlined in the last section, on a HPC computer cluster. In particular, we used the computer cluster named UBELIX of the University of Bern¹¹. Note that we also executed the correctness tests for the implementation of the Fribourg construction (see Section 4.1.3) and the analysis of the GOAL test set (see Section ??) on UBELIX.

Generally, a computer cluster consists of interlinked *nodes* (computers) on which the *jobs* (computation tasks) of the users of the clusters are run. Typically, cluster users submit their jobs to a cluster management software, which automatically dispatches these jobs to suitable and available nodes. In the case of UBELIX, this cluster management software is Oracle Grid Engine (formerly known as Sun Grid Engine) version 6.2¹².

We ensured that all our tasks are executed on similar nodes. These nodes have the following specifications:

- Processor: Intel Xeon E5-2665 2.40GHz
- Architecture: 64 bit
- CPUs (cores): 16
- Memory (RAM): 64 GB or 256 GB
- Operating System: Red Hat Enterprise Linux 6.6
- Java platform: OpenJDK Java 6u34
- Shell: GNU Bash 4.1.2

Note that the memory of the nodes is either 64 GB or 256 GB. The tasks on the GOAL test set were run on nodes with 64 GB RAM, the tasks on the Michel test set on special high-memory nodes with 256 GB RAM. Apart from that, the specifications of these nodes are identical. The use of the high-memory nodes for the Michel test set was required, because the maximally allocatable memory per CPU of the nodes with 64 GB memory is 4 GB, and this was not enough to complement the Michel automata. With the high-memory nodes, on the other hand, a total of 16 GB can be allocated per CPU which was sufficient to complement all the Michel automata.

Regarding multicore usage, the behaviour of our tasks depends on GOAL, and thus ultimately on Java. GOAL is multi-threaded and thus our tasks may use multiple CPUs. Theoretically, they can use up to the total number of 16 CPUs of a node. However, we observed that our tasks typically used between two and four CPUs¹³.

¹⁰Example usage: `java -Xmx1G`

¹¹<http://ubelix.unibe.ch>

¹²<http://www.oracle.com/us/products/tools/oracle-grid-engine-075549.html>

¹³We can just indirectly guess this number by comparing the measured CPU times and real times of concrete tasks.

We also measured the execution time of each complementation task as CPU time and real time (also known as wallclock time). The CPU time is the time a process is actually executed by the CPU. The real time is the time that passes from the start of a process until its termination, and thus includes the time the process is not executed by the CPU (idle time). These measurements were done with Bash's `time` reserved word. If a process runs on multiple CPUs, the CPU time is counted on each CPU separately and finally summed up. This means that for multicore execution (as in our case), the CPU time may be higher than the real time. For single-core execution, on the other hand, this is not possible, and the CPU time may only be equal to or lower than the real time. In the analysis of our results in Chapter 5, we sometimes present statistics of the execution times. These times are always *CPU times*.

The complementation tasks are executed sequentially via the command line interface of GOAL. For each complementation task the GOAL application is started separately, which includes the loading of the Java Virtual Machine (JVM). The JVM startup time is thus included in our measured execution times, and acts like a constant. According to our observations, the JVM startup time is approximately two CPU time seconds.

A computer cluster is a multi-user environment and a node can be used by multiple users at the same time. Thus, the total load of a node may vary, depending on number and intensity of other users' jobs. Our tasks were also subject to varying node loads. We do not know whether this has an influence on our experiments, in particular, on the measurement of the execution times, and on the imposed time limit (see next section). The load of a node has however no influence on the complements that are produced by the complementation constructions, which is after all our main interest. Hence, we leave the topic of the influence of node load on execution time for future work.¹⁴

This concludes the present chapter that described the setup of our empirical performance investigation of the Fribourg construction. We explained how we implemented the Fribourg construction as a part of the GOAL tool. Then, we presented our test data, consisting of the GOAL test set and the Michel test set. Finally, we presented our experimental setup, specifying exactly which constructions we are testing, under which constraints. In the next chapter, we present the results of this investigation.

¹⁴The idea of cluster usage is that each job has a requested number of CPUs exclusively for itself, in which case the used CPUs would not be affected of varying loads. However, jobs are not prevented from using more than the requested number of CPUs on the same node, what means that there is no guarantee that the used CPUs are not also used by other jobs, what may put them under varying load.

Chapter 5

Results and Discussion

Contents

5.1	Results of the Internal Tests	57
5.1.1	GOAL Test Set	57
5.1.2	Michel Test Set	65
5.2	Results of the External Tests	68
5.2.1	GOAL Test Set	68
5.2.2	Michel Test Set	71
5.3	Discussion	73
5.3.1	Summary of the Results	73
5.3.2	Gained Insights	74
5.3.3	Limitations of the Study	76

IN THIS CHAPTER, we present and discuss the results of our study. We present statistical evaluations for each of the tests that we outlined in Chapter 4. Note that the evaluations are mostly based on the complement sizes resulting from the different constructions. As mentioned, the complement size is our main performance measure, and the execution time is a secondary one (because it is not perfectly repeatable and comparable). We present a part of the measured execution times in this chapter (namely, from the Michel test set), however, the rest of the execution times can be found in Appendix C.

In Section 5.1, we present the results of the internal tests, and in Section 5.2, the results of the external tests. Both sections treat the results for the GOAL test set and the Michel test set separately. In Section 5.3 we present the discussion of the results, including a result summary, the gained insights, and the limitations of the study.

5.1 Results of the Internal Tests

The goal of the internal tests is to test different versions of the Fribourg constructions and to compare their results against each other. Below, we first present the results of the internal tests on the GOAL test set (Section 5.1.1), and then the results of the internal tests on the Michel test set (Section 5.1.2).

5.1.1 GOAL Test Set

For the GOAL test set, the tested versions of the Fribourg construction are the following (see Section 4.3.1):

1. Fribourg
2. Fribourg+R2C
3. Fribourg+R2C+C
4. Fribourg+M1
5. Fribourg+M1+R2C
6. Fribourg+M1+R2C+C
7. Fribourg+M1+M2
8. Fribourg+R

Below we present the results for all these eight versions from different perspectives. First, we determine the effective samples, on which the rest of the analysis is based. Then, we first analyse the results aggregated over the entire test set, and then split up by the 110 transition density/acceptance density classes of the GOAL test set. Finally, we try to generalise these per-class results by establishing a categorisation of the GOAL test set into *easy*, *medium*, and *hard* classes.

Effective Samples

First of all, we have to check how many timeouts and memory excesses there are in order to establish the effective samples, that is, the set of automata that have been successfully complemented by *all* of the tested constructions. All the remaining analyses in this section are based on this set of effective samples. Table 5.1 shows the number of timeouts and memory excesses for each of the tested Fribourg construction versions.

As can be seen in Table 5.1, there is no memory excess for any of the constructions. That is, none of the constructions required more than 1 GB Java heap for any of the 11,000 complementation tasks in the GOAL test set. However, as can be seen, there are some timeouts. For example, Fribourg, has 48 timeouts, which means that 48 of its complementation tasks were aborted, because they exceeded the time limit of 600 CPU time seconds per automaton.

Based on these results, we have to determine the effective samples, that is, the set of automata that have been successfully processed by *all* the eight constructions. The analysis yields a set of effective samples

Construction	Timeouts	Memory excesses
Fribourg	48	0
Fribourg+R2C	30	0
Fribourg+R2C+C	54	0
Fribourg+M1	2	0
Fribourg+M1+R2C	1	0
Fribourg+M1+R2C+C	8	0
Fribourg+M1+M2	1	0
Fribourg+R	48	0

Table 5.1: Number of timeouts and memory excesses in the internal tests on the GOAL test set.

of 10,939 automata. This means that 61 automata (0.55%) are excluded because their complementation failed for at least one of the constructions. All the following analyses will be based on these 10,939 automata of the effective samples.

Overall Results

Having determined the effective samples, we can turn to the analysis of our main performance measure, the sizes of the produced complements. To get a first impression, in Figure 5.1 we plot all the measured complement sizes in a stripchart. The complement sizes of all the tested constructions are represented as a strip containing a circle for each of the 10,939 automata of the effective samples. The number of states of the corresponding complements are indicated on the x -axis.




Figure 5.1: Complement sizes of the 10,939 effective samples for each tested version of the Fribourg construction on the GOAL test set.

The first thing to note is that the distribution of complement sizes is right-skewed (also known as positive-skewed). This means that there are many small complements, and then gradually fewer larger complements. This results in a long tail toward the right on the x -axis. Generally, in right-skewed distributions, the mean is larger than the median, because the mean is increased by the few number of larger or extremely large values. We will see that this is the case for our complement size distribution further below.

The stripchart allows to intuitively compare the density of large complements that are produced by the different versions. Going from top to bottom, the distributions of Fribourg and Fribourg+R2C have similarly long tails. Fribourg+R2C+C, however, has a considerably longer tail. This shows us that the

C option has a significant effect on the complement sizes, because it adds an additional state to the automata which are not complete. As we have seen in Section 4.2.1, only 9% of the automata of the GOAL test set are already complete. Thus 91% of the automata are affected by the C option.

Next, Fribourg+M1, Fribourg+M1+M2, and Fribourg+M1+R2C all have similarly long tails. However, these tails are significantly shorter than the ones of the previous three versions. This indicates that the M1 optimisation is very effective in reducing the complement sizes.

Fribourg+M1+R2C+C again has a longer tail than Fribourg+M1+R2C. The reason for this is most likely the effect of the C option that we already observed for Fribourg+R2C and Fribourg+R2C+C.

Finally, Fribourg+R has the shortest tail of all versions. Fribourg+R is a special case, because it is the same construction as Fribourg, but with the difference that at the end of the construction, all unreachable and dead states are removed. Comparing the strips of Fribourg and Fribourg+R thus gives an idea of how many unreachable and dead states the Fribourg construction produces for the tested automata.

The stripchart in Figure 5.1 gives a good first impression about the distribution of produced complement sizes of the different versions. However, for further insights, we need to statistically evaluate these distributions. Table 5.2 shows the mean complement size for each version, along with the classical five-number summary consisting of the minimum value, 25th percentile, median, 75th percentile and maximum value.

Construction	Mean	Min.	P25	Median	P75	Max.
Fribourg	2,004.6	2	222.0	761.0	2,175.0	37,904
Fribourg+R2C	1,955.9	2	180.0	689.0	2,127.5	37,904
Fribourg+R2C+C	2,424.6	2	85.0	451.0	2,329.0	54,648
Fribourg+M1	963.2	2	177.0	482.0	1,138.0	16,260
Fribourg+M1+R2C	937.7	2	152.0	447.0	1,118.0	16,260
Fribourg+M1+R2C+C	1,062.6	2	83.0	331.0	1,208.5	25,002
Fribourg+M1+M2	958.0	2	181.0	496.0	1,156.5	15,223
Fribourg+R	136.3	1	1.0	1.0	21.0	12,312

Table 5.2: Statistics of the complement sizes of the 10,939 effective samples for each tested version of the Fribourg construction on the GOAL test set.

As can be seen in Table 5.2, the mean is for all constructions significantly higher than the median. This is typical for right-skewed distributions. Generally, the median is said to be more robust than the mean, because it is not affected by extreme values. For example, for the distribution (1, 2, 3, 4, 5) the mean and the median are both 3. If this distribution happens to contain an extreme value, for example, (1, 2, 3, 4, 1000), then the mean is increased to 202, while the median is still 3. Thus, the median better represents the “typical” values of a distribution, whereas the mean better reflects the effect of extreme values. Both aspects may be important for the analysis of the data at hand. For our own analyses, we will mainly focus on the median, because we are more interested in the “typical” behaviour of the tested constructions. However, we will always put the median in relation with the mean and the other statistical indicators in Table 5.2.

The medians in Table 5.2 include provide interesting insights. Fribourg has a median complement size of 761. This means that half of all the complements produced by Fribourg have 761 or fewer states, and the other half has more than 761 states. For Fribourg+R2C, the median decreases from 761 to 689. This is because the R2C optimisation removes states, whose rightmost component has colour 2, from the output automata. However, this applies only to the input automata which are complete, and as we determined in Section ??, these are only 9% of the GOAL test set. Considering this, the R2C optimisation seems to be very effective in reducing complement sizes.

Regarding Fribourg+R2C+C the median complement size drops significantly to 451. This is surprising insofar as Fribourg+R2C+C has the highest mean of all versions, and in the stripchart in Figure 5.1, it has the longest tail. Thus at first glance, Fribourg+R2C+C might look like the version with the worst performance. However, regarding the median (as well as the 25th percentile) Fribourg+R2C+C belongs actually to the best versions. On the other hand, regarding the 75th percentile and the maximum,

Fribourg+R2C+C has actually the highest values. A possible explanation for this is that the R2C+C option combination makes easy complementation tasks easier and hard complementation tasks harder. We will further elaborate on this point when analysing the median complement sizes separately for all the transition density/acceptance density classes.

Regarding Fribourg+M1 and Fribourg+M1+M2, the median of Fribourg is lower than the median of the latter (482 against 496). The same applies to the 25th and 75th percentile. This means that the application of the M1 optimisation alone results in smaller complements than the application of the M1 and M2 optimisations together. At first sight this is surprising, because Fribourg+M1+M2 has a lower worst-case complexity than Fribourg+M1 (see Section 3.2). However, as we already hinted at in Section 1.2.2 in the introductory chapter, a lower worst-case state complexity does not mean smaller complements in concrete cases. With the results of Fribourg+M1 and Fribourg+M1+M2 we have an empirical evidence for this claim. Note that the mean of Fribourg+M1+M2 is slightly lower than the mean of Fribourg+M1, which might result from the cases where Fribourg+M1 produces larger complements than Fribourg+M1+M2. However, regarding the median, 25th percentile, and 75th percentile values, we still see Fribourg+M1 as the more performant version for the GOAL test set, as we have indicated in Section 4.3.1.

Fribourg+M1+R2C decreases the median of Fribourg+M1 from 482 to 447. Also the 25th and 75th percentile are decreased. This results from the removal of states that the R2C option allows for the 9% of complete automata in the GOAL test set.

Comparing Fribourg+M1+R2C with Fribourg+M1+R2C+C, we observe a similar phenomenon as for Fribourg+R2C and Fribourg+R2C+C. The median and 25th percentile drop significantly, whereas the mean, 75th percentile, and maximum value increase. Again it seems that making the 91% of incomplete automata preliminarily complete (by the C option) in order to apply the R2C optimisation to all of them, has a very positive effect on some automata, but a negative effect on others.

Finally, the results of Fribourg+R make apparent that this version is a special case. The median is 1, and further analysis reveals that all complements up to the 61st percentile have a size of 1. This number coincides well with the 61.8% of automata in the GOAL test set that are universal (see Section ??). These universal automata indeed have a minimum complement consisting of a single non-accepting state (an empty automaton). It seems that whatever complement the Fribourg construction produces for these automata, the R option prunes all but one state. Even though universal automata are special cases, the results of Fribourg+R still give an idea about the number of unreachable and dead states that are produced by the Fribourg construction.

Results by Transition Density/Acceptance Density Class

Up to now, we looked at the results that are aggregated over the entire test set. This mixes together all the automata of the 110 transition density/acceptance density classes. In this part, we analyse the results for each of these 110 classes separately. The goal is to reveal how the constructions perform on automata with very different characteristics, and to reveal the relative differences between them.

Such class-based analysis means that there is not just one value per statistical indicator for each construction, but 110 values, one for each class. For example, for every tested construction there are 110 means, 110 medians, and so on. In order to be able to present this amount of data, we restrict ourselves to the median values of each class. As already mentioned, our main focus is on the median, because it best reflects the “typical” results. Consequently, all the reported values in the following are median values.

The per-class median complement sizes result in a similar type of data that we encountered in the analysis of the complete, universal, and empty automata in the GOAL test set in Section 4.2.1. There, we presented this data in two forms, as matrices and as perspective plots. The advantage of matrices is that they show the explicit values, the advantage of perspective plots is that they better show the relative differences and the overall pattern. For the sake of brevity, we use only perspective plots in this chapter. However, we provide the matrices corresponding to all the perspective plots of this chapter in Appendix B.

Figures 5.2 and 5.3 show the perspective plots with the per-class median complement sizes for the eight tested Fribourg construction version. Note that looking at these perspective plots corresponds to looking at the corresponding matrix from the bottom right corner. We keep this orientation for all the perspective plots in this chapter. Note that this orientation is different from the orientation used for the perspective plots presenting the number of complete, universal, and empty automata of the GOAL test set in Section 4.2.1. The colours of the rectangles (called facets) in the perspective plots depend on the “height” of the plot at their four corners, and are chosen to draw an analogy with mountains.




Figure 5.2: Median complement sizes of the 10,939 effective samples of the internal tests for each of the 110 transition density/acceptance density classes of the GOAL test set.

As can be seen in the perspective plots of Figures 5.2 and 5.3, there are large differences of the median complement sizes between the 110 transition density/acceptance density classes. All constructions exhibit a mountain (in some it is only a hill) roughly in the area between transition densities of 1.2 and 2.4, and acceptance densities of 0.1 and 0.9. The mountain is oblong, and its ridge runs across the entire spectrum of acceptance densities. The summit of the ridge is roughly at a transition density of 1.6. In the direction of lower acceptance densities, the mountain keeps its high altitude. In the direction of higher acceptance densities, on the other hand, the mountain gradually flattens down until nearly ground level at an acceptance density of 1.0.

Comparing the height of the mountains in the perspective plots with the overall median complement sizes in Table 5.2 might be surprising at first. The mountains have a height of roughly 2,000–5,000 states (except Fribourg+R), which corresponds to the median complement sizes of the corresponding classes. On the other hand, the overall median complement size of all construction is 761 states or less. However,

a closer look at the perspective plots (and the corresponding matrices in Appendix B) reveals that around half of the classes have a rather low median complement size ($< 1,000$) what justifies the low overall median complement sizes in Table ???. At the same time, this highlights the importance of a per-class analysis of the results on the GOAL test set. It reveals much more accurately what to expect of the complementation of a specific automaton that resembles one of the classes of the GOAL test set.

Note that in the following, we will refer to specific points in the perspective plots as coordinates of the form (t, a) , where t stands for the transition density, and a for the acceptance density. For example, $(1.6, 0.3)$ means the point at a transition density of 1.6 and an acceptance density of 0.3.

Let us now turn to the relative differences between the tested constructions. The plots of Fribourg and Fribourg+R2C in Figure 5.2 (a) and (b) are rather similar. For both, the ridge has between 3,500 and 4,000 states with a peak of around 4,900 states at $(1.6, 0.3)$. In fact Fribourg+R2C only improves the performance on the complete input automata, compared to Fribourg. As we have analysed in Section 4.2.1, the density of complete automata is highest in classes with high transition densities. By looking closely at the two plots, the values for the high transition density classes are indeed slightly lower for Fribourg+R2C than for Fribourg. Obviously, for classes that do not contain any complete automata, the values of the two constructions are identical.

Fribourg+R2C+C in Figure 5.2 (c) has the highest mountain of all constructions. Its ridge has approximately 5,000 states, and it has a peak at $(1.6, 0.3)$ of more than 6,000 states. This means basically that with Fribourg+R2C+C the “mountain classes” (the classes that result in large complements), result in even larger complements than for the other constructions. This might be surprising again, because as can be seen in Table 5.2, Fribourg+R2C+C has one of the lowest *overall* median complement sizes. However, by comparing the classes with low median complement sizes of Fribourg+R2C+C with the other constructions, we see that they are indeed often even lower. This means that while the “mountain classes” have exceptionally high median complement sizes, the “flatland classes” have exceptionally low ones. This supports our previous proposition that the R2C+C combination, which includes the adding of a sink state to incomplete automata (which constitute 91% of the GOAL test set), makes easy complementation tasks easier and hard ones harder.

Fribourg+R in Figure 5.2 (d) has, as expected, extremely low values for all the classes. In particular, 68 of the 110 classes have a median complement size of 1 (see corresponding matrix in Figure B.2 (d) in Appendix B). We already hinted previously at the relation of a median complement size of 1 and the universal automata in the GOAL test set. The per-class results of Fribourg+R reveal that all the classes with a median complement size of 1 contain 50 or more universal automata (compare matrices in Figure B.2 (d) in Appendix B and Figure 4.2 (b)). If all these universal automata are reduced by the R option to a single-state automaton, then it explains the observed median complement size of 1 for these classes (remember that there are 100 automata in each class). However, also the classes which contain less than 50 universal automata have significantly lower complement median sizes than in the Fribourg construction without the R option. This indicates that the Fribourg construction also generates numerous unreachable and dead states for non-universal automata.

Figure 5.3 shows the perspective plots for Fribourg+M1, Fribourg+M1+R2C, Fribourg+M1+R2C+C, and Fribourg+M1+M2. All of these versions include the M1 optimisation. The first three versions, Fribourg+M1, Fribourg+M1+R2C, and Fribourg+M1+R2C+C can be seen as pairs to Fribourg, Fribourg+R2C, and Fribourg+R2C+C, that additionally include the M1 optimisation. Most apparent by comparing these three pairs is that for the versions including M1 the mountain is significantly flatter. For Fribourg+M1 and Fribourg+M1+R2C the ridge has less than 2,500 states, as opposed to around 3,500 states for Fribourg and Fribourg+R2C. For Fribourg+M1+R2C+C the top of the ridge has around 3,000 states, as opposed to 5,000 states for Fribourg+R2C+C. This shows that the M1 optimisation causes a significant improvement for the automata of the “mountain classes”, that is, the automata that generally result in larger complements. However, the M1 optimisation also has a positive effect on the automata of the “flatland classes”, that is, the automata that typically result in smaller complements. This can be best seen in the corresponding matrices in Appendix B, or by comparing the overall means and medians of the corresponding constructions with and without the M1 optimisation (see Table 5.2).

The relative differences between Fribourg+M1, Fribourg+M1+R2C, Fribourg+M1+R2C+C are very similar to the relative differences between Fribourg, Fribourg+R2C, and Fribourg+R2C+C. Also the reasons for these differences are basically the same, and caused by the R2C and C options, as we have




Figure 5.3: Median complement sizes of the 10,939 effective samples for each of the 110 transition density/acceptance density classes of the GOAL test set.

previously explained for Figure 5.2.

Regarding Fribourg+M1+M2 in Figure 5.3 (d), it is most interesting to compare it to Fribourg+M1, in order to isolate the effect of the M2 optimisation. From Table 5.2, we know that the overall median of Fribourg+M1+M2 is higher than the one of Fribourg+M1 (496 against 482), and the overall mean of Fribourg+M1+M2 is lower than the one of Fribourg+M1 (958.0 against 963.2). Comparing the per-class medians in the perspective plots of the two constructions does not reveal any salient differences, except that the highest values of Fribourg+M1+M2 seem to be slightly lower than the ones of Fribourg+M1. However, comparing the matrices of the two constructions in Figure B.2 (e) and (h) reveals that Fribourg+M1+M2 has lower median complement sizes for almost all classes with acceptance densities between 0.1 and 0.4. On the other hand, Fribourg+M1+M2 has higher median complement sizes for most classes with an acceptance density between 0.5 and 0.9. Thus, the M2 optimisation indeed has a positive effect on the performance, but only on *some* automata. These automata are those that contain not too many accepting states. On the other hand, for automata with 50% or more accepting states, the M2 optimisation decreases the performance, rather than increasing it. For many complementation constructions, automata with few accepting states are harder to complement than automata with a lot of accepting states [54]. This means that the M2 optimisation is especially suitable for these automata that are hard for other constructions. This makes it interesting to apply this optimisation selectively on automata on which a positive effect is expected.

Difficulty Categories

The per-class analysis shows that there are big difference in the complement sizes across the transition density/acceptance density classes of the GOAL test set. Furthermore, there is a certain pattern throughout the results of all construction, namely the mountain and the flatland regions. In the following, we attempt to categorise the classes of the GOAL test set into *easy*, *medium*, and *hard* categories.

To this end, we calculated for each transition density/acceptance density class the average of the median complement sizes of all the tested constructions. This results in a matrix containing the average median complement sizes for each class. Then, we defined two breakpoints at 500 and at 1,600 states that divide the classes into easy, medium, and hard ones. The result of this analysis is shown in Figure 5.4.




Figure 5.4: Difficulty categories *easy* (green), *medium* (yellow), and *hard* (red) of the 110 transition density/acceptance density classes of the GOAL test set.

Figure 5.4, shows a total of 53 easy, 36 medium, and 21 hard classes. The easy classes are mainly those with extreme values. All the classes with a transition density of 0.1, or 2.8 and 3.0, and a high acceptance density of 1.0 are easy. Furthermore, there is a “triangle” of easy classes between transition densities 2.0 and 2.6. and acceptance densities 0.5 and 0.9. The lower the transition density, the higher the acceptance density must be such that the class is classified as easy. The hard classes are roughly those with a transition density between 1.4 and 1.8 and an acceptance density between 0.1 and 0.6. The medium classes finally are located as a “belt” around the hard classes.

It is interesting to see that the hard automata are those with a “medium” transition density between 1.4 and 1.8. This range of transition densities means that every state has on average 1.4 to 1.8 outgoing transitions for each symbol of the alphabet. Somehow this specific level of connectivity seems to make the complementation of an automaton very hard. If the connectivity is lower, complementation becomes easier, and if the connectivity is harder, complementation becomes easier as well. A reason that complementation becomes easier with a lower transition density might be that such automata may contain unreachable and disconnected state. We let further investigations on the influence of the transition density on the complementation difficulty for future work.

Figure 5.4 also shows that automata with a larger number of accepting states are generally easier to complement. This proposition has been made by Tsai et al. [54], and it is the base of the MACC option (maximising acceptance set) of many complementation construction in GOAL (see Section 4.1.1). Summarising, we can say that a transition density between 1.4 and 1.8 makes automata hard to complement, but that this difficulty is alleviated by an increasing acceptance density.

Note that these results are specific to the Fribourg construction, and cannot be readily generalised to other complementation constructions. Furthermore the results are based exclusively on the GOAL test set with its specific types of automata. However, as we will see in the results of the external tests in Section 5.2.1, other complementation constructions show a similar behaviour on the automata of the GOAL test set.

5.1.2 Michel Test Set

The Michel test set consists of the four Michel automata Michel 1, Michel 2, Michel 3, and Michel 4 as pictured in Figure ?? of Chapter 4. These automata have of 3, 4, 5, and 6 states, and an alphabet size of 2, 3, 4, and 5, respectively. The versions of the Fribourg construction that we tested on the Michel test set are the following (see Section 4.3.1):

1. Fribourg
2. Fribourg+R2C
3. Fribourg+M1
4. Fribourg+M1+M2
5. Fribourg+M1+M2+R2C
6. Fribourg+R

Below, we first present the complement sizes of the Michel automata resulting from the different versions. Next, we discuss the corresponding state growths (number of states of the complement relative to the number of states of the input automaton) and bring them in relation with the worst-case state growths. Finally, we present the measured execution times for the individual complementation tasks.

Complement Sizes

In Table 5.3, we present the resulting complement sizes for all versions and all the Michel automata. The column *Fitted curve* contains a function of the form $(an)^n$. This function was obtained by fitting a curve of this form to the four xy -points of each construction, with the x -values being the sizes of the input automata, and the y -values the sizes of the corresponding complements. The used fitting method is the method of non-linear least squares. The column *Std. error* shows the standard error that results from the fit.

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Fribourg	57	843	14,535	287,907	$(1.35n)^n$	0.01%
Fribourg+R2C	33	467	8,271	168,291	$(1.24n)^n$	0.06%
Fribourg+M1	44	448	5,506	81,765	$(1.10n)^n$	0.07%
Fribourg+M1+M2	42	402	4,404	57,116	$(1.03n)^n$	0.12%
Fribourg+M1+M2+R2C	28	269	3,168	43,957	$(0.99n)^n$	0.04%
Fribourg+R	18	95	528	3,315	$(0.64n)^n$	0.35%

Table 5.3: Complement sizes for the Michel automata 1 to 4. These automata have 3, 4, 5, and 6 states, respectively.

The first four column of Table 5.3 reveal the massive state growth that the Michel automata provoke. This can be best seen in the results of Michel 4. This automaton has 6 states and an alphabet size of 5. The first tested version, Fribourg, complements Michel 4 to an automaton with 287,907 states. This is a much higher state growth than we observed for any of the automata in the GOAL test set. Note that the maximum complement size of Fribourg for the GOAL test set is 37,904 (see Table 5.2). However, the GOAL automata have 15 states, and Michel 4 has just 6 states, which makes a big difference with an exponential state growth. At the end of this section, we extrapolate complement sizes of Michel automata with 7 and more states, by assuming the state growth that we observed for the Michel automata up to 6 states.

Fribourg+R2C complements Michel 4 to an automaton with 168,291 states (note that Michel automata are complete, and thus the R2C optimisation has an effect). Compared to the result of Fribourg, this is a reduction of 41.5%. This means that 41.5% of the states produced by Fribourg have a rightmost component with colour 2, which can be removed.

The M1 optimisation of Fribourg+M1 has an even more positive effect than the R2C optimisation. The complement produced by Fribourg+M1 has 81,765 states, which is 71.6% less than the complement of

Fribourg. This indicates that the M1 optimisation is particularly efficient for hard automata.

Fribourg+M1+M2 complements Michel 4 to an automaton with only 57,116 states. This is smaller than the complement of Fribourg+M1, and the same applies to the other Michel automata. This is interesting, because for the GOAL test set, from an overall point of view, Fribourg+M1+M2 is less efficient than Fribourg+M1 (see Table 5.2). However, as we analysed for the individual transition density/acceptance density of the GOAL test set, Fribourg+M1+M2 is actually more efficient than Fribourg+M1 for a certain type of automata. These automata are the ones with an acceptance density up to 0.4. The four tested Michel automata all have an acceptance density of 0.33 or less, and thus fall in this category. Hence, the fact that Fribourg+M1+M2 is more efficient on the Michel automata than Fribourg+M1 coincides with this observation for the GOAL test set.

Fribourg+M1+M2+R2C further decreases the complement size of Michel 4 to 43,957. Again, this is because the R2C optimisation removes all the states whose rightmost component has colour 2. Compared to Fribourg, the complement produced by Fribourg+M1+M2+R2C is 84.7% smaller. In other words, the application of all the three optimisations to the Fribourg construction, cuts the complement size of Michel 4 in approximately six and a half.

Finally, the complement size of Fribourg+R for Michel 4 is 3,315. This means that 284,592, or 98.8%, of the states produced by Fribourg are unreachable or dead states, that are removed by the R option.

State Growth Functions

The heavy state growth of Michel automata make them suitable to be used as lower bounds for the worst-case state complexities of the tested constructions, in case these worst-case complexities have not yet been determined precisely. To this end, we calculated for each construction a state growth function in the second-last column of Table 5.3.

For Fribourg, the state growth of the Michel automata is $(1.35n)^n$, whereas the determined upper bound for the worst-case is $O((1.59n)^n)$ (see Section 3.1.3). This means that whatever the actual worst-case state growth of Fribourg is, it cannot be lower than $(1.35n)^n$, because with our empirically determined complement sizes of the Michel automata, we have a proof that Fribourg can generate $(1.35n)^n$ many states, with n being the number of states of the input automaton.

Regarding the other versions, for Fribourg+R2C this lower bound is $(1.24n)^n$. However, for Fribourg+R2C, we do not have a theoretically calculated upper bound of the worst-case state complexity. For Fribourg+M1, the state growth of the Michel automata is $(1.10n)^n$, whereas the calculated worst-case state growth is $O((1.195n)^n)$ ¹. This means that for Fribourg+M1, the gap between the upper and lower bound is relatively small, and the worst-case state growth cannot be further improved very much.

For Fribourg+M1+M2, the state growth of the Michel automata is $(1.03n)^n$. However, the worst-case state growth has been calculated to be $O((0.86n)^n)$, even with evidence that it might actually be $O((0.76n)^n)$. So, here we seem to have a discrepancy. Whereas we have the empirical proof that Fribourg+M1+M2 can produce $(1.03n)^n$ many states, with n being the number of states of the input automaton, the theoretically calculated results state that Fribourg+M1+M2 can generate at most $O((0.86n)^n)$ or $O((0.76n)^n)$ many states. The reason for this discrepancy might be hidden in the big- O notation. The functions $O((0.86n)^n)$ or $O((0.76n)^n)$ are not exact, but rather the result of omitting constant factors and lower order terms. We leave further investigations on this point for future work.

For the last two versions, Fribourg+M1+M2+R2C and Fribourg+R, the measured state growths are $(0.99n)^n$ and $(0.64n)^n$, respectively. However, for these versions we do not have theoretically calculated upper bounds for the worst-case state complexity.

¹This is actually the worst-case state growth of the lower part of the automaton. For the overall worst-case state growth, the maximum number of states of the upper part of $(0.53n)^n$ must be added to it. However, the resulting term is clearly dominated by $O((1.195n)^n)$.

Execution Times

In Table 5.4, we present the execution times of the different versions for complementing the Michel automata. The reported values are in CPU time seconds. As mentioned, the measurement of execution times is dependent on many factors, such as the implementation and the execution environment, and thus should be interpreted with care. However, for the case of the Michel automata, we think that the execution times still provide interesting insights.

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Fribourg	2.3	4.0	88.8	100,976.0	$(1.14n)^n$	0.64%
Fribourg+R2C	2.3	3.4	27.4	27,938.3	$(0.92n)^n$	0.64%
Fribourg+M1	2.2	3.6	17.9	6,508.4	$(0.72n)^n$	0.63%
Fribourg+M1+M2	2.3	3.5	13.8	2,707.4	$(0.62n)^n$	0.62%
Fribourg+M1+M2+R2C	2.5	3.5	10.8	2,332.6	$(0.61n)^n$	0.62%
Fribourg+R	2.4	3.7	86.0	101,809.6	$(1.14n)^n$	0.64%

Table 5.4: Execution times for the Michel automata 1 to 4. These automata have 3, 4, 5, and 6 states, respectively.

For Michel 1, all versions take slightly more than 2 seconds. Approximately 2 seconds are actually caused by the start-up of the Java Virtual Machine, which is also included in our time measurements. For Michel 2, all versions take 4 seconds or less, and for Michel 3 all versions are still under 1.5 minutes. For Michel 4, however, the execution times seem to explode. Fribourg takes 100,976.0 seconds, which are approximately 28 hours. However, this time is considerably lower Fribourg+R2C, which takes 27,938.3 seconds, or 7.76 hours. Fribourg+M1 takes 6,508.4 seconds, or 1.8 hours. The best version, Fribourg+M1+M2+R2C, takes 2,332.6 seconds, or 39 minutes. Note that Fribourg+R takes the same time as Fribourg, plus an additional delay to remove the unreachable and dead states.

There are two remarkable points in these execution times. First, the increase of the execution times between Michel 3 and Michel 4 is much bigger than the increase of the complement sizes between Michel 3 and Michel 4. For example, for Fribourg, the complement size of Michel 4 is about 20 times bigger than the complement size of Michel 3. However, the execution time for Michel 4 is approximately 1,138 times bigger than execution time for Michel 3. As another example, for Fribourg+M1+M2+R2C the factors are 14 for the complement size and 216 for the execution time. Second, the differences of the execution times between the different versions are larger than the corresponding differences between the complement sizes. For example, for Michel 4, the complement of Fribourg+M1+M2+R2C is around 6.5 times smaller than the complement of Fribourg. However the execution time of Fribourg+M1+M2+R2C is 43 times smaller than the execution time of Fribourg. In summary, this shows that with an increasing number of states the required computation time increases disproportionately.

As promised, below we extrapolate the execution times of larger Michel automata. To this end, we fitted functions of the form $(an)^n$ to the measured execution times in Table 5.4 in the same way as we did for the complement sizes. The parameter n is the number of states of the input automaton, and the value of the functions is the required execution time of the complementation task in CPU time seconds. Based on the fitted functions for Fribourg, we present in Table 5.5 the extrapolated values for the Michel automata 5 to 8.

Automaton	States (n)	Compl. size $(1.35n)^n$	Exec. time $(1.14n)^n$	\approx days/months/years
Michel 5	7	6,882,980	2,020,385	23 days
Michel 6	8	189,905,394	46,789,245	18 months
Michel 7	9	5,939,189,262	1,228,250,634	39 years
Michel 8	10	207,621,228,081	36,039,825,529	1,142 years

Table 5.5: Extrapolated complement sizes and execution times for the Michel automata 5 to 8 based on the fitted growth functions of Fribourg.

As can be seen, the complementation of Michel 5 would take 23 days with Fribourg. This is still a

conservative estimation, because, as we have seen, the execution time grows disproportionately fast with an increasing number of states. A duration of 23 days would already exceed the maximum running time of a job on the cluster we use for our experiments. This shows why we restricted ourselves to include only the Michel automata 1 to 4 in the Michel test set. And even if Michel 5 would still be doable, the execution times of 18 months, 39 years, and 1,142 years for Michel 6, 7, and 8, are most probably too long for an empirical study, even for a master's thesis.

5.2 Results of the External Tests

The external tests consist of the comparison of the best version of the Fribourg construction on each test set against three other complementation constructions. These three other constructions are the following (see Section ??):

- Piterman+EQ+RO
- Slice+P+RO+MADJ+EG
- Rank+TR+RO

Regarding the versions of the Fribourg construction, we declared the best we declared the best one on each test set to be the following:

- Fribourg+M1+R2C for the GOAL test set
- Fribourg+M1+M2+R2C for the Michel test set

In the subsequent two sections, we first present the results of these tests on the GOAL test set (Section 5.2.1), followed by the results on the Michel test set (Section 5.2.2).

5.2.1 GOAL Test Set

In this section, we analyse the results of the external tests on the GOAL test set in a similar way as we did for the internal tests. That is, we determine the effective samples, analyse the complement sizes from an overall point of view, and on a per-class basis for each of the 110 transition density/acceptance density of the GOAL test set.

However, there is an issue that caused us to slightly modify our scheme of analysis. One of the constructions, Rank, exhibits an exceptionally bad performance on the GOAL test set compared to the other constructions. The result is that the set of effective samples is reduced by more than one third. By using this reduced set of effective samples, we exclude a big part of the results of the other three constructions, Piterman, Slice, and Fribourg. To prevent this from influencing the interpretation of the results of these constructions, we present two separate analyses, one with Rank, and one without Rank. In the following, we first discuss the effective samples (with and without Rank), and then present the two versions of the result analysis with and without Rank.

Effective Samples

Table 5.6 shows the number of timeouts and memory excesses for all four tested construction. Like for the internal tests, the timeout is set to 600 seconds CPU time, and the memory limit to 1 GB Java heap size. As can be seen, there is a very high number of excesses for Rank. The construction timed out for 3,713 automata, and exceeded the memory for 83 automata. This gives a total of 3,796, or 34.5%, of aborted complementation tasks. This contrasts with the other constructions. There are only two timeouts for Piterman, and one for Fribourg.

By including all four constructions, there are 7,204 effective samples, which are 65.5% of the test set. On the other hand, by including only Piterman, Slice, and Fribourg, there are 10,998 effective samples. This means that by including Rank, we exclude more than one third of the results of Piterman, Slice, and Fribourg that are actually available. As mentioned, for this reason, in the following, we present two


Construction	Timeouts	Memory excesses
Piterman+EQ+RO	2	0
Slice+P+RO+MADJ+EG	0	0
Rank+TR+RO	3,713	83
Fribourg+M1+R2C	1	0

Table 5.6: Number of timeouts and memory excesses of the external tests on the GOAL test set.

separate analyses, one for the 7,204 effective samples with Rank, and another one for the 10,998 effective samples without Rank.

With Rank

Figure 5.5 shows a stripchart with the complement sizes of the 7,204 effective samples. As can be seen, Rank has a very long tail with a maximum reaching 120,000 states. This stripchart makes apparent that there is a very high number of automata for which Rank produces much larger complements than the other constructions. Regarding the other constructions, Piterman has clearly the highest density of small complements, whereas Slice and Fribourg have similar distribution.

Figure 5.5: Complement sizes of the 7,204 effective samples of the external tests *with Rank* on the GOAL test set.

In Table 5.7, we present the mean and the five-number summary of the complement sizes of the 7,204 effective samples. Regarding Rank, this table bears a surprise. The 25th percentile and the median of Rank are comparable to the other constructions, and even lower than for the Fribourg construction. However, this picture changes dramatically for the 75th percentile, which seems to explode compared to the other constructions. This means that, despite the bad image that Rank conveys in the stripchart and by the mean complement size, for at least half of the 7,204 effective samples, Rank performs perfectly comparable to the other constructions. Only for at most the other half of these automata, Rank performs poorly and results in very large complements. This is a remarkable fact, and it would be an interesting task for future work to investigate which characteristics cause an automaton to be hard for Rank.

Regarding the other constructions in Table 5.5, again Piterman is the clear winner. Slice and Fribourg have comparable results, although the values are in the favour of Slice.

We omit the analysis of the complement sizes for each of the transition density/acceptance density classes, because the low number of effective samples causes many classes to be not represented at all in the effective samples. Instead, in the following, we pass over directly to the second version of the analysis which excludes Rank.

Construction	Mean	Min.	P25	Median	P75	Max.
Piterman+EQ+RO	106.0	1	29.0	58.0	121.0	4,126
Slice+P+RO+MADJ+EG	555.4	2	70.0	202.0	596.0	41,081
Rank+TR+RO	5,255.6	2	81.0	254.5	3,178.2	120,674
Fribourg+M1+R2C	662.9	2	101.0	269.0	754.5	37,068

Table 5.7: Statistics of the complement sizes of the 7,204 effective samples of the external tests *with Rank* on the GOAL test set.

Without Rank

Considering only Piterman, Slice, and Fribourg, there are 10,998 effective samples. Figure 5.6 shows a stripchart with the complement sizes of these 10,998 automata for each construction. The stripchart presents a similar picture as we already saw in the previous analysis with the Rank construction. Piterman clearly has the highest density of small complements, whereas Slice and Fribourg have comparable distributions.




Figure 5.6: Complement sizes of the 10,998 effective samples of the external tests *without Rank* on the GOAL test set.

In Table 5.8, we present the mean and the five-number summary of the complement sizes of these constructions. These numbers confirm that Piterman is by far the most performant construction on the GOAL test set. Its values for mean, median, 25th and 75th percentile are multiple times lower than the values of the other two constructions. Regarding Slice and Fribourg, Slice has the better results than Fribourg. The mean, 25th percentile, median, and 75th percentile of Slice are 6.7%, 21.6%, 12.4%, and 11.6%, respectively, lower than for Fribourg. These results allow to rank the three constructions according to their overall performance on the GOAL test set. First is, by a large margin, Piterman, second is Slice, and third is Fribourg.

Construction	Mean	Min.	P25	Median	P75	Max.
Piterman+EQ+RO	209.6	1	38.0	80.0	183.0	20,349
Slice+P+RO+MADJ+EG	949.4	2	120.0	396.0	1,003.0	41,081
Fribourg+M1+R2C	1,017.3	2	153.0	452.0	1,134.0	37,068

Table 5.8: Statistics of the complement sizes of the 10,998 effective samples of the external tests *without Rank* on the GOAL test set.

Next, we analyse the complement sizes for each of the 110 transition density/acceptance density classes of the GOAL test set. Like for the internal tests, we consider the median complement size for each class. Figure 5.7 shows the perspective plots for Piterman, Slice, and Fribourg. Note that we provide the corresponding matrices with the exact values in Appendix B (Figure B.1). Also note that the perspective plot of Fribourg+M1+R2C in Figure 5.7 (c) is the same as the one we presented for the internal tests in Figure 5.3 (c), however, with a different z -axis scale.




Figure 5.7: Median complement sizes of the 10,998 effective samples of the external tests *without Rank* for each of the 110 transition density/acceptance density classes of the GOAL test set.

The perspective plots reveal for Piterman a significantly lower median complement size throughout all classes. It seems as if Piterman is more performant by a constant factor for all types of automata. However, the basic pattern, namely the mountain around a transition density of 1.6 and low acceptance densities, is the same for Piterman as for the other two constructions, merely in a reduced form. Regarding Slice and Fribourg, the two constructions exhibit a very similar pattern. A small difference is that Fribourg has slightly worse values for classes with a transition density up to 2.0, but often slightly better values for classes with a higher transition density. This means that, roughly, Fribourg performs slightly worse than Slice on the hard automata and slightly better on many easy automata.

The similar patterns of the different constructions in Figure 5.7 show that the classification into easy, medium, and hard automata of the GOAL test set, that we established for the Fribourg construction (see Section 5.1.1), is also applicable to other complementation constructions. At least for Piterman and Slice, the same automata seem to be hard and easy as for the Fribourg construction.

5.2.2 Michel Test Set

For the Michel test set, we compare Fribourg+M1+M2+R2C against the same three versions of Piterman, Slice, and Rank as for the external tests on the GOAL test set.

Complement Sizes

Table 5.8 shows the complement sizes of the Michel automata 1 to 4 for each of the tested constructions. Like for the internal tests, we fitted a function of the form $(an)^n$ to the measured complement sizes, and provide the standard error resulting from the fit in the last column.

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Piterman+EQ+RO	23	251	5,167	175,041	$(1.25n)^n$	0.29%
Slice+P+RO+MADJ+EG	35	431	6,786	123,180	$(1.18n)^n$	0.02%
Rank+TR+RO	23	181	1,884	25,985	$(0.91n)^n$	0.01%
Fribourg+M1+M2+R2C	28	269	3,168	43,957	$(0.99n)^n$	0.04%

Figure 5.8: Complement sizes for the Michel automata 1 to 4, with 3, 4, 5, and 6 states, respectively.

The values in Table 5.8 contain a series of surprising results. First, Rank produces the smallest complements for all Michel automata. Note that this is the same Rank construction which performs exceptionally bad on the GOAL test set, and which therefore we had to exclude from the result analysis of the GOAL test set. For the Michel automata, however, Rank beats all the other constructions by a large margin.

The next surprise is that Piterman has the worst result for Michel 4. Again, this is the same Piterman construction that has by far the best performance on the GOAL test set. However, note that the bad

performance of Piterman applies particularly to Michel 4. For Michel 3, Piterman has only the second-worst result, for Michel 2, it has the second-best, and for Michel 1, it has even the best result (together with Rank). It would be interesting to test how the performance of Piterman evolves with Michel 5 and beyond. However, as we have demonstrated in Section 5.1.2, for practical reasons this is nearly impossible.

Another surprise is that Fribourg performs relatively well on the Michel automata. For Michel 3, Fribourg is ranked second behind the unbeatable Rank with a large gap to the third-placed construction (Piterman). For Michel 4, Fribourg is also ranked second behind Rank, but with an even much larger gap to the third-placed construction (Slice). In particular, the complement of Fribourg for Michel 4 is 2.8 times smaller than the complement of Slice, and 4 times smaller than the complement of Piterman.

In summary, if we take the fitted curves (or the results on Michel 4) as the overall measure, then the ranking of the four construction is, first Rank, second Fribourg, third Slice, and fourth Piterman. Note that this is the reverse ranking of the same constructions on the GOAL test set, where the best is Piterman, second Slice, third Fribourg, and fourth Rank (however, the version of the Fribourg construction used for the GOAL test set is Fribourg+M1+R2C instead of Fribourg+M1+M2+R2C). This shows that it is generally not possible to declare a certain construction to be “better” than another in a universal sense. A construction can only be better than another with respect to a certain automaton, or set of automata.

Execution Times

As for the internal tests, we also present the execution times of the external tests on the Michel automata. Table 5.9 shows the measured execution times in CPU time seconds. The second-last column of the table again contains a growth function of the form $(an)^n$ that we fitted to the the input automaton sizes and corresponding execution times. The last column of the table contains the standard error of the fit.

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Piterman+EQ+RO	2.5	3.8	42.6	75,917.4	$(1.08n)^n$	0.64%
Slice+P+RO+MADJ+EG	2.3	3.6	11.4	159.5	$(0.39n)^n$	0.38%
Rank+TR+RO	2.2	3.0	6.4	30.0	$(0.29n)^n$	0.18%
Fribourg+M1+M2+R2C	2.5	3.5	10.8	2,332.6	$(0.61n)^n$	0.62%

Table 5.9: Execution times for the Michel automata 1 to 4. These automata have 3, 4, 5, and 6 states, respectively.

As can be seen in Table 5.9, the execution times of all constructions stay relatively low until Michel 3. This is basically the same phenomenon that we observed for the internal tests on the Michel test set. But then, for Michel 4, the differences between the constructions are enormous. Rank takes just 30 seconds to complement Michel 4. On the other hand, Piterman takes 75,917.4 seconds, which are approximately 21 hours. Slice is with 159.5 seconds also very fast. On the other hand Fribourg is with 2,332.6 seconds, or 39 minutes, significantly slower than Rank and Slice, but still much faster than Piterman.

From the results of the internal tests on the Michel test set, we learned that the execution time increases diproportionately with the number of generated states. This explains in parts the large time gap between Rank and Piterman for Michel 4. While Piterman’s complement of Michel 4 is around 6.7 times larger than Rank’s complement of Michel 4, Piterman’s execution time is approximately 2,530 times longer than Rank’s execution time. However, the extent of this disproportionate growing of the execution time is still surprising.

Another remarkable point is that Fribourg has a very long execution time for Michel 4 compared to Slice. This cannot be explained by the disproportionate growing of the execution time, because the complement produced by Fribourg is smaller than the complement produced by Slice. In detail, the complement of Fribourg is 2.8 times smaller than the complement of Slice, however, the execution time of Fribourg is 14.6 times greater than the execution time of Slice. There are two possible explanations for this. Either Slice is exceptionally fast on Michel 4, or Fribourg is exceptionally slow. To find this out, we would need to include further constructions in the test, which is a point that could be done in future work.

This concludes the presentation of the results of our empirical performance investigation of the Fribourg construction, consisting of the internal and external tests, on the GOAL test set and the Michel test set. In the next section, we discuss important aspects of these results, and of the study in general.

5.3 Discussion

5.3.1 Summary of the Results

Below we recapitulate the most important results from earlier in this chapter. We first review the results related to the GOAL test set, followed by the results related to the Michel test set.

GOAL Test Set

Regarding the different versions of the Fribourg construction, we identified Fribourg+M1+R2C as the version with the best overall performance on the GOAL test set. Interestingly, the application of the M1 and M2 optimisations together (Fribourg+M1+M2) results in a worse overall performance on the GOAL test set than the application of the M1 optimisations alone (Fribourg+M1). However, we found out that this penalty introduced by Fribourg+M1+M2 does not apply to all automata types. Actually, Fribourg+M1+M2 has a better performance than Fribourg+M1 for all automata classes with an acceptance density between 0.1 and 0.4. However, for most other automata classes, Fribourg+M1+M2 performs indeed worse than Fribourg+M1, and this is enough to tilt the overall performance in the favour of Fribourg+M1. Another result is that the application of the R2C optimisation together with the C options results in many more large complements, but at the same time in many more small complements than the corresponding version without R2C and C. A way to see this is that R2C and C together make easy complementation tasks easier, an harder complementation tasks harder.

Regarding the other tested constructions, Piterman has a significantly better performance on the GOAL test set than all the other constructions. This is true from an overall point of view and for each of the transition density/acceptance density classes of the GOAL test set. The other extreme is Rank, which exhibits significantly worse performance on the GOAL test set than all the other constructions. The best version of the Fribourg construction on the GOAL test set (Fribourg+M1+R2C) and Slice have a comparable performance on the GOAL test set, although there is a slight advantage for Slice.

We identified criteria that make the automata of the GOAL test set result in larger or smaller complements, or in other words, make automata harder or easier to complement. For the Fribourg construction, the automata with a transition density between 1.4 and 1.8, and an acceptance density between 0.1 and 0.6 are the hardest complementation tasks. A higher or lower transition density makes the complementation of an automaton by trend easier. Furthermore, a higher acceptance density makes the complementation of an automaton by trend easier as well. Our results of Piterman and Slice suggest that these criteria also apply to other complementation constructions.

Michel Test Set

We identified Fribourg+M1+M2+R2C as the best version of the Fribourg construction on the Michel test set. Interestingly, for the Michel automata, the application of the M1 and M2 optimisation together (Fribourg+M1+M2) results in a significantly better performance than the application of the M1 optimisation alone (Fribourg+M1). However, as we identified for the GOAL test set, there are indeed certain types of automata for which Fribourg+M1+M2 performs better than Fribourg+M1. The Michel automata seem to belong to these type of automata. Generally, the R2C, M1, and M2 optimisations significantly increase the performance of the Fribourg construction on the Michel automata. Their simultaneous application (Fribourg+M1+M2+R2C) cuts the complement size of the largest Michel automata (Michel 4) by 6.5, and the execution time by 43.3, compared to the Fribourg construction without optimisations.

Regarding the other tested complementation constructions, Rank has by far the best performance on the Michel automata. This is one of the most surprising results, because Rank has by far the worst overall

performance on the GOAL test set. On the other hand, Piterman has by far the worst performance on the largest Michel automaton (Michel 4). This is again surprising, because Piterman has by far the best performance on all the automata types in the GOAL test set. The best version of the Fribourg construction on the Michel test set (Fribourg+M1+M2+R2C) has the second-best overall performance on the Michel test set after Rank. The gap between Fribourg and the first-placed Rank is smaller than the gap between Fribourg and the third-placed Slice. However, the execution time of Fribourg for Michel 4 is significantly greater than the execution time of Rank and Slice. In general, Rank and Slice can be seen as extremely fast on the largest tested Michel automaton (Michel 4).

5.3.2 Gained Insights

Reasonable Expectations for Complement Sizes

A question to which our results provide insights is: given an example automaton with 15 states and an alphabet size of 2 (similar to an automaton of the GOAL test set), what can we expect from its complementation? In which order of magnitude can we expect the size of the complement to be? With our results on the GOAL test set, we have a large set of empirical data that allow us to draw such conclusions. Our data suggests that a reasonable expectation for the complement size is a couple of hundred states. This corresponds to the median complement sizes of the tested constructions (except Piterman, which has a median complement size below 100 states). A reasonable worst-case to expect, given such an automaton, is a couple of ten-thousand states. This corresponds to the maximum complement sizes of the tested constructions (however, note that we do not know how big the complements of the aborted complementation tasks would be).

Such guidelines can only be provided by empirical studies like our own. Merely knowing the worst-case state complexity of a construction does not allow to make a reasonable expectation of the complement size for a concrete automaton. For example, if we would have a so-called “optimal” construction with a worst-case state complexity matching the lower bound of $(0.76n)^n$ [68], then we could only deduce that in the worst case the complement has 7.14 quadrillion states (7.14×10^{15}). This is neither a reasonable expectation (the probability that the current automaton is such a worst case is extremely low), nor is it a helpful information from a practical point of view. It is not a worst case we could prepare our computation environment for, because such a case is most probably infeasible in practice anyway.

Our results on the GOAL test set allow furthermore to refine such expectations based on the transition density and acceptance density of the automaton at hand. For example, if the automaton has a transition density of 3.0 and an acceptance density of 1.0, then a reasonable expectation of the complement size is below 100 states (based on the median complement size of this class of the GOAL test set). On the other hand, if the automaton has, for example, a transition density of 1.6 and an acceptance density of 0.3, then a reasonable expectation would be a couple of thousand states (except for Piterman, in which case it would be below one thousand states).

What Makes Automata Hard or Easy

Another question of practical interest is: can we tell whether an automaton is likely to be hard or easy to complement, before actually complementing it? Our results of the complement sizes for the different transition density/acceptance density classes of the GOAL test set provide evidence that this is possible. We identified a range of transition densities that cause an automaton to be harder to complement, and we discovered that a lower acceptance density by trend makes the complementation task harder too. These results are based only on the Fribourg construction and on the specific automata of the GOAL test set. However, it seems likely that different constructions exhibit a similarly structured behaviour on different automata. Hence, by further investigations on this point, it could be possible to establish generally valid, or probably construction-specific, rules for assessing the difficulty of a given automaton.

Worst-Case Performance Unrelated to Actual Performance

Our results confirm the proposition that “worst-case bounds are poor guides to actual performance” [54]. This means that how good or bad a construction’s worst-case state complexity is, does not allow to conclude on how well or bad the construction performs on concrete cases. For example, Fribourg+M1+M2 has a lower worst-case state complexity than Fribourg+M1, however, we discovered that Fribourg+M1+M2 performs worse on the GOAL test set (from an overall point of view) than Fribourg+M1. Another example is Schewe’s rank-based construction [45], which is included in our tests as Rank. This construction has a very low worst-case state complexity of $(0.76n)^n n!$. However, on the concrete automata of the GOAL test set, Rank performs by far worse than all other tested constructions, even though all these other constructions have a higher worst-case state complexity than Rank. On the other hand, Rank has by far the best performance on the Michel automata. This indicates that the worst-case performance is really unrelated to the actual performance, and there seem to be no rules relating the worst-case performance of a construction with its actual performance.

No Overall-Best Construction

Another insight that our results provide is that there seems to be no “best” construction from an overall point of view. A construction can only be better than another construction with respect to a specific automaton or set of automata. However, this is no guarantee that the construction performs also better on different automata. For example, in our experiments, Piterman had clearly the best performance on the GOAL test set, throughout all transition density/acceptance density classes. However, for a different set of automata, such as the Michel automata, Piterman is not the best construction, and even by far the worst construction for the Michel automaton with six states (Michel 4).

Therefore, a more adequate characterisation for the merits of complementation construction is that they have relative strengths and weaknesses. Consequently, instead of trying to find an overall-best construction, a better approach to practical Büchi complementation might be to identify a best-suitable construction for the concrete automaton, or set of automata, at hand. This leads to a portfolio approach to algorithm selection [27], which has been proposed for Büchi complementation by Vardi [62]. This means that we match complementation constructions to complementation tasks, based on different criteria, such as complement size, running time, or memory usage.

Our results support this selective approach. For example, for automata that are similar to the automata in the GOAL test set, we would choose Piterman as the best-suitable construction. On the other hand, for the Michel automata we would clearly choose Rank. If we want to use only the Fribourg construction, then for automata similar to the ones in the GOAL test set with an acceptance density up to 0.4, we would choose Fribourg+M1+M2. On the other hand, for automata with a higher acceptance density, we would choose Fribourg+M1. Furthermore, as we have seen, the R2C and C options together make some complementation tasks easier and some harder. If we can identify the criteria that make automata easier with the R2C+C combination, then we could choose Fribourg+R2C+C (or a variation of it with M1 and M2) for such automata, and obtain smaller complements than without R2C+C. With such a selective approach to construction selection, we can harness the strengths of many different constructions and avoid their weaknesses. In the end, this results in better overall results, than relying on a single construction.

Limited Understanding of Büchi Complementation

A general insight is that we know little about the mechanisms that make Büchi complementation constructions behave and perform the way they do. The concrete behaviour and performance of a construction is very hard to predict, and usually not revealed until it is empirically tested. For example, we would not have anticipated that Rank produces so many very large complements for the automata in the GOAL test set. On the other hand, it produces the smallest complement for the Michel test set. As another example, it is not clear from the outset why the M2 optimisation of the Fribourg construction results in degradation of the performance in many cases (but not all cases). To find explanations for these behaviours would be a topic for further research. It shows that at the moment we are at the very beginning of understanding the behaviour of Büchi complementation constructions, and thus of Büchi complementation in general.

5.3.3 Limitations of the Study

In this section, we discuss the limitations of our study. Some of these limitations also affect and relativise the insights that we presented in the previous section. At the same time of being limitations, we see the following points also as suggestions for future work. That is, the shortcomings of our study can give an idea of how future similar studies can be improved.

Test Data

The data we use for evaluating the different constructions is limited in several ways. Regarding the GOAL test set, all automata have 15 states and an alphabet size of 2. This is a very specific type of automata (even though the number of transitions and accepting states of these automata vary considerably). We think that it is not readily possible to generalise the results on automata with 15 states to automata with, for example, 50, 100, or 500 states. The same applies to the alphabet size. This affects the reasonable expectations of the complement sizes of concrete automata, that we identified as one of the insights gained by our study in the last section. Our results allow such expectations only on automata with similar specifications as the GOAL test set automata. If, for example, an automaton with 15 states has an alphabet size of 3 instead of 2, then we have no clear hints about what to expect from its complement. The same applies to the criteria that make an automaton hard or easy to complement, which we presented as another insight of the study. These results are also exclusively based on the GOAL test set, and we do not know whether they can be generalised to other types of automata.

Another limitation of the test data are the characteristics of the automata in the GOAL test set itself. We have identified that 61.8% of all automata in the GOAL test set are universal (see Section 4.2.1). Of the 110 transition density/acceptance density classes, 30 classes even contain 95% or more universal automata (see matrix in Figure 4.2 (b)). Universal automata are a special case of automata, and we do not know if and how their high representation affects our results. Furthermore, In the practical application of Büchi complementation, it might be rarely the case that we have to complement a universal automata. For example, in automata-theoretic model checking, we have to complement the automaton representing the property to verify, and it is unlikely that this automaton is universal. Consequently, this bias of the GOAL test set towards universal automata might affect the significance of our study for practical applications of Büchi complementation.

Regarding the Michel test set, we were only able to include the first four Michel automata with up to 6 states (Michel 1 to 4). This might not be enough to represent the behaviour of the tested constructions on other Michel. For example, for claiming more reliably that Rank is very good and Piterman is very bad for Michel automata in general, we would require results of more Michel automata.

Generally, our test data does not include, or resemble, real-world examples. Automata that need to be complemented in practice, for example, in automata-theoretic model checking, are likely to have more than 15 states and more than 2 alphabet symbols. According to Tasiran et al. [50], an automaton for representing a real-world model checking property of the form “if X , then eventually Y ” has several hundred states, and this is still considered a small automaton. As mentioned, we do not know if and how our results of automata with 15 states can be generalised to automata with several hundreds of states as they are used in practice.

Reliance on Implementation

Naturally, the results of our empirical investigations rely on implementations of the involved constructions. Regarding the constructions that are pre-implemented in GOAL, we do not have any details of their implementation (as GOAL is not open-source). Concerning the correctness, we know that the constructions are cross-checked by their authors [57] (in a similar way as we checked the correctness of our implementation of the Fribourg construction), and it is reasonably to assume that they are correct, in the sense of producing correct complements. However, we do not know reliably the specifications according to which the implementations have been created. All the basic constructions in GOAL are mapped to a specific specification in the literature (see Table 4.1). However, it is still possible that they contain modifications and variants. Furthermore, most constructions contain options which rely on propositions

from various sources. This makes it even more difficult to map different versions of an implementation to a specific specification from the literature. This means that, in general, the mapping of the results of an implementation back to a construction, as devised by its authors, must be done with care.

Another limitation resulting from the reliance on implementations is that measures such as the execution time or memory usage are not comparable. An implementation can be more or less efficient in terms of speed or memory usage. Consequently, if, for example, an implementation of a construction has a lower execution time than another, we cannot tell whether the construction itself is more efficient in terms of time than the other constructions. The same applies to the comparison of the results of one study with the results of another study. If the studies use different implementations of the same construction, then implementation-dependent measures like execution time and memory usage can be very different. Even if the studies use the same implementation but execute it in a different execution time, these results may vary considerably (as it is, for example, the case for our own study and the study by Tsai et al. [54]).

Last, the implementations we used for our study are in Java (because GOAL is written in Java). However, Java might not be an optimal language for tasks that are as computationally heavy as Büchi complementation. An implementation in a different language might greatly reduce the execution time and memory usage. For example, Göttel [14] implemented the Fribourg construction (with the R2C optimisation) in the C programming language and tested this implementation with the GOAL test set. The execution times of his implementation are on average tens of times faster than the execution times of our Java implementation (note that the execution time also depends on the execution environment). Such an efficient implementation has the advantage that the need for setting time and memory limits could probably be overcome. It would become more easily possible to test larger automata (for example, up to some hundred states).

Test Setup

For the Fribourg construction, we first evaluated different versions against each other (internal tests), and then chose the best one on each test set for comparing it with other constructions (external tests). However, for the other constructions, we did not do such an evaluation. We just selected the version that includes the default options of the constructions. Consequently, we do not know whether these versions are also the best ones for the respective test sets. A fairer setup would consider the performance differences between the different versions of all the constructions. For example, if we would compare the least performant version of the Fribourg construction on the two test sets, Fribourg, to the other constructions, then the Fribourg construction would be the least efficient construction most of the time in the external test. We have to keep this in mind when concluding from the results of a specific version to the construction in general. For example, it could be possible that Rank has a better performance on the GOAL test set with a different set of optimisations.

Statistical Evaluation

The statistical evaluations that we did for the results of our investigations (especially for the GOAL test set) are not exhaustive. More in-depth analyses might lead to new or different conclusions. For example, for the GOAL test set, we considered only aggregated statistics, such as the mean or the median. A more detailed evaluation could include considering individual automata. For example, we could determine for which specific automata a construction produces smaller complements than another construction. In this way, the relative strengths and weaknesses of different constructions could be analysed more concretely. This would be interesting, for example, for Fribourg+M1+M2 and Fribourg+M1, or for Fribourg+R2C+C and Fribourg. Tracking individual automata could reveal what exactly makes an automaton better suited for Fribourg+M1+M2 or for Fribourg+R2C+C. Altogether, such investigations would improve our understanding of the relative merits of complementation constructions. This would in turn support the portfolio approach of selecting suitable constructions for specific automata, that we explained above.

Chapter 6

Conclusions

Appendix A

GOAL and Plugin Installation

Below we provide instructions about how to install GOAL (if not yet installed), and our plugin containing the Fribourg construction.

GOAL

GOAL can be downloaded from <http://goal.im.ntu.edu.tw/wiki/doku.php> (requires registration of name and e-mail address). After the download, GOAL does not require an actual installation. It is run directly from within the downloaded directory tree. The root directory `<root>` of this directory tree has the format `GOAL-YYYYMMDD`, where `YYYYMMDD` is the release date of the downloaded GOAL version.

There are two executables for running GOAL:

- `<root>/goal`
- `<root>/gc`

The first one, `<root>/goal`, starts the graphical user interface of GOAL. The second one, `<root>/gc` is the entry point to the command line interface of GOAL. For example, running `<root>/gc help complement` shows the help page for GOAL's complementation command.

Note for Mac OS X users: the GOAL website provides a special Mac OS X version that can be downloaded as a `.dmg` file. This file installs GOAL as an application bundle under `/Applications/GOAL.app`. The root directory `<root>` of the GOAL directory tree is in this case `GOAL.app/Contents/Resources/Java`.

Fribourg Construction Plugin

For obtaining and installing our plugin to an existing GOAL installation, execute the following steps:

1. Download the plugin from https://frico.s3.amazonaws.com/goal_plugins/ch.unifr.goal.complement.zip
2. Unzip the downloaded file `ch.unifr.goal.complement.zip`
3. Move the unzipped directory `ch.unifr.goal.complement` to `<root>/plugins`

This is all. When GOAL is started for the next time, the Fribourg construction will be included in the graphical user interface, as well as in the command line interface of GOAL.

Important note: the Fribourg construction plugin requires GOAL version 2014-11-17 or later. The plugin does not work with older versions of GOAL. This is because between the 2014-08-08 and the 2014-11-17 version, the plugin interfaces of GOAL have changed.

Note that our additional plugin `ch.unifr.goal.util`, which contains two commands for testing completeness and universality of automata (only for the command line interface), can be downloaded from https://frico.s3.amazonaws.com/goal_plugins/ch.unifr.goal.util.zip and installed according to the above instructions.

Appendix B

Matrices Corresponding to Perspective Plots

In the following two figures, we provide the matrices corresponding to all the perspective plots in Chapter 5 (that is, the following matrices contain the same data as the corresponding perspective plots, but just visualised in a different way).

The matrices in Figure B.1 correspond to the perspective plots in Figure 5.7 (a) and (b) (external tests on the GOAL test set). The eight matrices in Figure B.2 correspond to the perspective plots of Figure 5.2 and 5.3 (internal tests on the GOAL test set).

The cells of the matrices represent the median complement sizes of all the transition density/acceptance density classes of the GOAL test set. The rows of a matrix represent the transition densities, and the columns represent the acceptance densities.

Note that in Figure B.1, we do not provide the matrix for Fribourg+M1+R2C, because this is the same matrix as the one in Figure B.2 (f).

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0		0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	130	117	109	77	69	61	56	40	40	29	1.0	171	174	166	124	118	117	100	67	84	35
1.2	387	456	352	281	155	136	101	105	75	45	1.2	622	833	803	877	529	398	320	372	215	53
1.4	822	683	394	376	230	204	151	120	105	63	1.4	2,086	1,618	1,367	1,676	1,065	967	664	682	494	78
1.6	890	594	458	321	237	178	134	114	113	61	1.6	2,465	2,073	2,182	1,959	1,518	1,259	767	545	623	78
1.8	624	507	324	275	196	136	110	92	89	41	1.8	2,310	1,963	1,950	1,988	1,485	1,095	746	418	346	57
2.0	362	286	211	176	117	103	79	64	59	34	2.0	1,318	1,482	1,393	1,461	981	871	434	338	228	50
2.2	248	222	124	116	82	73	56	52	50	28	2.2	1,068	1,145	1,085	1,067	772	747	263	235	158	40
2.4	147	145	114	87	56	48	43	39	35	19	2.4	689	838	809	751	524	466	240	159	93	30
2.6	115	117	67	61	47	42	32	29	29	15	2.6	469	531	555	565	437	360	169	94	71	23
2.8	95	71	52	45	38	29	27	25	23	13	2.8	369	421	536	405	329	224	130	81	58	21
3.0	59	60	47	35	32	27	22	21	20	10	3.0	244	327	360	322	219	176	85	64	49	16

(a) Piterman+EQ+RO
(b) Slice+P+RO+MADJ+EG

Figure B.1: Median complement sizes based on the 10,998 effective samples of the external tests (without Rank) for each of the 110 transition density/acceptance density classes of the GOAL test set.

Appendix B. Matrices Corresponding to Perspective Plots

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	269	308	254	236	238	297	266	156	207	68
1.2	960	1,407	1,479	2,150	1,152	1,090	942	1,206	718	104
1.4	3,426	2,915	2,752	3,393	2,693	3,265	2,263	2,425	1,844	154
1.6	3,799	3,698	4,901	3,926	3,960	3,655	2,580	1,905	2,124	155
1.8	3,375	3,169	3,420	3,967	3,943	3,132	2,246	1,144	971	114
2.0	1,906	2,261	2,383	2,884	2,354	2,096	1,169	932	568	98
2.2	1,467	1,633	1,795	1,942	1,611	1,640	569	499	330	78
2.4	924	1,232	1,319	1,317	1,056	886	514	314	182	59
2.6	625	763	880	945	828	684	316	175	132	44
2.8	483	584	836	690	575	395	240	151	103	41
3.0	319	450	557	523	367	313	155	116	84	32

(a) Fribourg

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	390	438	434	324	328	459	337	204	227	40
1.2	1,576	2,394	2,505	2,996	1,613	1,551	1,166	1,542	1,002	58
1.4	5,007	4,336	4,652	4,877	3,458	3,956	3,169	3,380	1,868	86
1.6	5,067	5,032	6,444	4,868	4,575	3,864	3,211	1,731	1,892	85
1.8	4,016	3,701	3,647	4,523	3,548	3,009	1,808	451	336	62
2.0	1,663	2,276	2,676	3,035	1,925	1,932	464	307	150	54
2.2	989	1,514	1,621	1,826	1,121	846	155	127	93	45
2.4	560	821	919	771	529	267	133	87	55	32
2.6	388	519	524	441	259	219	84	50	41	26
2.8	311	317	396	242	165	95	64	44	33	22
3.0	173	224	211	169	102	72	41	34	27	18

(b) Fribourg+R2C

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	126	118	97	60	51	52	62	36	48	30
1.2	432	517	345	262	160	126	92	120	109	40
1.4	1,044	331	133	89	45	22	19	31	27	20
1.6	358	24	11	5	4	6	5	3	3	4
1.8	19	5	1	1	1	1	1	1	1	1
2.0	1	1	1	1	1	1	1	1	1	1
2.2	1	1	1	1	1	1	1	1	1	1
2.4	1	1	1	1	1	1	1	1	1	1
2.6	1	1	1	1	1	1	1	1	1	1
2.8	1	1	1	1	1	1	1	1	1	1
3.0	1	1	1	1	1	1	1	1	1	1

(c) Fribourg+R2C+C

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	225	223	195	181	187	199	189	124	161	68
1.2	731	971	946	1,071	629	562	488	568	388	104
1.4	2,228	1,701	1,543	1,732	1,241	1,287	945	944	727	154
1.6	2,489	2,263	2,331	2,133	1,777	1,443	964	757	889	155
1.8	2,381	2,027	2,009	2,075	1,618	1,243	1,005	592	515	114
2.0	1,390	1,569	1,416	1,573	1,093	1,008	594	464	330	98
2.2	1,118	1,197	1,150	1,151	879	809	317	330	241	78
2.4	712	885	836	809	580	535	316	231	145	59
2.6	498	569	601	627	497	412	217	137	113	44
2.8	391	455	578	456	374	263	173	119	90	41
3.0	258	350	392	354	253	208	119	97	74	32

(d) Fribourg+R

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	225	223	195	181	187	199	189	124	161	68
1.2	731	971	946	1,071	629	562	488	568	388	104
1.4	2,228	1,701	1,543	1,732	1,241	1,287	945	944	727	154
1.6	2,489	2,263	2,331	2,133	1,777	1,443	964	757	889	155
1.8	2,381	2,027	2,009	2,075	1,618	1,215	1,005	592	515	114
2.0	1,390	1,513	1,416	1,542	1,093	1,003	594	441	330	97
2.2	1,019	1,156	1,064	1,104	859	785	304	303	221	78
2.4	672	867	789	772	544	478	269	191	139	55
2.6	466	542	572	568	452	348	183	129	99	43
2.8	368	407	480	337	260	197	129	96	75	36
3.0	201	261	266	272	199	136	83	74	50	27

(e) Fribourg+M1

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	329	303	279	240	229	288	230	157	160	40
1.2	988	1,392	1,356	1,352	751	741	608	704	516	58
1.4	2,939	2,581	2,066	2,190	1,351	1,622	1,132	1,261	932	86
1.6	3,150	2,900	2,842	2,218	1,885	1,563	1,177	821	896	85
1.8	2,782	2,485	2,047	2,180	1,625	1,269	855	395	309	62
2.0	1,338	1,638	1,544	1,566	979	957	349	261	147	54
2.2	838	1,125	993	1,027	667	521	153	125	93	45
2.4	494	700	624	524	296	214	126	87	55	32
2.6	327	434	383	334	212	163	82	50	41	26
2.8	283	273	305	202	144	95	60	44	33	22
3.0	164	200	173	142	92	72	41	34	27	18

(f) Fribourg+M1+R2C

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	215	213	189	174	175	192	186	121	156	68
1.2	712	914	913	1,075	619	563	526	620	416	104
1.4	2,075	1,620	1,503	1,650	1,254	1,339	1,003	1,006	848	154
1.6	2,344	2,062	2,340	2,016	1,755	1,520	1,053	858	986	155
1.8	2,205	1,873	1,920	2,040	1,689	1,315	1,080	664	598	114
2.0	1,290	1,485	1,405	1,522	1,134	1,044	652	531	392	98
2.2	1,023	1,119	1,092	1,127	868	875	376	359	262	78
2.4	674	849	790	807	617	544	355	251	156	59
2.6	478	549	594	597	510	431	231	147	116	44
2.8	370	439	559	455	382	283	182	124	93	41
3.0	249	341	388	348	260	225	123	101	77	32

(g) Fribourg+M1+R2C+C

(h) Fribourg+M1+M2

Figure B.2: Median complement sizes based on the 10,939 effective samples of the internal tests for each of the 110 transition density/acceptance density classes of the GOAL test set.

Appendix C

Execution Times

Below we present statistics about the measured execution times of the complementation tasks on the GOAL test set. Note that the execution times for the Michel test set are presented in Chapter 5.

The reported values are in CPU time seconds. All the execution time measurements in this thesis have been done by the `time` keyword of Bash.

The second-last column of each table contains the total of the execution times of all the effective samples in the corresponding test. The last column converts this total number of seconds to hours.

Note that in general measurements of execution times are dependent on the implementation, execution environment, and on other factors, such as the way the time measurement is done by the operating system, or the current load of the processor. This is why we use the execution times only as an auxiliary performance measure.

Construction	Mean	Min.	P25	Median	P75	Max.	Total	\approx hours
Fribourg	8.5	2.5	3.3	4.9	7.3	586.0	93,351.2	259
Fribourg+R2C	6.6	2.2	2.9	4.2	6.4	219.7	72,545.7	202
Fribourg+R2C+C	8.5	2.2	2.6	3.5	6.4	582.9	93,396.2	259
Fribourg+M1	4.9	2.5	3.2	4.1	5.9	55.1	54,061.3	150
Fribourg+M1+R2C	4.4	2.2	2.8	3.6	5.3	42.5	48,572.0	135
Fribourg+M1+R2C+C	5.6	2.5	3.2	4.0	6.5	147.4	60,918.9	169
Fribourg+M1+M2	4.6	2.2	2.9	3.8	5.1	38.4	49,848.0	138
Fribourg+R	7.5	2.2	3.0	3.9	6.3	470.5	82,387.3	229

Table C.1: Execution times of the 10,939 effective samples of the internal tests on the GOAL test set. The reported values are in CPU time seconds.

Construction	Mean	Min.	P25	Median	P75	Max.	Total	\approx hours
Piterman+EQ+RO	3.0	2.2	2.6	2.8	3.0	42.9	21,410.6	59
Slice+P+RO+MADJ+EG	3.7	2.2	2.7	3.2	4.1	36.7	26,398.9	73
Rank+TR+RO	16.0	2.3	2.8	3.7	9.3	443.3	115,563.9	321
Fribourg+M1+R2C	4.0	2.2	2.7	3.1	4.4	410.4	28,970.8	80

Table C.2: Execution times of the 7,204 effective samples of the external tests (with Rank) on the GOAL test set. The reported values are in CPU time seconds.

Construction	Mean	Min.	P25	Median	P75	Max.	Total	\approx hours
Piterman+EQ+RO	3.6	2.2	2.7	2.9	3.4	365.7	39,663.4	110
Slice+P+RO+MADJ+EG	4.3	2.2	2.9	3.7	5.0	42.4	47,418.2	132
Fribourg+M1+R2C	4.7	2.2	2.8	3.6	5.3	410.4	52,149.0	145

Table C.3: Execution times of the 10,998 effective samples of the external tests (without Rank) on the GOAL test set. The reported values are in CPU time seconds.

Bibliography

- [1] J. Allred, U. Ultes-Nitsche. Complementing Büchi Automata with a Subset-Tuple Construction. Tech. rep.. University of Fribourg, Switzerland. 2014.
- [2] C. Althoff, W. Thomas, N. Wallmeier. Observations on Determinization of Büchi Automata. In J. Farré, I. Litovsky, S. Schmitz, eds., *Implementation and Application of Automata*. vol. 3845 of *Lecture Notes in Computer Science*. pp. 262–272. Springer Berlin Heidelberg. 2006.
- [3] A. Aziz, F. Balarin, S.-T. Cheng, et al. HSIS: A BDD-Based Environment for Formal Verification. In *Design Automation, 1994. 31st Conference on*. pp. 454–459. June 1994.
- [4] C. Baier, J.-P. Katoen, et al. *Principles of Model Checking*. MIT press Cambridge. 2008.
- [5] M. Ben-Ari. *Mathematical Logic for Computer Science*. Springer. 3 ed.. 2012.
- [6] S. Breuers. Optimization and experimental analysis of the algebraic complementation construction for Büchi automata. Bachelor’s thesis, RWTH Aachen. October 2010.
- [7] S. Breuers, C. Löding, J. Olschewski. Improved Ramsey-Based Büchi Complementation. In L. Birkedal, ed., *Foundations of Software Science and Computational Structures*. vol. 7213 of *Lecture Notes in Computer Science*. pp. 150–164. Springer Berlin Heidelberg. 2012.
- [8] J. R. Büchi. On a Decision Method in Restricted Second Order Arithmetic. In *Proc. International Congress on Logic, Method, and Philosophy of Science, 1960*. Stanford University Press. 1962.
- [9] A. Cimatti, E. Clarke, E. Giunchiglia, et al. NuSMV 2: An OpenSource Tool for Symbolic Model Checking. In E. Brinksma, K. Larsen, eds., *Computer Aided Verification*. vol. 2404 of *Lecture Notes in Computer Science*. pp. 359–364. Springer Berlin Heidelberg. 2002.
- [10] A. Cimatti, E. Clarke, F. Giunchiglia, et al. NuSMV: A New Symbolic Model Verifier. In N. Halbwachs, D. Peled, eds., *Computer Aided Verification*. vol. 1633 of *Lecture Notes in Computer Science*. pp. 495–499. Springer Berlin Heidelberg. 1999.
- [11] S. J. Fogarty, O. Kupferman, T. Wilke, et al. Unifying Büchi Complementation Constructions. *Logical Methods in Computer Science*. 9(1). 2013.
- [12] E. Friedgut, O. Kupferman, M. Vardi. Büchi Complementation Made Tighter. In F. Wang, ed., *Automated Technology for Verification and Analysis*. vol. 3299 of *Lecture Notes in Computer Science*. pp. 64–78. Springer Berlin Heidelberg. 2004.
- [13] E. Friedgut, O. Kupferman, M. Y. Vardi. Büchi Complementation Made Tighter. *International Journal of Foundations of Computer Science*. 17(04):pp. 851–867. 2006.
- [14] C. Göttel. Implementation of an Algorithm for Büchi Complementation. Bachelor’s thesis, University of Fribourg, Switzerland. November 2013.
- [15] R. L. Graham, B. L. Rothschild, J. H. Spencer. *Ramsey theory*. vol. 20. John Wiley & Sons. 1990.
- [16] S. Gurumurthy, O. Kupferman, F. Somenzi, et al. On Complementing Nondeterministic Büchi Automata. In D. Geist, E. Tronci, eds., *Correct Hardware Design and Verification Methods*. vol. 2860 of *Lecture Notes in Computer Science*. pp. 96–110. Springer Berlin Heidelberg. 2003.
- [17] G. J. Holzmann. The Model Checker SPIN. *IEEE Transactions on Software Engineering*. 23(5):pp. 279–295. 1997.

- [18] J. E. Hopcroft, R. Motwani, J. D. Ullman. *Introduction to Automata Theory, Languages, and Computation*. Addison-Wesley. 2nd edition ed.. 2001.
- [19] M. Huth, M. Ryan. *Logic in Computer Science: Modelling and Reasoning about Systems*. Cambridge University Press. 2004.
- [20] D. Kähler, T. Wilke. Complementation, Disambiguation, and Determinization of Büchi Automata Unified. In L. Aceto, I. Damgård, L. Goldberg, et al, eds., *Automata, Languages and Programming*. vol. 5125 of *Lecture Notes in Computer Science*. pp. 724–735. Springer Berlin Heidelberg. 2008.
- [21] H. Karmarkar, S. Chakraborty. On Minimal Odd Rankings for Büchi Complementation. In Z. Liu, A. Ravn, eds., *Automated Technology for Verification and Analysis*. vol. 5799 of *Lecture Notes in Computer Science*. pp. 228–243. Springer Berlin Heidelberg. 2009.
- [22] N. Klarlund. Progress measures for complementation of omega-automata with applications to temporal logic. In *Foundations of Computer Science, 1991. Proceedings., 32nd Annual Symposium on*. pp. 358–367. Oct 1991.
- [23] J. Klein. Linear Time Logic and Deterministic Omega-Automata. *Master's thesis, Universität Bonn*. 2005.
- [24] O. Kupferman, M. Y. Vardi. Weak Alternating Automata Are Not that Weak. In *Proceedings of the 5th Israeli Symposium on Theory of Computing and Systems*. pp. 147–158. IEEE Computer Society Press. 1997.
- [25] O. Kupferman, M. Y. Vardi. Weak Alternating Automata Are Not that Weak. *ACM Trans. Comput. Logic*. 2(3):pp. 408–429. Jul. 2001.
- [26] R. Kurshan. Complementing Deterministic Büchi Automata in Polynomial Time. *Journal of Computer and System Sciences*. 35(1):pp. 59 – 71. 1987.
- [27] K. Leyton-Brown, E. Nudelman, G. Andrew, et al. A portfolio approach to algorithm selection. In *IJCAI*. vol. 1543. p. 2003. 2003.
- [28] C. Löding. *Methods for the Transformation of ω -Automata: Complexity and Connection to Second Order Logic*. Master's thesis. Christian-Albrechts-University of Kiel. 1998.
- [29] C. Löding. Optimal Bounds for Transformations of ω -Automata. In C. Rangan, V. Raman, R. Ramanujam, eds., *Foundations of Software Technology and Theoretical Computer Science*. vol. 1738 of *Lecture Notes in Computer Science*. pp. 97–109. Springer Berlin Heidelberg. 1999.
- [30] M. Michel. Complementation is more difficult with automata on infinite words. *CNET, Paris*. 15. 1988.
- [31] M. Mohri. A Disambiguation Algorithm for Finite Automata and Functional Transducers. In N. Moreira, R. Reis, eds., *Implementation and Application of Automata*. vol. 7381 of *Lecture Notes in Computer Science*. pp. 265–277. Springer Berlin Heidelberg. 2012.
- [32] A. Mostowski. Regular expressions for infinite trees and a standard form of automata. In A. Skowron, ed., *Computation Theory*. vol. 208 of *Lecture Notes in Computer Science*. pp. 157–168. Springer Berlin Heidelberg. 1985.
- [33] D. E. Muller. Infinite Sequences and Finite Machines. In *Switching Circuit Theory and Logical Design, Proceedings of the Fourth Annual Symposium on*. pp. 3–16. Oct 1963.
- [34] D. E. Muller, A. Saoudi, P. E. Schupp. Alternating automata, the weak monadic theory of the tree, and its complexity. In L. Kott, ed., *Automata, Languages and Programming*. vol. 226 of *Lecture Notes in Computer Science*. pp. 275–283. Springer Berlin Heidelberg. 1986.
- [35] D. E. Muller, P. E. Schupp. Simulating Alternating Tree Automata by Nondeterministic Automata: New Results and New Proofs of the Theorems of Rabin, McNaughton and Safra. *Theoretical Computer Science*. 141(1–2):pp. 69 – 107. 1995.
- [36] F. Nießner, U. Nitsche, P. Ochsenschläger. Deterministic Omega-Regular Liveness Properties. In S. Bozapalidis, ed., *Preproceedings of the 3rd International Conference on Developments in Language Theory, DLT'97*. pp. 237–247. Citeseer. 1997.

- [37] N. Piterman. From Nondeterministic Buchi and Streett Automata to Deterministic Parity Automata. In *Logic in Computer Science, 2006 21st Annual IEEE Symposium on*. pp. 255–264. 2006.
- [38] N. Piterman. From Nondeterministic Buchi and Streett Automata to Deterministic Parity Automata. *Logical Methods in Computer Science*. 3(5):pp. 1–21. 2007.
- [39] M. Rabin, D. Scott. Finite Automata and Their Decision Problems. *IBM Journal of Research and Development*. 3(2):pp. 114–125. April 1959.
- [40] M. O. Rabin. Decidability of second-order theories and automata on infinite trees. *Transactions of the American Mathematical Society*. 141:pp. 1–35. July 1969.
- [41] F. P. Ramsey. On a Problem of Formal Logic. *Proceedings of the London Mathematical Society*. s2-30(1):pp. 264–286. 1930.
- [42] M. Roggenbach. Determinization of Büchi-Automata. In E. Grädel, W. Thomas, T. Wilke, eds., *Automata Logics, and Infinite Games*. vol. 2500 of *Lecture Notes in Computer Science*. pp. 43–60. Springer Berlin Heidelberg. 2002.
- [43] S. Safra. On the Complexity of Omega-Automata. *Journal of Computer and System Science*. 1988.
- [44] S. Safra. On the Complexity of Omega-Automata. In *Foundations of Computer Science, 1988., 29th Annual Symposium on*. pp. 319–327. Oct 1988.
- [45] S. Schewe. Büchi Complementation Made Tight. In *26th International Symposium on Theoretical Aspects of Computer Science-STACS 2009*. pp. 661–672. 2009.
- [46] A. Sistla, M. Vardi, P. Wolper. The complementation problem for Büchi automata with applications to temporal logic. In W. Brauer, ed., *Automata, Languages and Programming*. vol. 194 of *Lecture Notes in Computer Science*. pp. 465–474. Springer Berlin Heidelberg. 1985.
- [47] A. P. Sistla, M. Y. Vardi, P. Wolper. The Complementation Problem for Büchi Automata with Applications to Temporal Logic. *Theoretical Computer Science*. 49(2–3):pp. 217 – 237. 1987.
- [48] F. Somenzi, R. Bloem. Efficient Büchi automata from LTL formulae. In *Computer Aided Verification*. pp. 248–263. Springer. 2000.
- [49] R. S. Streett. Propositional dynamic logic of looping and converse is elementarily decidable. *Information and Control*. 54(1–2):pp. 121 – 141. 1982.
- [50] S. Taşiran, R. Hojati, R. Brayton. Language Containment of Non-Deterministic ω -Automata. In P. Camurati, H. Ekeing, eds., *Correct Hardware Design and Verification Methods*. vol. 987 of *Lecture Notes in Computer Science*. pp. 261–277. Springer Berlin Heidelberg. 1995.
- [51] W. Thomas. Automata on Infinite Objects. In J. van Leeuwen, ed., *Handbook of Theoretical Computer Science (Vol. B)*. chap. Automata on Infinite Objects, pp. 133–191. MIT Press, Cambridge, MA, USA. 1990.
- [52] W. Thomas. Languages, Automata, and Logic. In G. Rozenberg, A. Salomaa, eds., *Handbook of Formal Languages*. pp. 389–455. Springer Berlin Heidelberg. 1997.
- [53] W. Thomas. Complementation of Büchi Automata Revisited. In J. Karhumäki, H. Maurer, G. Păun, et al, eds., *Jewels are Forever*. pp. 109–120. Springer Berlin Heidelberg. 1999.
- [54] M.-H. Tsai, S. Fogarty, M. Vardi, et al. State of Büchi Complementation. In M. Domaratzki, K. Salomaa, eds., *Implementation and Application of Automata*. vol. 6482 of *Lecture Notes in Computer Science*. pp. 261–271. Springer Berlin Heidelberg. 2011.
- [55] M.-H. Tsai, Y.-K. Tsay, Y.-S. Hwang. GOAL for Games, Omega-Automata, and Logics. In N. Sharygina, H. Veith, eds., *Computer Aided Verification*. vol. 8044 of *Lecture Notes in Computer Science*. pp. 883–889. Springer Berlin Heidelberg. 2013.
- [56] Y.-K. Tsay, Y.-F. Chen, M.-H. Tsai, et al. Goal: A Graphical Tool for Manipulating Büchi Automata and Temporal Formulae. In O. Grumberg, M. Huth, eds., *Tools and Algorithms for the Construction and Analysis of Systems*. vol. 4424 of *Lecture Notes in Computer Science*. pp. 466–471. Springer Berlin Heidelberg. 2007.

- [57] Y.-K. Tsay, Y.-F. Chen, M.-H. Tsai, et al. Goal Extended: Towards a Research Tool for Omega Automata and Temporal Logic. In C. Ramakrishnan, J. Rehof, eds., *Tools and Algorithms for the Construction and Analysis of Systems*. vol. 4963 of *Lecture Notes in Computer Science*. pp. 346–350. Springer Berlin Heidelberg. 2008.
- [58] Y.-K. Tsay, Y.-F. Chen, M.-H. Tsai, et al. Tool support for learning Büchi automata and linear temporal logic. *Formal Aspects of Computing*. 21(3):pp. 259–275. 2009.
- [59] Y.-K. Tsay, M.-H. Tsai, J.-S. Chang, et al. Büchi Store: An Open Repository of Büchi Automata. In P. Abdulla, K. Leino, eds., *Tools and Algorithms for the Construction and Analysis of Systems*. vol. 6605 of *Lecture Notes in Computer Science*. pp. 262–266. Springer Berlin Heidelberg. 2011.
- [60] U. Ultes-Nitsche. A Power-Set Construction for Reducing Büchi Automata to Non-Determinism Degree Two. *Information Processing Letters*. 101(3):pp. 107 – 111. 2007.
- [61] M. Vardi. An automata-theoretic approach to linear temporal logic. In F. Moller, G. Birtwistle, eds., *Logics for Concurrency*. vol. 1043 of *Lecture Notes in Computer Science*. pp. 238–266. Springer Berlin Heidelberg. 1996.
- [62] M. Vardi. Automata-Theoretic Model Checking Revisited. In B. Cook, A. Podelski, eds., *Verification, Model Checking, and Abstract Interpretation*. vol. 4349 of *Lecture Notes in Computer Science*. pp. 137–150. Springer Berlin Heidelberg. 2007.
- [63] M. Vardi. The Büchi Complementation Saga. In W. Thomas, P. Weil, eds., *STACS 2007*. vol. 4393 of *Lecture Notes in Computer Science*. pp. 12–22. Springer Berlin Heidelberg. 2007.
- [64] M. Y. Vardi. Buchi Complementation: A Forty-Year Saga. In *5th symposium on Atomic Level Characterizations (ALC'05)*. 2005.
- [65] M. Y. Vardi, T. Wilke. Automata: From Logics to Algorithms. In J. Flum, E. Grädel, T. Wilke, eds., *Logic and Automata: History and Perspectives*. vol. 2 of *Texts in Logic and Games*. pp. 629–736. Amsterdam University Press. 2007.
- [66] T. Wilke. ω -Automata. In J.-E. Pin, ed., *Handbook of Automata Theory*. European Mathematical Society. To appear, 2015.
- [67] Q. Yan. Lower Bounds for Complementation of ω -Automata Via the Full Automata Technique. In M. Bugliesi, B. Preneel, V. Sassone, et al, eds., *Automata, Languages and Programming*. vol. 4052 of *Lecture Notes in Computer Science*. pp. 589–600. Springer Berlin Heidelberg. 2006.
- [68] Q. Yan. Lower Bounds for Complementation of ω -Automata Via the Full Automata Technique. *CoRR*. abs/0802.1226. 2008.