

Empirical Performance Investigation of a Büchi Complementation Construction

Daniel Weibel

June 22, 2015

Abstract

This will be the abstract.

Acknowledgements

First of all, I would like to thank my supervisors Prof. Dr. Ulrich Ultes-Nitsche and Joel Allred to bring me to the topic of this thesis in the first place, and for the many enlightening discussions and explanations. A very special thank goes to Ming-Hsien Tsai, the main author of the GOAL tool, from the National Taiwan University. Without his very prompt and patient responses to my many questions regarding implementation details of GOAL, the plugin would not have been possible in its current form. I also would like to thank Michael Rolli and Nico Färber from the Grid Admin Team of the University of Bern for their very friendly and helpful responses to my questions about the Linux cluster where the experiments of this thesis have been executed.

Contents

1	Performance Investigation of the Fribourg Construction	2
1.1	Implementation	3
1.1.1	GOAL	3
1.1.2	Implementation of the Construction	4
1.1.3	Verification of the Implementation	6
1.2	Test Data	7
1.2.1	GOAL Test Set	8
1.2.2	Michel Test Set	11
1.3	Experimental Setup	12
1.3.1	Constructions for the Internal Tests	12
1.3.2	Constructions for the External Tests	14
1.3.3	Execution Environment	14
1.3.4	Time and Memory Limits	15
2	Results and Discussion	17
2.1	Internal Tests	17
2.1.1	GOAL Test Set	17
2.1.2	Michel Test Set	24
2.2	External Tests	26
2.2.1	GOAL Test Set	26
2.2.2	Michel Automata	28
2.3	Summary and Discussion of the Results	30
2.4	Limitations of the Approach	30
A	Plugin Installation and Usage	31
B	Median Complement Sizes of the GOAL Test Set	32
C	Execution Times	34

Chapter 1

Performance Investigation of the Fribourg Construction

Contents

1.1 Implementation	3
1.1.1 GOAL	3
1.1.2 Implementation of the Construction	4
1.1.3 Verification of the Implementation	6
1.2 Test Data	7
1.2.1 GOAL Test Set	8
1.2.2 Michel Test Set	11
1.3 Experimental Setup	12
1.3.1 Constructions for the Internal Tests	12
1.3.2 Constructions for the External Tests	14
1.3.3 Execution Environment	14
1.3.4 Time and Memory Limits	15

In this chapter we come to the core of this thesis, namely the empirical performance investigation of the Fribourg construction. We are interested in two things. First, how the different versions of the Fribourg construction compare to each other. That is, how do different combinations of optimisations influence the performance of the construction. Second, we want to know how the Fribourg construction performs compared to existing complementation constructions. Our main measure for the performance of a construction is the number of states of the produced complement. Throughout this thesis, we will refer to the first question as the *internal* tests, and to the second question as the *external* tests.

To do an empirical performance investigation we need an implementation of the Fribourg construction. We decided to create this implementation in the framework of an existing tool called GOAL. This is a Java tool with a graphical user interface for manipulating ω -automata, and it contains implementations of various Büchi complementation constructions. In this way we can easily compare the Fribourg construction to these other construction (see external tests).

The next thing we need for an empirical performance investigation is test data. These are specific sets of automata on which all the tested construction are run. We defined two test sets. The first one, called the GOAL test set, contains a large number of randomly generated automata. The second one, called the Michel test set, contains just a small number of automata that have a special property.

Having an implementation and test data, the experiments need to be executed. Our chosen implementation approach and test data results in heavy computation tasks, that require a lot of computation power

and time. We therefore decided to execute the experiments in a professional high-performance computing (HPC) environment. This environment is the Linux-based HPC computing cluster, called UBELIX, at the University of Bern.¹

In this chapter, we describe each of these points in a separate section. Section 1.3 also includes our experimental setup, that is, a listing of the concrete construction versions that we tested, the allocated computing resources, and so on. The results of the experiments will finally be presented in Chapter 2.

1.1 Implementation

As mentioned, we implemented the Fribourg construction as part of the GOAL tool. This is possible thanks to the extensible plugin architecture of GOAL which allows to write plugins that contain additional functionality for GOAL. Our implementation of the Fribourg construction has therefore the form of a GOAL-plugin.

In this section, we first present the GOAL tool in a general way (Section 1.1.1). In Section 1.1.2, we give some more details about the plugin architecture of GOAL, and describe some properties of our implementation. Finally, in Section 1.1.3, we describe how we verified the correctness of our implementation.

1.1.1 GOAL

GOAL stands for *Graphical Tool for Omega-Automata and Logics* and is being developed since 2007 by the Department of Information Management at the National Taiwan University (NTU)². The tool has been presented in various scientific publications [15][16][17][14]. The executables of GOAL can be freely downloaded on <http://goal.im.ntu.edu.tw>.

GOAL allows to graphically and interactively create and manipulate different types of ω -automata, including non-deterministic Büchi automata (NBW). The palette of provided manipulations is vast and ranges from input testing, over conversions to other types of non-deterministic ω -automata, to implementations of graph layout algorithms. Figure 1.1 shows a screenshot of GOAL's graphical user interface with an open menu showing the breadth of possible manipulations for ω -automata.

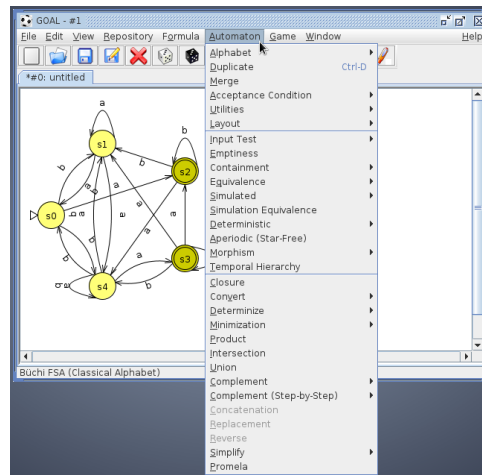


Figure 1.1: The graphical user interface of GOAL (version 2014-11-17). The open menu item gives an idea about the different types of manipulations that GOAL provides for ω -automata.

Relevant for our purposes are the complementation constructions for NBW that GOAL provides. GOAL comes up with a comprehensive offer in this regard too. In the 2014-11-17 version of GOAL, there are 10 pre-implemented complementation constructions for NBW. Table 1.1 summarises these constructions and indicates the authors and the publications on which the implementations are based.

¹ <http://ubelix.unibe.ch>

² <http://exp.management.ntu.edu.tw/en/IM>

#	Identifier	Name/description	Authors (year)	Ref.
1	Ramsey	Ramsey-based construction	Sistla, Vardi, Wolper (1987)	[10]
2	Safra	Safra construction	Safra (1988)	[8]
3	ModifiedSafra	Modification by Althoff	Althoff (2006)	[1]
4	Piterman	Safra-Piterman construction	Piterman (2007)	[7]
5	MS	Muller-Schupp construction	Muller, Schupp (1995)	[6]
6	Rank	Rank-based construction	Schewe (2009)	[9]
7	WAPA	Via weak alternating parity automata	Thomas (1999)	[12]
8	WAA	Via weak alternating automata	Kupferman, Vardi (2001)	[4]
9	Slice+P	Slice-based construction (earlier)	Vardi, Wilke (2007)	[18]
10	Slice	Slice-based construction (later)	Kähler, Wilke (2008)	[3]

Table 1.1: The pre-implemented NBW complementation constructions in GOAL (version 2014-11-17).

The construction in Table 1.1 are sorted by the four fundamental complementation approaches. The first one, Ramsey, is the only construction belonging to the Ramsey-based complementation approach. The following four constructions, Safra, ModifiedSafra, Piterman, and MS, belong to the determinization-based approach. Rank, WAPA, and WAA are representants of the Rank-based approach. Finally, Slice and Slice+P belong to the slice-based approach.

Slice and Slice+P are combined in a single construction in GOAL but the two specific constructions can be selected by the means of an option. If the P option of Slice is activated, then the Slice+P construction is used, and otherwise the Slice construction is used. In our experiments, we will use the Slice+P version, however, we will still refer to this construction as simply Slice in its short form.

Almost the entire functionality of GOAL is also available via a command line interface. This makes it suitable for automatic batch operation, as it is needed for our experiments. For storing automata, GOAL defines its own file format called GFF (GOAL file format) which is based on XML and is typically used with the filename extension `.gff`.

At this point, it is worth pointing at a related project of the same research group, called the Büchi Store. This is an online repository of classified and tagged ω -automata that can be downloaded in different formats (including GFF). The Büchi Store is located on <http://buchi.im.ntu.edu.tw/> and has also been described in a scientific publication [?]. Furthermore, there is a binding in GOAL to the Büchi Store, so that the contents of the store can be directly accessed from GOAL. For our project we did not make use the Büchi Store, but it is might be an interesting option for related projects.

1.1.2 Implementation of the Construction

The GOAL Plugin

GOAL has been designed from the ground up to be modular and extensible. To this end, it has been created with the Java Plugin Framework (JPF)³. This framework allows to build applications whose functionality can be easily and seamlessly extended by writing additional plugins for it. These plugins can be installed in the main application without the need to recompile the whole application. Rather, the plugin is compiled separately and the resulting bytecode files are copied to the directory tree of the main application. It is not even necessary to know the source code of the main application in order to write a plugin. The interfaces of JPF itself, and the documentations of the relevant classes of the main application are all that a plugin writer needs to know.

In some more detail, JPF requires an application to define so called *extension points*. For any extension point, multiple *extensions* can be provided. These extensions contain the actual functionality of the application. A JPF application basically consists of extensions that are plugged into their corresponding

³<http://jpf.sourceforge.net/>

Option	Description
R2C	Apply R2C optimisation if input automaton is complete
M1	Apply M1 optimisation (component mergin)
M2	Apply M2 optimisation (colour 2 reduction)
C	Make input automaton complete before start of construction
R	Remove unreachable and dead states from output automaton
RR	Remove unreachable and dead states from input automaton
MAcc	Maximise accepting states of input automaton
B	Use the “bracket notation” for state labels

Table 1.2: The options of the Fribourg construction in GOAL.

extension points. A plugin is an arbitrary bundle of extensions and extension points. It is the basic unit of organisation in the Java Plugin Framework.

One of the extension points of GOAL is called **ComplementConstruction**. The extensions to **ComplementConstruction** contain the actual complementation constructions that GOAL provides. For adding a new complementation construction to GOAL, one has thus to ceate a new extension to **ComplementConstruction**. This extension can then be wrapped in a plugin, and the plugin can be compiled and installed in the main application, what makes it an integral part of it. This means that once the plugin is installed, the new construction is included in GOAL in the same way as all the other constructions.

This is how we added the Fribourg construction to GOAL. The name of our plugin is `ch.unifr.goal.complement`⁴. It is publicly available and can be installed by anybody in their GOAL application. We give instructions on how to get, install, and use the plugin in Appendix A.

In reality, there is more than just the extension point **ComplementConstruction** that can be extended to add a new complementation construction to GOAL. There are separate extension points for, for example, the command line binding, menu inclusion, or step-by-step execution support. We created extensions to all these extension points as well and included them in our plugin. Our aim was to make the integration of the Fribourg construction in GOAL as complete as possible so that it provides the same facilities as the pre-implemented constructions.

The Fribourg Construction Options

In our implementation of the Fribourg construction we also included the three optimisations, R2C, M1, and M2, described in Section ???. We implemented these optimisation as user-selectable complementation construction options. In the GUI, these options are presented to the user as a list of checkboxes immediately before the start of each complementation task. In the command line mode, there is a command line flag for each option that can be set or not set by the user.

In addition to the three optimisations, we added further options to our construction. Table 1.2 lists all the available options for the Fribourg construction. Each option has an identifier consisting of upper-case letters that we will use troughout the rest of this thesis to refer to teh corresponding options.

The first three options in Table 1.2 represent the three optimisations from Section ???. The R2C optmisation is implemented so that it applies only to input automata that are complete. That is, selecting R2C for the complementation of an automaton that is not complete has no effect, and the result is the same as if R2C would not have been selected. The options M1 and M2 implement the M1 and M2 optimisations. Since M2 is dependent on M1, it is not possible to select M2 without also selecting M1. This restriction is enforced in both the GUI and the command line interface.

The C option is one of the options that modifies the input automaton before the actual complementation starts. This option first checks if the input automaton is complete⁵, and if this is not the case, makes it complete by adding a sink state. This means that an additional non-accepting state, the sink state, is

⁴By convention, JPF plugins are named after the base package name of their implementation files.

⁵An automaton is complete if every state has at least one outgoing transition for every symbol of the alphabet.

added to the automaton, and from every incomplete state the “missing” transitions are added from this state to the sink state. The sink state itself has loop transitions for all symbols of the alphabet.

The purpose of the C option is to be used in conjunction with the R2C optimisation. By making an automaton complete before the start of the construction, we can ensure that the R2C optimisation will be applied. The question then arises whether, in terms of performance, it is worth to do is. Because for making an automaton complete, we have to add an additional state to the automaton what generally increases the complexity of the complementation. This question has been investigated in previous work about the Fribourg construction by Göttel [2]. In this thesis will re-investigate this point in an extended form.

The R option modifies the output automaton at the end of the construction. In particular, it removes all the so called unreachable and dead states from the complement. Unreachable states are states that cannot be reached from the initial state. Dead states are states from which it is not possible to reach an accepting state. These states can be removed from any automaton without changing the language of the automaton. The pre-implemented complementation constructions Ramsey, Piterman, Rank, and Slice also contain a similar R option.

One usage case of the R option is to determine the number of unreachable and dead state that a complementation construction produces. Complementing the same automaton with and without the R option, and taking the difference of the complement sizes will yield this number. Investigations in this direction have been done with GOAL by Tsai et al. [13]. In our own investigations we will use the R option to determine the number of unreachable and dead states the plain Fribourg construction produces.

The RR option is similar to the R option, except that it removes the unreachable and dead states from the input automaton rather than from the output automaton. This option is a custom creation by us, and the pre-implemented complementation constructions in GOAL do not contain a similar option. We are not using the RR option in our investigations.

The MACC option again modifies the input automaton before the start of the construction. Namely, it applies the technique of the “maximisation of the accepting set”. This means that as many states as possible are made accepting without changing the language of the automaton. The larger number of accepting state should then simplify the complementation task. This technique has been introduced and empirically investigated by Tsai et al. in [13]. The pre-implemented complementation constructions Ramsey, Piterman, Rank, and Slice contain a similar MACC option. In our own investigations we will however not use the MACC option.

Finally, the B option is a pure display option and does not alter the automata or the construction itself. Its effect is to use the bracket notation for state labels, that indicates component colours by different types of brackets, instead of the default notation that uses numbers to specify the component colours.

1.1.3 Verification of the Implementation

Having an implementation, it is important to verify that it is correct. With correct we mean in this section that the construction produces a valid complement for a given input automaton, that is, that the language of the output automaton is indeed the complement of the language of the input automaton. We verified this point of our implementation as described below. In this sense, we know that our implementation is a valid construction. This is not the same as the verification that our implementation faithfully represents the Fribourg construction as it has been devised by its creators. Rather, this latter point has been informally verified during the whole development period of the implementation, which was possible thanks to the close collaboration with the construction creators.

In order to verify that our implementation produces correct results, we performed so called complementation-equivalence tests in GOAL against one of the pre-implemented construction. This works as follows. Take an automaton A and complement it with one of the GOAL constructions. This yields automaton A' . Then, complement the same automaton A with our implementation of the Fribourg construction. This yields automaton A'' . Now, by the means of GOAL's equivalence operation then test whether A' and A'' accept the same language.

This approach makes a couple of assumptions. First, it relies on the correctness of the pre-implemented complementation constructions, and on the equivalence operation of GOAL (which is also based on complementation). Second, with this empirical approach it is of course not possible to conclusively verify the correctness of our implementation. Every passed complementation-equivalence test just further confirms the hypothesis that our implementation is correct, but it can never be proved. Despite these points, we think that this approach is the best we can do, and that, with a sufficient number of test cases, it can confirm the correctness of our implementation with a high probability.

We tested the Fribourg construction in different versions with all the options that we described in the last section (except the B option as it does not influence the construction). The tested versions are the following.

- Fribourg
- Fribourg+R2C+C
- Fribourg+M1
- Fribourg+M1+M2
- Fribourg+R
- Fribourg+RR
- Fribourg+MACC

We tested each version with 1,000 randomly generated automata. The automata had a size of 4 and an alphabet size between 2 and 4. The pre-implemented construction we tested against, was the Piterman construction. The computations were executed on the UBELIX computer cluster that is described in Section 1.3.3. The result was that all the tests were successful, that is, there was not a single counterexample.

With a size of 4, the automata used for the tests are rather small, and it would be interesting to do the tests with bigger automata. However, our limited computing and time resources prevented us from doing so. The equivalence test of GOAL is implemented as reciprocal containment tests, which include complementation. That is, the overall test includes the complementation of the complements of the test automata. By using bigger test automata, their complements might be already so big that a further complementation is practically infeasible with our available resources. Nevertheless, we think that the high number of performed tests confirms the correctness of our construction with a high probability.

1.2 Test Data

Test data is the sample data that is given to an algorithm as input in order to measure and evaluate certain properties of the output. The test data is an important part of every empirical study and should meet several requirements. It should include enough test cases so that the results are statistically significant. It should not be biased in favour of the tested algorithm. It should cover test scenarios that are relevant for evaluating the performance of the algorithm. Ideally, publicly available test data is used that has also been used for other empirical studies. In this way, the data is “objective” and results of different studies are comparable.

In our case the test data consists of a set of non-deterministic Büchi automata which are provided as input to the Fribourg construction. We chose two different sets of automata of which we believe that together they form a meaningful test set for our study.

The first set of test automata, which we call GOAL test set, consists of 11,000 NBW of size 15. This test set has been used by a previous empirical study with GOAL by Tsai et al. [?], and is publicly available. The second set of automata, which we call Michel test set, consists of a family of automata that show a very pronounced state growth for complementation. These automata have been introduced and described by Michel [5], hence the name. Below, we describe these two test sets with their relevant properties in separate sections.

1.2.1 GOAL Test Set

The GOAL test set is the larger and more complex one of the two test sets. In this section, we first introduce the GOAL test set and describe its basic structure. In the second part, we present the results of an analysis that we did in order to reveal further properties of the GOAL test set, namely the number and distribution of complete, universal, and empty automata.

Introduction and Basic Structure

The GOAL test set has been created by Tsai et al. for an empirical study evaluation the effects of several optimisations on existing Büchi complementation constructions [?]. This study has been executed in GOAL and the automata in the test set are available in the GOAL file format, hence the name GOAL test set.

The entire test set consists of 11,000 automata of size 15, and 11,000 similar automata of size 20. For our study we used however only the automata of size 15. Hence, when we refer to the GOAL test set in the rest of this thesis, we specifically mean the 11,000 automata of size 15. The GOAL test set is publicly available through the following link: <https://fribourg.s3.amazonaws.com/testset/goal.zip>⁶.

All the automata in the test set have 15 states and an alphabet size of 2. The further properties of the automata, namely accepting states and transitions, are determined by the values of two parameters called *transition density* and *acceptance density*

The transition density determines the number of transitions of an automaton. In some more detail, the transition density t is defined as follows. Let n be the number of states of automaton A , and t its transition density. Then A contains tn transitions for each symbol of the alphabet. In the case that tn is not an integer, it is rounded up to the next integer. That is, if one of our automata with 15 states and the alphabet $0, 1$ has a transition density of 2, then it contains exactly 30 transitions for symbol 0 and 30 transitions for symbol 1.

The acceptance density a is defined as the ratio of accepting states to non-accepting states in the automaton. It is thus a number between 0 and 1. If automaton A has n states and an acceptance density of a , then it has an accepting states. In the case that an is not an integer, it is rounded up to the next integer.


The GOAL test set is structured into 110 classes with different transition density/acceptance density pairs. These 110 pairs result from the cartesian product of 11 transition densities and 10 acceptance densities. The concrete transition densities t' and acceptance densities a' are the following:

$$\begin{aligned} t' &= (1.0, 1.2, 1.4, 1.6, 1.8, 2.0, 2.2, 2.4, 2.6, 2.8, 3.0) \\ a' &= (0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0) \end{aligned}$$


Thus, there is one class whose automata have a transition density of 1.0 and an acceptance density of 0.1, another class with a transition density of 1.0 and an acceptance density of 0.2, and so on. Each of the 110 classes contains 100 automata

When Tsai et al. created the GOAL test set, they created the automata within the given constraints of any class at random. They chose this structure and parameter values in order to generate a broad range of complementation problems ranging from easy to hard [?].

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	0	0	0	0	0	0	0	0	0	0
1.2	0	0	0	0	0	0	0	0	0	0
1.4	0	0	0	0	0	0	0	0	0	0
1.6	0	0	0	0	0	0	0	0	0	0
1.8	1	1	0	0	0	1	1	1	0	0
2.0	0	5	1	2	2	3	1	2	2	2
2.2	5	10	8	5	3	5	8	6	7	1
2.4	10	6	11	11	8	6	10	20	9	7
2.6	17	17	12	16	14	19	22	21	19	19
2.8	27	20	29	32	26	27	30	25	24	19
3.0	37	37	40	39	34	37	38	35	38	39

(a) Number of *complete* automata per class.

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	4	5	5	7	8	4	6	10	4	3
1.2	1	3	5	8	8	12	10	13	4	14
1.4	2	17	13	17	20	24	22	21	27	26
1.6	16	28	30	37	49	42	42	49	45	45
1.8	31	40	55	59	64	67	76	70	63	78
2.0	60	64	85	75	83	83	79	90	87	83
2.2	67	87	86	88	89	91	89	89	89	86
2.4	88	89	86	92	95	95	94	97	96	97
2.6	86	93	92	97	97	97	98	96	98	96
2.8	94	97	95	94	97	99	98	97	97	100
3.0	99	99	99	97	99	98	100	100	100	99

(b) Number of *universal* automata per class

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	17	7	4	5	2	4	3	1	1	0
1.2	4	2	1	1	0	1	0	0	0	0
1.4	2	1	0	0	0	0	0	0	1	2
1.6	0	0	0	0	0	0	1	0	0	0
1.8	1	0	0	0	1	0	0	0	0	0
2.0	0	0	0	0	0	0	0	0	0	0
2.2	0	0	0	0	0	0	0	0	0	0
2.4	0	0	0	0	0	0	0	0	0	0
2.6	0	0	0	0	0	0	0	0	0	0
2.8	0	0	0	0	0	0	0	0	0	0
3.0	0	0	0	0	0	0	0	0	0	0


(c) Number of *empty* automata per class

Figure 1.2: Number of complete, universal, and empty automata for each of the 110 classes of transition density/acceptance density combinations. Each class contains 100 automata.

Further Properties: Completeness, Universality, and Emptiness

We analysed the properties of completeness, universality, and emptiness⁷ of the automata in the GOAL test set. To know about these properties is useful for the interpretation of the results of our study. For

⁶This link is maintained by the author of the thesis. The original link of the GOAL tests set that is maintained by the authors of [?] is http://goal.im.ntu.edu.tw/wiki/lib/exe/fetch.php?media=goal:c1aa2010_automata.tar.gz. This package

example, the R2C optimisation applies only to complete automata, thus it is interesting to know how many automata are complete. The smallest possible complements of universal and empty automata have a size of 1. Thus, we can see how many “superfluous” states a construction produces when complementing a universal or empty automaton.

GOAL provides a command for testing emptiness. However, it does not provide commands for testing completeness and universality. We therefore implemented these commands on our own and bundled them as a separate GOAL plugin. The plugin is called `ch.unifr.goal.util` and also publicly available as described in Appendix A.

With these GOAL commands we tested each of the 11,000 automata for the three properties. On one hand, we want to know how many complete, universal, and empty automata are there in total. On the other hand, we also want to know how these properties are distributed across the 110 classes of transition density/acceptance density combinations. Thus, we also determined the number of complete, universal, and empty automata for each class. The computations were executed on the UBELIX computer cluster that is described in Section 1.3.3.

The resulting overall numbers are the following:

- 990 of the 11,000 automata are complete (9%)
- 6,796 of the 11,000 automata are universal (61.8%)
- 63 of the 11,000 automata are empty (0.6%)

The fact that only 9% of the automata are complete means that the R2C optimisation of the Fribourg construction affects only 9% of the test data. This is an interesting fact for the analysis of the effect the R2C optimisation has on the overall performance of the construction. A surprisingly high number of 61.8% of the automata are universal. A reason might be the small alphabet of the GOAL test set automata, which has just two symbols. With a certain number of transitions in the automata, there seems to be a high probability that the automata are universal. Conversely, the number of empty automata is very low. This can be seen as the reverse of the same effect that causes the number of universal automata to be high.

In Figure 1.2, we show the number of complete, universal, and empty automata per class. In this figure we introduce by the way two ways for representing per-class data that we will use throughout this thesis. On the left side of Figure 1.2 the per-class data is represented as matrices. These matrices always have 11 rows and 10 columns, and the rows always represent the transition densities and the columns represent the acceptance densities. On the right side of Figure 1.2, the same data is visualised as so called perspective plots. The corner of the perspective plots that is closest to the viewer corresponds to the upper-left corner of the corresponding matrices. Thus, looking at a perspective plots is like looking at the corresponding matrix from the upper-left corner. The advantage of matrices is that they show all the data values explicitly. The advantage of perspective plots is that they show the patterns in the data more intuitively. When we present the results of our study in Chapter 2, we will mainly use perspective plots, however, we will give all the corresponding matrices in Appendix B.2.

Regarding the complete automata per class in Figure 1.2 (a), we can see their number increases with the transition density. Up to a transition density of 1.6 there are no complete automata at all, and then it starts to increase up to a number between 34 and 40 for the transition density of 3.0. Since each class contains exactly 100 automata, these numbers are percentages at the same time. That the number of complete automata increases with the transition density is because a higher number of transitions per alphabet symbol in the automaton increases the probability that each state has at least one outgoing transition for each alphabet symbol. For example, with a transition density of 1.0 and 15 states, the automaton contains exactly 15 transitions for each alphabet symbol. It is still possible that this automaton is complete, but the probability is very low, because there must be a one-to-one mapping of transitions and states. On the other hand, with a transition density of 3.0, there would be 45 transitions per alphabet symbol, and the probability that each state gets one of them is much higher.

contains additionally the 11,000 automata of size 20. In the package of the first link, the files have been renamed, however, the content of the files has not been changed.

⁷An automaton is *complete* if every state has at least one outgoing transition for every symbol of the alphabet. An automaton is *universal* if it accepts every word that can be generated from its alphabet. An automaton is *empty* if it does not accept anything (except the empty word ϵ).

The number of universal automata per class in Figure 1.2 (b) also increases with the transition density, although much stronger. While in the classes with a transition density of 1.0, there are between 3 and 10 universal automata, in the classes with a transition density of 3.0 there are between 97 and 100. As already mentioned, the small alphabet size of the GOAL test set automata and a sufficiently high number of transitions results in a high probability that an automaton accepts every possible word, and thus is universal. In Figure 1.2 (b) we can also see that low acceptance densities result by trend in slightly fewer universal automata. This is because with fewer accepting states there is less chance that a given word is accepted. As we identified the small alphabet size as a possible reason for the high number of universal automata, it would be interesting to test how many universal automata there are in similar automata with a bigger alphabet size.

Conversely to the high number of universal automata, the number of empty automata is very low. The totally 63 empty automata are mainly concentrated in the upper-left corner of the matrix in Figure 1.2 (c). That is, the automata with a low transition density and a low acceptance density have the highest probability to not accept any word, and thus being empty. The reasons for this are basically the opposite reasons for the distribution of the universal automata.

1.2.2 Michel Test Set

The Michel test set is very different from the GOAL test set. It consists of a family of automata, the Michel automata, which exhibit an especially heavy state growth for complementation.

Michel automata have been introduced in 1988 by Max Michel in order to prove a lower bound for the state growth of Büchi complementation of $(n - 2)!$, where n is the number of states of the input automaton [5][11]. Michel constructed a family of automata, characterised by the parameter m , that have $m + 1$ alphabet symbols, and $m + 2$ states. He proved that the complements of these automata cannot have less than $m!$ states. Since the number of states of the input automata is $n = m + 2$, the state growth in terms of input and output states is $(n - 2)!$, which is around $(0.36n)^n$.


The state growth of Michel automata is so heavy that for practical reasons we are restricted to include only the first four Michel automata, that is the ones with $m = \{1, \dots, 4\}$, in our test set. For the Michel automata with $m \geq 5$, the required time and computing power for complementing them with our implementation would by far exceed our available resources. We present some extrapolations in this direction in Section 2.1.2. Despite the small number of automata in the test set, we can still obtain very interesting results from them, as we will see in Chapter 2.

The four Michel automata in our test set are shown in Figure 1.3. We will call them Michel 1, Michel 2, Michel 3, and Michel 4, respectively. As mentioned, Michel automata have $m + 2$ states and an alphabet size of $m + 1$. Furthermore, they all have a single accepting state. Our Michel automata 1 to 4 have thus 3, 4, 5, and 6 states, and alphabet sizes of 2, 3, 4, and 5, respectively.

The interesting thing about the high complementation state-complexity of Michel automata is that they allow to “elicit” large number of states from the complementation constructions that we investigate in our study. In the theoretical approach to Büchi complementation, the main performance metric is the worst-case state complexity. This is the maximum number of states that a construction can produce, in function of the number of states of the input automaton. If for example a construction has a worst-case complexity of $(0.76n)^n$, then, if the input automaton has size n , the maximum size of the complement is $(0.76n)^n$.

Thus, if with the complementation of a Michel automata we measure a certain state growth for one of the constructions, say $(0.99n)^n$, then we can deduce that the worst-case complexity of this construction must be greater than or equal to $(0.99n)^n$. In this way, we can get an idea about the lower bounds for the worst-case complexities of our investigated constructions.

This concludes the presentation of the test data for our study. In the next section we describe the experimental setup in which this test data is used.

Figure 1.3: The Michel automata with $m = \{1, \dots, 4\}$, an alphabet size of $m + 1$, and $m + 2$ states.

1.3 Experimental Setup

In this section we describe the concrete experiments that we executed, including the allocated resources and imposed constraints. As mentioned, the experiments are divided into the internal tests and external tests. In the internal tests we compare different versions of the Fribourg construction with each other. In the external tests, we compare one (the most performant) version of the Fribourg construction with other well-known complementation constructions.

The internal and external test, are done with both, the GOAL and the Michel test set. Thus, there are four groups of experiments: internal-GOAL, internal-Michel, external-GOAL, and external-Michel.

In Section 1.3.1, we present the versions of the Fribourg construction used for the internal tests. These versions differ for the GOAL and the Michel test set, and we present them separately. In Section ??, we present the version of the Fribourg construction used for the external tests (which is also different for the GOAL and the Michel test set), and the concrete versions of the third-party construction against which we compare the Fribourg construction. In Section 1.3.3, we describe the computing environment in which the experiments were executed. Finally, in Section 1.3.4, we present the time and memory limits that were imposed on the experiments.

1.3.1 Constructions for the Internal Tests

The versions of the Fribourg construction used for the internal tests consist of combinations of the three optimisations R2C, M1, and M2, and of the additional options C and R (see list of options for the Fribourg construction in Table 1.2). The sets of versions are different for the two test sets. Our aim in choosing specific versions was to find the most performant version of the Fribourg construction for each test set.

GOAL Test Set

For the internal test with GOAL test set, we use the following eight versions of the Fribourg construction:

1. Fribourg
2. Fribourg+R2C
3. Fribourg+R2C+C
4. Fribourg+M1

5. Fribourg+M1+M2
6. Fribourg+M1+R2C
7. Fribourg+M1+R2C+C
8. Fribourg+R

Version 1 is the plain Fribourg construction without any optimisations or options. Version 2 and 3 aim at investigating the R2C optimisation. In Version 2, the R2C optimisation is applied only to complete input automata, and as we have seen in Section 1.2.1 these are just 9% of the automata. Version 3, on the other hand, makes all input automata complete so that the R2C optimisation can be applied to all automata. The question is whether it is worth to increase the size of an automaton by one (for adding the sink state) and then being able to apply the R2C optimisation, or not.

A very similar question has been investigated in previous work about the Fribourg construction by Göttel [2]. In terms of our above listing, he compared Version 1 with Version 3, by also using the GOAL test set as the test data. His result was that the overall mean complement size of Version 3 is higher than for Version 1. By looking closely at his results we suppose however that the median complement size (which was not recorded by Göttel) might be lower for Version 3 than for Version 1. This would be an interesting relation, and therefore, we decided to reinvestigate this question.

Versions 4 and 5 aim at investigating the M1 and M2 optimisations. As M2 can only be applied together with M1, there are only these two possible combinations. As we will see in Chapter 2, Version 4 shows a better performance for the GOAL test set than Version 5. Therefore, we do not further investigate any versions containing Fribourg+M1+M2. However, in Version 6, we further improve Version 4 by adding R2C to Fribourg+M1. In Version 7, we replace R2C by the R2C+C variant. Finally, Version 8 is the same as Version 1, but all unreachable and dead states are removed from the output automaton. This allows to determine the number of unreachable and dead states that have been produced by Version 1.

Versions 6 and 7 then enhance the “better” one of Version 4 and 5 with R2C and its alternative R2C+C. As we will see in Chapter 2, the better one of Version 4 and 5 in terms of median complement sizes is Version 4. That is, the application of M2 results in a decline, rather than a gain, in performance compared to the application of M1 alone. We have to note at this point that such results are always specific to the used test set, and not universally valid. With a different test set, Version 5 might indeed be better than Version 4. As we will see in the next section, this is the case for our alternative test set consisting of the first four Michel automata.

Version 8, finally, is again the plain Fribourg construction, but this time the output automata are reduced by removing their unreachable and dead states. Comparing the results of Version 8 with Version 1 gives an idea of how many unreachable and dead states the Fribourg construction produces. This is inspired by the paper of the GOAL authors [13] in which the number of unreachable and dead states is one of the main metrics for assessing the performance of a construction.

Michel Test Set

For the internal tests with the Michel test set, we use the following six versions of the Fribourg construction:

1. Fribourg
2. Fribourg+R2C
3. Fribourg+M1
4. Fribourg+M1+M2
5. Fribourg+M1+M2+R2C
6. Fribourg+R

The reasons for selecting these versions is basically the same as for the GOAL test set. However, there are the following differences. First, the Michel automata are complete, thus there is no need to include the C option. Second, for the Michel test set, Fribourg+M1+M2 is more performant than Fribourg+M1. For the GOAL test set, the contrary is the case. This is why in Version 5 we add R2C to Fribourg+M1+M2

rather than to Fribourg+M1, because, as mentioned, our aim is to identify the most performant version for each test set.

1.3.2 Constructions for the External Tests

The constructions used for the external tests consist of the most performant version of the Fribourg construction for each test set, and a fixed set of third-party constructions that are implemented in GOAL.

Regarding the third-party constructions, theoretically all the constructions listed in Table ?? could be used. However, practical reasons prevent us from doing so. In preliminary tests we observed that most of these constructions are very inefficient, or inefficiently implemented, for the automata in the GOAL test set. Using these constructions for our external tests would cause the required memory and time resources to be prohibitively high. According to our tests, this excludes all but the Piterman, Slice, Rank, and Safra constructions from being used. A similar experience has been made by Tsai et al. in their own empirical study with GOAL [13]. They observed that the Ramsey construction could not complete the complementation of any automata in the GOAL test set within the time limit of 10 minutes and memory limit of 1 GB.

Considering these restrictions, we decided to include only the Piterman, Slice, and Rank construction in our external tests. These constructions are furthermore the main representatives of three of the four main complementation approaches, determinization-based, rank-based, and slice-based. The fourth approach would be Ramsey-based, but as mentioned, the Ramsey construction in GOAL is not efficient enough. It would have been possible to also include the Safra construction, but as it belongs to the determinization-based approach and we already have the Piterman construction, we decided to not include it.

For the Slice construction, we chose the Slice+P version (see Table ??) by Vardi and Wilke [18]. According to Tsai et al. [13] this version has a lower worst-case complexity than the alternative Slice version by Kähler and Wilke [3].

The Piterman, Rank, and Slice constructions also have a bunch of options in GOAL (for a complete list of their options it is best to consult the help page for the `complement` command in the command line interface of GOAL⁸). For each construction we included those options that are set by default in the GOAL GUI, except the `MACC` and `R` options. The reason to exclude these options is that they are not part of the actual construction, but they just modify the input and output automata, respectively.

We made an exception for the Piterman construction where we also excluded the `SIM` option. The reason for this is that the `SIM` option simplifies the intermediate NPW of the Piterman construction, which can also be seen as a modification of an (intermediate) output automaton.

Altogether, this gives the following three constructions that we used for the external tests:

1. Piterman+EQ+RO
2. Slice+P+RO+MADJ+EG
3. Rank+TR+RO

Regarding the Fribourg construction, we chose the most performant version for each test set. These versions are:

1. Fribourg+M1+R2C for the GOAL test set
2. Fribourg+M1+M2+R2C for the Michel test set

1.3.3 Execution Environment

As mentioned, we ran all the experiments on the high performance computing (HPC) computer cluster UBELIX of the University of Bern⁹. UBELIX consists of different types of computers (called *nodes*) on which the tasks (called *jobs*) of the cluster users are run.

⁸Type `gc help complement`.

⁹<http://ubelix.unibe.ch>

We ensured that all our experiments run on similar nodes. These nodes have the following specifications:

- Processor: Intel Xeon E5-2665 2.40GHz
- Architecture: 64 bit
- CPUs (cores): 16
- Memory (RAM): 64 GB or 256 GB
- Operating System: Red Hat Enterprise Linux 6.6
- Java platform: OpenJDK Java 6u34
- Shell: GNU Bash 4.1.2

The experiments with the GOAL test set were run on nodes with 64 GB RAM, and the experiments with the Michel test set on special high-memory nodes with 256 GB RAM. Apart from the memory, the specifications of these nodes are identical. The use of the high-memory nodes for the Michel test set was required, because the maximally allocatable memory per CPU core of the nodes with 64 GB memory is 4 GB, and this was not enough to complement the Michel automata. With the high-memory nodes, on the other hand, a total of 16 GB can be allocated per CPU core which is enough to complement all the Michel automata in our test set with all the constructions.

Regarding multicore usage, the behaviour of our runs depends on GOAL, and thus ultimately on Java. GOAL is programmed multi-threaded and thus can use multiple CPUs. Theoretically, our tasks can use up to the total number of 16 CPUs of a node. However, we observed that our GOAL complementation task most likely use 2–4 CPUs¹⁰.

We also measured the execution time of each complementation task as CPU time and real time (also known as wallclock time)¹¹. The CPU time is the time a process is actually executed by the CPU. The real time is the time that passes from the start of a process until its termination, and thus includes the time the process is not executed by the CPU (idle time). If a process runs on multiple CPUs, the CPU time is counted on each CPU separately and finally summed up. This means that for multicore execution (as in our case), the CPU time may be higher than the real time. For single-core execution this is not possible, and the CPU time can only be equal to or lower than the real time. In the analysis of our results, we sometimes present statistics of the execution times. These times are always *CPU times*.

The complementation tasks are executed via the command line interface of GOAL. The complementation of each automata is run by a separate invocation of GOAL. That means that for each automaton a new Java Virtual Machine (JVM) is started up. Consequently, this JVM startup time is included in the measured execution times. According to our observations, this JVM startup time is a constant of approximately two CPU time seconds.

The cluster itself is managed by Oracle Grid Engine (formerly known as Sun Grid Engine) version 6.2¹². This is a load scheduler that automatically dispatches incoming jobs from cluster users to nodes with enough free resources and capacity.

A computer cluster is a multi-user environment and a node can be used by multiple users at the same time. Thus the used nodes may be under different levels of loads for different experiments, depending on how many other users are using the same node and how heavy their computation tasks are. At the moment we don't know whether this has an influence on our experiments, especially on the measurement of the execution time. We observed variations in the measured CPU time for similar tasks. This would also have an influence on the time limit that we describe in the next section. For the moment, we leave further investigations on this topic for future work.

1.3.4 Time and Memory Limits

We defined a time limit of 600 seconds CPU time and a memory limit of 1 GB per complementation task in the GOAL test set. That means, if the complementation of a single automaton is not finished after

¹⁰We can just indirectly guess this number by comparing the measured CPU times and real times.

¹¹These measurements were done with the `time` reserved word of Bash.

¹²<http://www.oracle.com/us/products/tools/oracle-grid-engine-075549.html>

600 seconds CPU time or uses more than 1 GB memory, the task is aborted and marked as a *timeout* or *memory excess*.

These limits correspond to the ones used by the experiments of the GOAL authors [13]. However, from their paper it is not clear if their time limit is in CPU time or wallclock time. But since they used different computing nodes, our results regarding the number of timeouts will anyway differ from theirs.

The reason for these limits is simply the restricted amount of time and memory resources that we have available for the experiments. In an ideal world, we would let every complementation task run to its end, no matter how long it takes and how much memory it uses. This would give a perfectly unbiased picture of the results. By setting time and memory limits, we basically exclude the most difficult automata from the experiment. However, as mentioned, the practical reasons of limited time and computing power force us to make this compromise.

The timeout and memory limit of 1 GB apply just to the automata of the GOAL test set. For the Michel test set we did not set a timeout because we wanted each one of the four automata to finish. The longest complementation task of a Michel automaton consequently lasted 109,810 seconds which is about 28 hours. We set a very high memory limit of 14 GB for the Michel test set to avoid memory excesses as we wanted each automaton to successfully complete. The number of 14 GB is determined by the physically available memory on the used computing nodes. All Michel automata successfully completed with this amount of memory.

The timeout was implemented by the means of the `ulimit` Bash builtin¹³, which allows to set a maximum time after which running processes are killed. The memory limit was implemented by setting the maximum size of the Java heap, which can be done by the `-Xmx` option to the Java Virtual Machine (JVM). The heap is the main memory area of Java and the place where all the objects reside. Note that since our memory limit defines actually the size of the Java heap, the total amount of memory used by the process is higher than our limit, as Java has some other memory areas, for example for the JVM itself. However, this is a rather constant amount of memory and independent from the current automaton, so it does not disturb the relative comparisons of the results.

The presence of aborted complementation tasks requires the consideration of the so called effective samples in the result analysis, as introduced in the experiment paper of the GOAL authors. The effective samples are those automata which have been successfully completed by *all* constructions that are to be compared to each other. Imagine two constructions *A* and *B* where *A* is successful in complementing all the automata, whereas *B* has timeouts or memory excesses at 100 of the automata. If we would now take, for example, the median complement sizes of the two result sets without first extracting the effective samples, then *B* is likely to be assessed as too good relative to *A*, because *B*'s results do not include the 100 automata at which it failed, and which are thus likely to have large complement sizes with *B*. The same 100 automata would however be included in the results of *A*. Therefore, all the result analysis of the experiments with the GOAL test sets, that we present in Chapter 2, are based on the effective samples of the result sets.

¹³<http://linux.die.net/man/1/bash>

Chapter 2

Results and Discussion

Contents

2.1 Internal Tests	17
2.1.1 GOAL Test Set	17
2.1.2 Michel Test Set	24
2.2 External Tests	26
2.2.1 GOAL Test Set	26
2.2.2 Michel Automata	28
2.3 Summary and Discussion of the Results	30
2.4 Limitations of the Approach	30

2.1 Internal Tests

2.1.1 GOAL Test Set

First of all, let us see how many unsuccessful complementation tasks there were with the internal tests on the GOAL test set. Table 2.1 shows the number of timeouts and memory excesses for each of the eight tested versions of the Fribourg construction.

Construction	Timeouts	Memory excesses
Fribourg	48	0
Fribourg+R2C	30	0
Fribourg+R2C+C	54	0
Fribourg+M1	2	0
Fribourg+M1+M2	1	0
Fribourg+M1+R2C	1	0
Fribourg+M1+R2C+C	8	0
Fribourg+R	48	0

Table 2.1: Number of timeouts and memory excesses in the internal tests on the GOAL test set.

As we can see in the table, there were no memory excesses, that is, none of the 11,000 complementation tasks required more than 1 GB memory (actually, Java heap size). On the other hand, there is quite a number of timeouts. If we extract the effective samples from these result sets, we get a set of 10,939 automata. That means that these 10,939 automata have been successfully completed by all the versions of the Fribourg construction, whereas the remaining 61 automata (0.55%) provoked a timeout in at least one of the versions.

Our main analysis is now about the sizes of the complements of these 10,939 automata. With size we mean the number of states of an automaton. A stripchart as the one in Figure 2.1 is a good way to get a first glance of the complement sizes. Each horizontal strip contains 10,939 dots, each one corresponding to a produced complement automaton. The x-position of a dot indicates the size of the corresponding complement automaton.




Figure 2.1: Stripchart with the complement sizes of the 10,939 effective samples of the GOAL test set.

The first thing to note is that the distribution of complement sizes is extremely right-skewed (also known as positive-skewed). The peak close to the left end of the x-axis and there is a long tail toward the right. This means that most of the complements are very small and the frequency decreases with increasing complement size. Finally, there are some few complements which are very large. This distribution implies that the mean is generally higher than the median, because the mean is “dragged” to the right by the few very large complements.

Next, we can compare the distributions of the individual version to each other. While Fribourg and Fribourg+R2C have a similarly long tail, Fribourg+R2C+C has a considerably longer one. This clearly is the effect of increasing the size of 91% of the automata by adding a sink state in order to make them complete. Next, Fribourg+M1, Fribourg+M1+M2, and Fribourg+M1+R2C have similarly long tails that are significantly shorter than the ones of the previous versions. This indicates that the M1 optimisation is effective in reducing the size of otherwise large complements. Fribourg+M1+R2C+C, as expected, again increases the size of the tail due to the C option. Finally, Fribourg+R has a very short tail and an even higher concentration of very small complements. This is because the complements are pruned by removing their unreachable and dead states, and thus at least the complements of the 61.8% universal input automata are reduced to empty automata of size 1.

Having a first impression of the distribution of complement sizes, we can look at some statistics characterising this distribution. Table 2.2 shows the mean complement size along with the classic five-number summary consisting of minimum value, 25th percentile, median, 75th percentile and maximum value for each version of the Fribourg construction. As expected, the mean is always higher than the median. Generally, the median is a more robust metrics than the mean, because it is not affected by the value of outliers. This applies specifically to our distribution where some few very large complements might significantly increase the mean whereas they leave the median unaffected. Therefore, the median will be our main metric, and the subsequent analyses will be based on the median.

Going through the median values in Table 2.2 we encounter some surprises. To begin with, as expected, there is a decrease from Fribourg (761) to Fribourg+R2C (689). Then, however, there is a significant drop to 451 with Fribourg+R2C+C. This is a surprise insofar as by looking at Figure 2.1, Fribourg+R2C+C seems to have the worst performance at a first glance. Indeed it also has the highest mean, which is due to the group of extremely large complements. The median, however, is very low, even lower than the one of Fribourg+M1 with its significantly shorter tail in Figure 2.1. Also the 25th percentile of

Construction	Mean	Min.	P25	Median	P75	Max.
Fribourg	2,004.6	2	222.0	761.0	2,175.0	37,904
Fribourg+R2C	1,955.9	2	180.0	689.0	2,127.5	37,904
Fribourg+R2C+C	2,424.6	2	85.0	451.0	2,329.0	54,648
Fribourg+M1	963.2	2	177.0	482.0	1,138.0	16,260
Fribourg+M1+M2	958.0	2	181.0	496.0	1,156.5	15,223
Fribourg+M1+R2C	937.7	2	152.0	447.0	1,118.0	16,260
Fribourg+M1+R2C+C	1,062.6	2	83.0	331.0	1,208.5	25,002
Fribourg+R	136.3	1	1.0	1.0	21.0	12,312

Table 2.2: Statistics of the complement sizes of the 10,939 effective samples of the GOAL test set.

Fribourg+R2C+C is with 85 one of the lowest. Going to the other side of the median, however, the 75th percentile (2,329) is the largest of all versions. A possible characterisation of this phenomenon is that the C option (together with R2C) makes small complements smaller, and large complements larger. The diminishment of small complements is far-reaching enough that the median is affected by it and decreased significantly.

The next thing we see in Table 2.2 is that the median of Fribourg+M1 (482) is slightly lower than the median of Fribourg+M1+M2 (496). The same applies to the 25th and 75th percentile. This backs up our statement from Section 1.3.1 that Fribourg+M1 performs better on the GOAL test set than Fribourg+M1+M2. The difference is rather small (the median increase is by 2.9%), but it is enough to consider Fribourg+M1 as the better of the two versions, and to combine it therefore with R2C and R2C+C.

Fribourg+M1+R2C brings down the median from 482 to 447, with respect to Fribourg+M1. Also the 25th and 75th percentile are decreased. Adding the C option to Fribourg+M1+R2C, again causes the median to drop dramatically, from 447 to 331. The 25th percentile decreases from 152 to 83. The 75th percentile however increases from 1,118 to 1,208.5. Here we have again the same picture of the effect of adding the C option that we had before. Namely that small complements are made smaller, and large complements are made larger.

Finally, the last row in Table 2.2 with Fribourg+R shows the extent of unreachable and dead states that the Fribourg construction produces. The median is 1, and a further analysis reveals that the single-state complements go up to the 61st percentile. That is, 61% of the complements have a size of 1. This is likely to correspond to the 61.8% of universal automata in the GOAL test set whose complements can be reduced to empty automata with a single state.

Up to now, we only looked at statistics aggregated over the entire test set. But as we know from Section 1.2.1, the 11,000 automata of the GOAL test set are divided into 110 classes which consist of combinations of 11 transition densities and 10 acceptance densities. The transition densities range from 1 to 3 in steps of 0.2, and the acceptance densities range from 0.1 to 1 in steps of 0.1. In each class there are 100 automata. It is assumable that the median complement sizes are different for different classes. In the next step of the analysis, we want to analyse exactly this point. How do the median complement sizes of the different classes compare to each other, and can we identify a pattern of the classes that result in large and small complements?

This type of analysis naturally results in three-dimensional data. The classes themselves have two dimensions (the transition densities and the acceptance densities), and each of them has a third dimension consisting of the median complement size of this class. One way to present such data is by, for example, a 11x10 matrix with the rows being the transition densities, the columns the acceptance densities, and the cells the median values. Another way is by so called perspective plots. We used both of these ways in Section 1.2.1 when we analysed the number of complete and universal automata in each class of the GOAL test set. In the present analysis, we will use perspective plots and omit the matrices for space reasons. However, we present the corresponding matrices in Appendix B.2. In this way, the interested reader may consult the exact values of the median complement sizes for each class. In the same way, we restrict ourselves to the median complement sizes as the third dimension. The same analysis could

also be done, for example, for the mean, the 25th percentile, or the 75th percentile of complement sizes. However, as mentioned, we consider the median to be the most meaningful statistics, and thus exclusively focus on it for the matter of conciseness.

Figures ?? and ?? show the perspective plots of the median complement sizes for the 110 classes of the GOAL test set. Each crossing of two lines on the xy-plane represents a class, and the z-value (height) of this crossing indicates the median complement size of this class. As a help for the orientation, if one imagines the corresponding matrices, as they are presented in Appendix B.2, then in the perspective plots we are looking at these matrices from the bottom-right corner. That is, the class that is in the top-left corner of a matrix (transition density 1.0 and acceptance density 0.1) is the one that is the farthest away from the viewer in the corresponding perspective plot. This orientation holds for all the remaining perspective plots in this thesis. The colours of the squares in the perspective plot (called facets) are a function of the z-values of their four corners and are selected to draw an analogy with topographical terrain maps.




Figure 2.2: Median complement sizes of the 110 transition density/acceptance density classes of the 10,939 effective samples of the GOAL test set.

The first thing we notice when looking at the plots in Figure 2.2 and ?? is that there are indeed large differences in the median complement sizes across the transition/acceptance density classes. Figuratively speaking, there is an oblong mountain (or hill) in the northern part of the area in east-west orientation, whose ridge increases in height from east to west and stays high until the western edge of the area. The mountain is located roughly in the area of transition densities between 1.2 and 2.2 and acceptance densities from 0.1 up to 0.9. This means that the automata of these classes are apparently harder to

complement than the automata of other classes, as they result more frequently in larger complements.

We will further elaborate on the reasons of these different complement sizes across classes, and also try to characterise easy, medium, and hard automata for the Fribourg construction, at the end of this subsection. For now, we focus on the relative differences of median complement sizes between the different versions of the Fribourg construction.


Looking at Figure 2.2, the perspective plots for Fribourg and Fribourg+R2C are rather similar. The top of the mountain ridge is at between 3,500 and 4,000 states with a single peak of around 4900 states in the class with transition density 1.6 and acceptance density 0.3. From Table 2.2 we can learn that the overall median complement size is 761 for Fribourg and 689 for Fribourg+R2C. These low values might surprise at first as the mountain, which is much higher, seems to dominate. However, by taking a closer look, it becomes apparent that around half of the classes are in rather low terrain (less than 1,000 states). Furthermore, the heights of the mountain peak do not allow to deduce anything about the overall median, because the median is not affected by the actual values of the data points which are greater than the median. This is the main characteristics of the median and applies to all the subsequent perspective plots in this chapter. The means in turn are 2,004.6 for Fribourg and 1955.9 for Fribourg+R2C.

Fribourg+R2C+C in Figure 2.2 (c) has an even considerably higher mountain than the previous two versions. The top of the ridge is at around 5,000 states and the peak at the class with transition density 1.6 and acceptance density 0.3 has close to 6,500 states. This increase in height is however not reflected in the overall median of Fribourg+R2C+C relative to Fribourg+R2C. As we have seen in in Table 2.2, the median of Fribourg+R2C+C is 451, and thus 34.5% lower than the median of Fribourg+R2C. By taking a closer look at the perspective plots, the reason for this can be seen, as the low areas of Fribourg+R2C+C are indeed slightly lower than the low areas of Fribourg+R2C. The significantly higher high areas of Fribourg+R2C+C do not have an effect on the overall median. However, they do have an effect on the overall mean which with 2,424.6 is indeed 24% higher than the one of Fribourg+R2C (1,955.9).


Going from the first three perspective plot in Figure 2.2 to the perspective plot of Fribourg+R is like going from the Swiss Alps to a Dutch polder. The mountain shrinks to a small hillock and the rest of the terrain is low and flat. This is because so many complements of the Fribourg construction can be reduced to very small sizes by removing their unreachable and dead states. A look at the corresponding matrix in Appendix B.2 reveals that 68 of the 110 classes have a median complement size of 1. If we further compare this matrix to the matrix with the number of universal automata in Figure ?? (b), we see that all the classes with a median of 1 contain more than 50 universal automata, and the classes with a median greater than 1 contain less than 50 universal automata. There is a total of 100 automata per class. This makes sense as the complements of universal automata are empty automata, and every empty automaton can be reduced to an automaton with a single non-accepting state. Looking at the classes with a median greater than 1, we see that their values are still considerably lower than the ones of the plain Fribourg construction, which indicates that the Fribourg construction generates a large number of unreachable and dead states. However, as already mentioned, this thread of analysis is not the focus of the present thesis, and we leave it for future work.

Figure ?? shows the perspective plots of the remaining four versions of the Fribourg construction, all of which include the M1 optimisation. The first thing we note is that, while the mountain is still there, its height shrinks significantly. For Fribourg+M1, and Fribourg+M1+R2C, the top of the ridge is between is around 2,500 states. This is reflected by the overall means of these two versions compared to their counterparts without the M1 optimisation, Fribourg, and Fribourg+R2C. The decrease of the overall mean from Fribourg to Fribourg+M1 is by 52% (from 2004.6 to 963.2) and from Fribourg+R2C to Fribourg+M1+R2C by 52.1% (from 1955.9 to 937.7). The decreases of the overall medians are by 36.6% (from 761 to 482), and 35.1% (from 689 to 447) for the same two pairs of versions. With this we can confirm that the M1 optimisation brings a significant performance gain.


Regarding the M2 optimisation, we can see that the mountain ridge in the Fribourg+M1+M2 perspective plot is slightly lower than the one in the Fribourg+M1 perspective plot. The flatland regions, however, seem to not change much. This is reflected by the overall mean of Fribourg+M1+M2 which is slightly lower than in Fribourg+M1 (958.9 opposed to 963.2). The overall median, on the other hand, is higher for Fribourg+M1+M2 than for Fribourg+M1 (496 opposed to 482). An interpretation of this behaviour is that the application of the M2 optimisation results in smaller complements for *some* input automata. These automata are especially the “hard” ones that produce large complements. This positive effect of




(a) Fribourg+M1



(b) Fribourg+M1+R2C



(c) Fribourg+M1+M2



(d) Fribourg+M1+R2C+C

M2 does however not affect enough input automata, especially not the “easy” automata, as to improve the overall performance of the construction in terms of median complement size. As already stated previously, we consider therefore Fribourg+M1 as the better construction on the GOAL test set than Fribourg+M1+M2.

Finally, Fribourg+M1+R2C+C differs from Fribourg+M1+R2C in a similar way Fribourg+R2C+C differs from Fribourg+R2C. The higher regions get higher and the lower regions get lower, that is, a performance decline on “hard” automata, but a performance gain on easy automata. The performance gain on the easy automata is however effective enough to decrease the overall median from 447 to 331, which is minus 26%.

With 331 states, Fribourg+M1+R2C+C has the lowest median of all the versions (except Fribourg+R which is a special case because it modifies the output of the construction). However, we still declare Fribourg+M1+R2C as the winner on the GOAL test set, mainly for two reasons. First, while Fribourg+M1+R2C+C has a lower median, the mean is still higher (1062.6 to 937.7 which is a plus of 13.3%). This results from the complements of the hard automata, which are larger than with Fribourg+M1+R2C. From a practical point of view, the mean might be relevant, because it relates more directly to the required computing resources than the median. For example the execution CPU time per complementation task, that we also measured along with our experiments, is 25.4% higher for Fribourg+M1+R2C+C than for Fribourg+M1+R2C. The increase in the average execution time per automaton is from 4.44 to 5.57 seconds and in the total execution time from 48,572 seconds (≈ 135 hours) to 60,919 seconds (≈ 169 hours). Fribourg+M1+R2C, on the other hand, has the lowest mean of all versions. The second reason that we choose Fribourg+M1+R2C as the winner and not Fribourg+M1+R2C+C is that the C option

is not a real part of the construction. It actually modifies the input automata before the construction starts in order to make them better suited for the construction. Fribourg+M1+R2C, on the other hand, includes only native options, and input and output of the construction are not modified. This should also make it fairer to compare the Fribourg construction to other constructions in the external tests (Section ??).

As we have seen, there are big difference in the complement sizes across the different classes of the GOAL test set. Furthermore, there is a certain pattern, namely the mountain in the north-western region of the class matrix. Automata of classes that are in the mountain region seem to be harder to complement than automata from the classes in the flatland region. We attempted to categorise the classes of the GOAL test set into the three groups “easy”, “medium”, and “hard”. To do so, we first averaged the matrices with the median complement sizes of all the eight versions of the Fribourg construction. In this way, we have a mean median complement size for each class. Then we defined two breakpoints that divide the classes into easy, medium, and hard groups. The breakpoints 500 and 1,600 result in an appropriate groups that seem to capture the reality well. The result with these breakpoints can be seen in Figure ??.




Figure 2.4: Difficulty categories of the transition/acceptance density classes of the GOAL test set. Green: easy; yellow: medium; red: hard.

As can be seen in Figure ??, there are 53 easy, 36 medium, and 21 hard classes. The easy classes are mainly those with extreme values. In particular, all the classes with a low or high transition density of 0.1, or 2.8 and 3.0, and a high acceptance density of 1.0 are easy. Furthermore, there is a “triangle” of easy classes between transition densities 2.0 and 2.6. and acceptance densities 0.5 and 0.9. The higher the transition density, the lower the acceptance density may be so that the corresponding class is still easy. The hard classes are roughly those with a transition density between 1.4 and 1.8 and an acceptance density between 0.1 and 0.6. The medium classes finally are situated as a “belt” around the hard classes.

It is interesting that the extreme values of transition density and acceptance density result in easy automata. With a transition density of 1.0 and an alphabet size of 2, each of the 15 states has on average two outgoing and two incoming transitions¹. With a transition density of 3.0, each state has on average 6 outgoing and 6 incoming transitions. These low or high connectivity seems to considerably simplify the complementation task. The same applies to a high acceptance density of 1.0, which means that every state is an accepting state². Generally, we can say that automata with high acceptance densities are easier to complement than automata with lower acceptance densities. This also means that the pattern of easy automata at the extreme values of transition and acceptance density, does not apply to the lower extreme of the acceptance density. Automata with a very low acceptance density of 0.1 are hard to complement—unless they are made easy by a low or high transition density.

Another interesting point is that the hard automata have transition densities between 1.4 and 1.8. It seems that this range of transition densities is the crucial factor in the hardness of a complementation task, and that it is only alleviated by a growing acceptance density. This explains the decline of the mountain ridge from west to east.

¹The transition density multiplied with the number of states defines the number of transition for each symbol of the alphabet in the automaton, see Section 1.2.1.

²The acceptance density defines the percentage of states that are accepting, see Section 1.2.1 as well.

Summarising we can say that transition densities between 1.4 and 1.8 produce the hardest complementation tasks, and that to the both sides the difficulty steadily decreases with declining or growing transition density. Furthermore, a growing acceptance density generally implies easier complementation tasks.

2.1.2 Michel Test Set

Our second test set consists of the four Michel automata with $m = \{1, \dots, 4\}$ that are listed in Figure ???. They have 3, 4, 5, and 6 states, respectively. Our main motivation to use these automata is to test our complementation construction on very difficult automata. In this way we can get an idea about the worst-case state complexity of our construction.

We complemented the four Michel automata with the six versions of the Fribourg construction which are listed in Section 1.3.1. We did not impose a time or memory limitation for the individual complementation tasks, as we wanted every task to finish. The resulting complement sizes are listed in Table 2.3.

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Fribourg	57	843	14,535	287,907	$(1.35n)^n$	0.01%
Fribourg+R2C	33	467	8,271	168,291	$(1.24n)^n$	0.06%
Fribourg+M1	44	448	5,506	81,765	$(1.10n)^n$	0.07%
Fribourg+M1+M2	42	402	4,404	57,116	$(1.03n)^n$	0.12%
Fribourg+M1+M2+R2C	28	269	3,168	43,957	$(0.99n)^n$	0.04%
Fribourg+R	18	95	528	3,315	$(0.64n)^n$	0.35%

Table 2.3: Complement sizes of the Michel automata with $m = \{1, \dots, 4\}$ and 3, 4, 5, and 6 states, respectively.

In the second-last column “Fitted curve” of Table 2.3, we fitted a function of the form $(an)^n$ to the measured data points of the previous four columns. These data points consist of the sizes of the four Michel automata (3, 4, 5, and 6) as. The fitted function $(an)^n$ can be seen as an approximated generalisation of the state growth, where n is the size of the input automaton. The last column “Std. error” shows the standard error that resulted from the fit of the previous column.

We can see in Table 2.3 that the state growths are indeed very large. For example, complementing Michel 4, which has six states, with the plain Fribourg construction results in a complement of 287,907 states. However, the optimisations R2C, M1, and M2 have a large influence on the complement sizes. If we consider Michel 4, then the R2C optimisation alone reduces the complement size from 287,907 to 168,291 which is a reduction of 51.5%. The M1 optimisation has an even larger influence as it reduces the complement size from 287,907 to 81,765 which is a reduction of 71.6%. Adding M2 to M1 further reduces the complement size of Fribourg+M1 by 30.1% (from 81,765 to 57,116). Finally, adding R2C on top of M1 and M2 brings a further reduction of 23% (from 57,116 to 43,957). If we compare the most efficient version (Fribourg+M1+M2+R2C) with the least efficient one (Fribourg), then the three optimisations reduce the complement size by 85.7%, or in other words, the complement size of Fribourg+M1+M2+R2C is 15.3% of the complement size of Fribourg.

It is interesting to see that for the Michel automata Fribourg+M1+M2 is more efficient than Fribourg+M1. For the GOAL test set, Fribourg+M1+M2 had a slightly higher overall median than Fribourg+M1 although it had a slightly lower overall mean. We identified in Section 2.1.1 that the M2 optimisation has a positive effect only on some automata, and that these are mostly the hard automata. Michel automata are very hard automata, and indeed the M2 optimisation has a considerably positive effect. These results support thus the observation we made in Section 2.1.1.

The special version Fribourg+R yields very small complements compared to the other versions. This tells us that the complements of the other versions contain a large number of unreachable and dead states. For example, the complement of Michel 4 of Fribourg+R (3,315 states) is 1.2% of the size of the complement of Fribourg. This means that 98.8% of the 287,907 states of the complement of Fribourg are unreachable and dead states. This is actually not surprising, because, following the proof of Michel [5][11], the smallest possible complement of Michel 4 has 24 states. This is because Michel 4 has $m = 4$ and Michel proved

that the complement has at least size $m!$. This means that the Fribourg construction is very far from reaching these optimal complement sizes.

Up to now we just looked at the specific results of Michel 4. The fitted functions of the form $(an)^n$ summarise the results of all the four Michel automata. These functions give us reference points for the worst-case state complexities of the different versions of the Fribourg construction. For example, for the plain Fribourg construction with its fitted function of $(1.35n)^n$, we know now empirically that this construction produces complements of size $(1.35n)^n$, where n is the size of the input automaton. This means that the real (theoretical) worst-case complexity cannot be lower than $(1.35n)^n$ (but it can still be higher). This bound decreases for the different versions of the Fribourg construction until $(0.99n)^n$ for Fribourg+M1+M2+R2C. This value is still greater than the lower bound of $(0.76n)^n$ for Büchi complementation determined by Yan [?].

We also measured the execution times for the individual complementation tasks. This time was measured in CPU time, that is, the time the task is running on the CPU. Table 2.4 shows the measured values in seconds. We can see that the difference between the least and most efficient version is bigger than for the complement sizes. For example for Michel 4, Fribourg+M1+M2+R2C is more than 43 times faster than Fribourg (2,332.6 seconds compared to 100,976 seconds). In more familiar unities, this corresponds to approximately 39 minutes for Fribourg+M1+M2+R2C against 28 hours for Fribourg. We also fitted functions of the form $(an)^n$ to the measured execution times where n is the number of states of the input automaton, and the value of the function is the execution time of the tasks in CPU time seconds.

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Fribourg	2.3	4.0	88.8	100,976.0	$(1.14n)^n$	0.64%
Fribourg+R2C	2.3	3.4	27.4	27,938.3	$(0.92n)^n$	0.64%
Fribourg+M1	2.2	3.6	17.9	6,508.4	$(0.72n)^n$	0.63%
Fribourg+M1+M2	2.3	3.5	13.8	2,707.4	$(0.62n)^n$	0.62%
Fribourg+M1+M2+R2C	2.5	3.5	10.8	2,332.6	$(0.61n)^n$	0.62%
Fribourg+R	2.4	3.7	86.0	101,809.6	$(1.14n)^n$	0.64%

Table 2.4: Execution times.

The fitted functions that we calculated for the measured complement sizes and execution times are based on four data points and basically only valid for the first four Michel automata. They can thus not be used to reliably extrapolate values for larger Michel automata. However, it is still interesting to do such an extrapolation in order to see the involved complexity and to show why we were restricted to include only the first four Michel automata in the test set. In Table 2.5, we show extrapolated values for the complement sizes and execution times for the plain Fribourg construction (the least efficient one), based on the corresponding fitted functions. The table includes the values for the Michel automata 5 to 8, which have 7 to 10 states.

Automaton	Compl. size $(1.35n)^n$	Exec. time $(1.14n)^n$	\approx days/months/years
Michel 5	6,882,980	2,020,385	23 days
Michel 6	189,905,394	46,789,245	18 months
Michel 7	5,939,189,262	1,228,250,634	39 years
Michel 8	207,621,228,081	36,039,825,529	1,142 years

Table 2.5: Extrapolated values for the complement sizes and execution times for the Michel automata with $m = \{5, \dots, 10\}$ with the plain Fribourg construction.

According to the fitted state growth function, the complement of Michel 5 would have nearly 7 million states, and the complement of Michel 8 even more than 207 billions. Already the computation of the 7 million states of Michel 5 would most probably exceed the available memory resources in our computing environment. Regarding the extrapolated execution times, the complementation of Michel 5 would take 23 days. This would similarly exceed the maximum running time of a job on the computer cluster on which we executed the computations. And even without these restriction, at least starting from Michel 6,

the execution times between 18 months and 1,142 years are clearly too long to be included in a master's thesis.

2.2 External Tests

In the external tests we compare the results of the most efficient version of the Fribourg construction to the results of three pre-implemented constructions in GOAL. These constructions are Piterman+EQ+RO, Slice+P+RO+MADJ+EG, and Rank+TR+RO. The most efficient version of the Fribourg construction differs for the two test sets. For the GOAL test set it is Fribourg+M1+R2C, and for the Michel test set it is Fribourg+M1+M2+R2C. In the following two sections we present the results of all these constructions on the two test sets.

2.2.1 GOAL Test Set

For the GOAL test set we compared Fribourg+M1+R2C with Piterman+EQ+RO, Slice+P+RO+MADJ+EG, and Rank+TR+RO. As for the internal tests, we set a time limit of 600 seconds CPU time, and a memory limit determined by a Java heap size of 1 GB per complementation task. If a task does not finish within these limits, it is aborted and marked as either a timed out or memory exceeded. Table 2.6 shows the number of timeouts and memory excesses that we observed for the four constructions.

Construction	Timeouts	Memory excesses
Piterman+EQ+RO	2	0
Slice+P+RO+MADJ+EG	0	0
Rank+TR+RO	3,713	83
Fribourg+M1+R2C	1	0

Table 2.6: Number of timeouts and memory excesses.

Most striking in Table 2.6 is the high number of unsuccessful tasks for the Rank construction. 3,317 of the 11,000 automata (33.8%) were aborted due to a timeout, and further 83 (0.8%) due to a memory excess. Regarding the other constructions, there are just two timeouts for the Piterman, and one timeout for the Fribourg construction.

Determining the effective samples of these runs gives a number of 7,204 automata (65.5%). In Figure 2.5 we present the complement sizes of these 7,204 effective samples as a stripchart. Looking at this picture makes the reason for the high number of aborted tasks of the Rank construction apparent. Rank+TR+RO produces a very big number of very large complements compared to the other constructions.




Figure 2.5: Complement sizes of the 7,204 effective samples.

But from the stripchart in Figure 2.6 alone we cannot yet tell whether the Rank construction *generally* produces larger complements than the other constructions, or if this holds just for *some* automata. Therefore, we calculated further statistics of the complement sizes of the 7,204 effective samples in Table 2.7.

Construction	Mean	Min.	P25	Median	P75	Max.
Piterman+EQ+RO	106.0	1	29.0	58.0	121.0	4,126
Slice+P+RO+MADJ+EG	555.4	2	70.0	202.0	596.0	41,081
Rank+TR+RO	5,255.6	2	81.0	254.5	3,178.2	120,674
Fribourg+M1+R2C	662.9	2	101.0	269.0	754.5	37,068

Table 2.7: Aggregated statistics of complement sizes of the 7,204 effective samples.

And indeed, the 25th percentile and the median of Rank are higher than for Piterman and Slice, but still lower than for our Fribourg construction. However, the picture changes dramatically for the 75th percentile where the value of Rank is more than four times higher than the value for Fribourg. Also the mean of Rank is many times higher than the means of all the other constructions. A possible explanation for this is that the Rank construction performs comparably with the other constructions on easy automata. For the one half of easier automata of the 7,204 effective samples, Rank performs even better than our Fribourg construction, as shown by the median values. For harder automata, however, the performance of Rank is decreased tremendously, and the construction results in very large complements that are beyond any comparison with the complements of the other constructions. In addition, the automata that are hardest for Rank are not even included in this analysis. These are the 34.5% of the test set that Rank could not complete within the given resource restrictions. If we would have these results and include them in the analysis, then the picture would probably look even much more distorted.

What we cannot tell is whether the automata which are hard for Rank are the same that are hard for the other constructions. However, as we will see later, we think that this is not necessarily the case.

Given the large number of unsuccessful complementation tasks of Rank we decided to do the main analysis and comparison of the results without the Rank construction. Because with the Rank construction would basically exclude more than one third of the tasks that have been successfully completed by the other constructions from the result analysis. Our main interest is to compare the performance of the Fribourg construction to the other constructions. Above we have already compared it to the special behaviour of the Rank construction, and by excluding Rank from the further analysis, we can compare the Fribourg constructions in full detail with the Piterman and Slice constructions.

Without Rank there are 10,998 effective samples, that is, there are only two automata that have not been completed by all the three of Piterman, Slice, and Fribourg. In Figure 2.6 we display the complement sizes of these 10,998 effective samples as a stripchart.




Figure 2.6: Complement sizes of the 10,998 effective samples.

From the stripchart we can see that Fribourg and Slice have a comparable distribution of complement sizes, whereas Piterman has a considerably higher concentration of small complement sizes. We can say that Piterman generally produces smaller complement than Fribourg and Slice.

We present the statistics of these distributions in Table 2.8. Indeed, for all statistics Piterman has values that are in the order of magnitudes lower than the ones of Fribourg and Slice. It is interesting that this order of magnitude for the mean, 25th percentile, median, and 75th percentile is very roughly five. It seems that Piterman produces throughout complements that are around five times smaller than the complements of Fribourg and Slice.

Construction	Mean	Min.	P25	Median	P75	Max.
Piterman+EQ+RO	209.6	1	38.0	80.0	183.0	20,349
Slice+P+RO+MADJ+EG	949.4	2	120.0	396.0	1,003.0	41,081
Fribourg+M1+R2C	1,017.3	2	153.0	452.0	1,134.0	37,068

Table 2.8: Aggregated statistics of complement sizes of the 10,998 effective samples without Rank.

Comparing Fribourg and Slice, there is a slight favour for Slice. Mean, 25th percentile, median, and 75th percentile are lower for Slice than for Fribourg by 6.7%, 21.6%, 12.4%, and 11.6%, respectively. We have to conclude that from an overall point of view, the Fribourg construction has the second-worst performance for the GOAL test set after Piterman and Slice, and before Rank.

But now let us look at the results for the 110 classes of transition density and acceptance density classes of the GOAL test set. As already for the internal tests, we take the median complement size as our statistic of interest and calculate it for each of the 110 classes. The results as perspective plots are shown in Figure 2.7. The same data in matrix form can be found in Appendix B.2.




Figure 2.7: Median complement sizes (10,998 samples)

First of all, note that the data displayed in the perspective plot of Fribourg+M1+R2C in Figure 2.7 (c) is the same as the one from the internal tests in Figure ?? (b). The only difference is the range of the z -axis, which in Figure ?? ranges from 0 to 6,500, and in Figure 2.7 from 0 to only 3,000.

We can see that the pattern for Fribourg and Slice are very similar. The median complement sizes in the individual classes do not differ a lot, both relatively and absolutely. However, the medians of Fribourg seem to be throughout (with some exceptions) slightly higher than the ones of Slice. This means that Fribourg and Slice seem to have similar strengths and weaknesses, but Slice is slightly more efficient on the tested automata.

Piterman, as expected, has medians that are throughout multiple times lower than the corresponding medians of Fribourg and Slice. The basic pattern, however, is still similar. There is a mountain ridge along the classes with a transition density of 1.6 with its top in the class with transition density 1.6 and acceptance density 0.1. Thus, this supports our claim from before that throughout all of the tested automata, Piterman produces complements that are multiple times smaller than the corresponding ones of Fribourg and Slice.

2.2.2 Michel Automata

We tested the same three third-party constructions, Piterman+EQ+RO, Slice+P+RO+MADJ+EG, and Rank+TR+RO on the Michel automata 1 to 4 and compared the results with the version of the

Fribourg construction that was most efficient on the Michel automata in the internal tests. This version is Fribourg+M1+M2+R2C. As for the internal tests, we did not set a time or memory limitation as we wanted every task to successfully finish.

The resulting complement sizes are shown in Table 2.8. Again, we fitted a function of the form $(an)^n$ to the four measured data points of each construction and calculated the standard error of this fit.

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Piterman+EQ+RO	23	251	5,167	175,041	$(1.25n)^n$	0.29%
Slice+P+RO+MADJ+EG	35	431	6,786	123,180	$(1.18n)^n$	0.02%
Rank+TR+RO	23	181	1,884	25,985	$(0.91n)^n$	0.01%
Fribourg+M1+M2+R2C	28	269	3,168	43,957	$(0.99n)^n$	0.04%

Figure 2.8: Complement sizes of the first four Michel automata.

Considering the results of the GOAL test set, the results in Table 2.8 are surprising. Rank is the most efficient construction. It produces the smallest complements for all Michel automata, and with $(0.91n)^n$ it has the flattest fitted curve of all constructions. This is surprising because for the GOAL test set, Rank produced by far the largest complements, and 34.5% of the test data could not even be completed within the given time and memory limits. With the Michel automata, however, the case seems to be reversed and Rank produces by far the smallest complements.

Rank is followed by the Fribourg construction, which has the second-smallest complements for Michel 3 and 4, and with $(0.99n)^n$ the second-flattest fitted curve. The complements of Michel 2, 3, and 4 of the Fribourg construction are bigger than the ones of the Rank construction by 48.6%, 68.2%, and 69.2%, respectively. With this, our Fribourg construction is relatively close to the clear winner, which is the Rank construction.

The next in the ranking is the Slice construction with a fitted curve of $(1.18n)^n$. for Michel 1 to 3, this is actually the worst construction, but then for Michel 4, the complement is smaller than the one of Piterman what results in the flatter fitted curve. The gap to the second-ranked construction, Fribourg, is big. The complement sizes of Michel 2 to 4 exceed the ones of Fribourg by 60.2%, 114.2%, and 180.2%, respectively. This is also a remarkable point, because for the GOAL test set, Fribourg and Slice showed a very similar performance.

The last in the ranking is Piterman with a fitted curve of $(1.25n)^n$. However, special for Piterman is that it has the smallest complement for Michel 1 (together with Rank), the second-smallest for Michel 2, the third-smallest for Michel 3, and the largest for Michel 4. It is actually the large complement of Michel 4 that makes Piterman having the steepest fitted curve. However, it is still remarkable that the construction which is by far the most efficient on the GOAL test set produces so much worse results for the Michel automata than all the other constructions. Compared with the Rank construction, Piterman's complements of Michel 2 to 4 are 38.7%, 174.3%, and 573.6%, respectively, bigger. And this is exactly the Rank construction that performed so much worse on the GOAL test set that it made a comparison with the Piterman construction nearly impossible. Compared to the Fribourg construction, Piterman produces slightly smaller complements for Michel 1 and 2, but larger ones for Michel 3 and 4. Namely, they are 63.1% and 298.2% larger than the corresponding ones of the Fribourg construction.

Summarising we can say that the ranking of the constructions for the Michel test set is exactly the reverse of the ranking for the GOAL test set. The by far worst construction for the GOAL test set (Rank) is the best one for the Michel test set, and the by far best construction for the GOAL test set (Piterman) is the worst one for the Michel test set (at least for Michel 4). For the Fribourg construction this means that it “advances” from rank 3 of 4 for the GOAL test set to rank 2 of 4 for the Michel test set.

We also measured the execution times of the individual complementation tasks in CPU time seconds. The results are shown in Table 2.9.

Most interesting in this table is the column with the times for Michel 4. The time difference between the best and the worst construction is enormous. While the Rank construction took just 30 seconds to complement Michel 4, the Piterman construction took 75,917.4 seconds which is approximately 21 hours. This is more than 2500 times longer than the Rank construction. Of course the Piterman construction

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Piterman+EQ+RO	2.5	3.8	42.6	75,917.4	$(1.08n)^n$	0.64%
Slice+P+RO+MADJ+EG	2.3	3.6	11.4	159.5	$(0.39n)^n$	0.38%
Rank+TR+RO	2.2	3.0	6.4	30.0	$(0.29n)^n$	0.18%
Fribourg+M1+M2+R2C	2.5	3.5	10.8	2,332.6	$(0.61n)^n$	0.62%

Table 2.9: Execution times for the first four Michel automata.

produced a bigger automata, which naturally requires more time, however, the automaton produced by the Piterman construction is just around 6.7 times bigger than the one of the Rank construction. This means that the Piterman construction must include very inefficient processes before finally arriving at the output automaton.

Furthermore, we can see in Table 2.9 that also the Fribourg construction took relatively long to complement Michel 4 compared to Rank, namely 2,332.6 seconds which are approximately 39 minutes. This is 77.8 times longer than the 30 seconds of Rank. At the same time, Fribourg's complement has just 68.2% more states than Rank's complement. Also compared to the Slice construction the Fribourg construction is slow for Michel 4. Slice's complement is 2.8 times bigger than Fribourg's complement, but with 159.5 seconds the complementation of slice was 14.6 times faster than the complementation of Fribourg.

So there seems to be an inefficiency in the Fribourg construction in terms of execution time for the complementation of Michel 4. However, this inefficiency is by far not as pronounced as for Piterman. While the complement of Piterman is just 4 times bigger, the execution time of Piterman is 32.5 times longer than the one of the Fribourg construction. One could also look at it from the other side and say that not the Fribourg construction is inefficient on Michel 4, but that Rank and Slice are extraordinarily efficient on this automaton.

Finally, these interesting differences in the execution time can only be observed for Michel 4. For Michel 3 there are also differences but they are by far not as pronounced as for Michel 4. If the computational resources would only allow it, it would be very interesting to see how the behaviour is for Michel 5 and beyond. One thing that stays the same for all the four Michel automata is that Rank is always the fastest and Piterman always the slowest construction.

2.3 Summary and Discussion of the Results

2.4 Limitations of the Approach

Appendix A

Plugin Installation and Usage

Since between the 2014-08-08 and 2014-11-17 releases of GOAL certain parts of the plugin interfaces have changed, and we adapted our plugin accordingly, the currently maintained version of the plugin works only with GOAL versions 2014-11-17 or newer. It is thus essential for any GOAL user to update to this version in order to use our plugin.

Appendix B

Median Complement Sizes of the GOAL Test Set

Bla bla bla

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0		0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	130	117	109	77	69	61	56	40	40	29	1.0	171	174	166	124	118	117	100	67	84	35
1.2	387	456	352	281	155	136	101	105	75	45	1.2	622	833	803	877	529	398	320	372	215	53
1.4	822	683	394	376	230	204	151	120	105	63	1.4	2,086	1,618	1,367	1,676	1,065	967	664	682	494	78
1.6	890	594	458	321	237	178	134	114	113	61	1.6	2,465	2,073	2,182	1,959	1,518	1,259	767	545	623	78
1.8	624	507	324	275	196	136	110	92	89	41	1.8	2,310	1,963	1,950	1,988	1,485	1,095	746	418	346	57
2.0	362	286	211	176	117	103	79	64	59	34	2.0	1,318	1,482	1,393	1,461	981	871	434	338	228	50
2.2	248	222	124	116	82	73	56	52	50	28	2.2	1,068	1,145	1,085	1,067	772	747	263	235	158	40
2.4	147	145	114	87	56	48	43	39	35	19	2.4	689	838	809	751	524	466	240	159	93	30
2.6	115	117	67	61	47	42	32	29	29	15	2.6	469	531	555	565	437	360	169	94	71	23
2.8	95	71	52	45	38	29	27	25	23	13	2.8	369	421	536	405	329	224	130	81	58	21
3.0	59	60	47	35	32	27	22	21	20	10	3.0	244	327	360	322	219	176	85	64	49	16

(a) Piterman+EQ+RO

(b) Slice+P+RO+MADJ+EG

Figure B.1: Median complement sizes of the 10,998 effective samples of the external tests without the Rank construction. The rows (1.0 to 3.0) are the transition densities, and the columns (0.1 to 1.0) are the acceptance densities.

Appendix B. Median Complement Sizes of the GOAL Test Set

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	269	308	254	236	238	297	266	156	207	68
1.2	960	1,407	1,479	2,150	1,152	1,090	942	1,206	718	104
1.4	3,426	2,915	2,752	3,393	2,693	3,265	2,263	2,425	1,844	154
1.6	3,799	3,698	4,901	3,926	3,960	3,655	2,580	1,905	2,124	155
1.8	3,375	3,169	3,420	3,967	3,943	3,132	2,246	1,144	971	114
2.0	1,906	2,261	2,383	2,884	2,354	2,096	1,169	932	568	98
2.2	1,467	1,633	1,795	1,942	1,611	1,640	569	499	330	78
2.4	924	1,232	1,319	1,317	1,056	886	514	314	182	59
2.6	625	763	880	945	828	684	316	175	132	44
2.8	483	584	836	690	575	395	240	151	103	41
3.0	319	450	557	523	367	313	155	116	84	32

(a) Fribourg

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	390	438	434	324	328	459	337	204	227	40
1.2	1,576	2,394	2,505	2,996	1,613	1,551	1,166	1,542	1,002	58
1.4	5,007	4,336	4,652	4,877	3,458	3,956	3,169	3,380	1,868	86
1.6	5,067	5,032	6,444	4,868	4,575	3,864	3,211	1,731	1,892	85
1.8	4,016	3,701	3,647	4,523	3,548	3,009	1,808	451	336	62
2.0	1,663	2,276	2,676	3,035	1,925	1,932	464	307	150	54
2.2	989	1,514	1,621	1,826	1,121	846	155	127	93	45
2.4	560	821	919	771	529	267	133	87	55	32
2.6	388	519	524	441	259	219	84	50	41	26
2.8	311	317	396	242	165	95	64	44	33	22
3.0	173	224	211	169	102	72	41	34	27	18

(b) Fribourg+R2C

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	225	223	195	181	187	199	189	124	161	68
1.2	731	971	946	1,071	629	562	488	568	388	104
1.4	2,228	1,701	1,543	1,732	1,241	1,287	945	944	727	154
1.6	2,489	2,263	2,331	2,133	1,777	1,443	964	757	889	155
1.8	2,381	2,027	2,009	2,075	1,618	1,243	1,005	592	515	114
2.0	1,390	1,569	1,416	1,573	1,093	1,008	594	464	330	98
2.2	1,118	1,197	1,150	1,151	879	809	317	330	241	78
2.4	712	885	836	809	580	535	316	231	145	59
2.6	498	569	601	627	497	412	217	137	113	44
2.8	391	455	578	456	374	263	173	119	90	41
3.0	258	350	392	354	253	208	119	97	74	32

(c) Fribourg+R2C+C

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	215	213	189	174	175	192	186	121	156	68
1.2	712	914	913	1,075	619	563	526	620	416	104
1.4	2,075	1,620	1,503	1,650	1,254	1,339	1,003	1,006	848	154
1.6	2,344	2,062	2,340	2,016	1,755	1,520	1,053	858	986	155
1.8	2,205	1,873	1,920	2,040	1,689	1,315	1,080	664	598	114
2.0	1,290	1,485	1,405	1,522	1,134	1,044	652	531	392	98
2.2	1,023	1,119	1,092	1,127	868	875	376	359	262	78
2.4	674	849	790	807	617	544	355	251	156	59
2.6	478	549	594	597	510	431	231	147	116	44
2.8	370	439	559	455	382	283	182	124	93	41
3.0	249	341	388	348	260	225	123	101	77	32

(d) Fribourg+M1

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	225	223	195	181	187	199	189	124	161	68
1.2	731	971	946	1,071	629	562	488	568	388	104
1.4	2,228	1,701	1,543	1,732	1,241	1,287	945	944	727	154
1.6	2,489	2,263	2,331	2,133	1,777	1,443	964	757	889	155
1.8	2,381	2,027	2,009	2,075	1,618	1,215	1,005	592	515	114
2.0	1,390	1,513	1,416	1,542	1,093	1,003	594	441	330	97
2.2	1,019	1,156	1,064	1,104	859	785	304	303	221	78
2.4	672	867	789	772	544	478	269	191	139	55
2.6	466	542	572	568	452	348	183	129	99	43
2.8	368	407	480	337	260	197	129	96	75	36
3.0	201	261	266	272	199	136	83	74	50	27

(e) Fribourg+M1+M2

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	329	303	279	240	229	288	230	157	160	40
1.2	988	1,392	1,356	1,352	751	741	608	704	516	58
1.4	2,939	2,581	2,066	2,190	1,351	1,622	1,132	1,261	932	86
1.6	3,150	2,900	2,842	2,218	1,885	1,563	1,177	821	896	85
1.8	2,782	2,485	2,047	2,180	1,625	1,269	855	395	309	62
2.0	1,338	1,638	1,544	1,566	979	957	349	261	147	54
2.2	838	1,125	993	1,027	667	521	153	125	93	45
2.4	494	700	624	524	296	214	126	87	55	32
2.6	327	434	383	334	212	163	82	50	41	26
2.8	283	273	305	202	144	95	60	44	33	22
3.0	164	200	173	142	92	72	41	34	27	18

(f) Fribourg+M1+R2C

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
1.0	126	118	97	60	51	52	62	36	48	30
1.2	432	517	345	262	160	126	92	120	109	40
1.4	1,044	331	133	89	45	22	19	31	27	20
1.6	358	24	11	5	4	6	5	3	3	4
1.8	19	5	1	1	1	1	1	1	1	1
2.0	1	1	1	1	1	1	1	1	1	1
2.2	1	1	1	1	1	1	1	1	1	1
2.4	1	1	1	1	1	1	1	1	1	1
2.6	1	1	1	1	1	1	1	1	1	1
2.8	1	1	1	1	1	1	1	1	1	1
3.0	1	1	1	1	1	1	1	1	1	1

(g) Fribourg+M1+R2C+C

(h) Fribourg+R

Figure B.2: Median complement sizes of the 10,939 effective samples of the internal tests on the GOAL test set. The rows (1.0 to 3.0) are the transition densities, and the columns (0.1 to 1.0) are the acceptance densities.

Appendix C

Execution Times

Construction	Mean	Min.	P25	Median	P75	Max.	Total	\approx hours
Fribourg	8.5	2.5	3.3	4.9	7.3	586.0	93,351.2	259
Fribourg+R2C	6.6	2.2	2.9	4.2	6.4	219.7	72,545.7	202
Fribourg+R2C+C	8.5	2.2	2.6	3.5	6.4	582.9	93,396.2	259
Fribourg+M1	4.9	2.5	3.2	4.1	5.9	55.1	54,061.3	150
Fribourg+M1+M2	4.6	2.2	2.9	3.8	5.1	38.4	49,848.0	138
Fribourg+M1+R2C	4.4	2.2	2.8	3.6	5.3	42.5	48,572.0	135
Fribourg+M1+R2C+C	5.6	2.5	3.2	4.0	6.5	147.4	60,918.9	169
Fribourg+R	7.5	2.2	3.0	3.9	6.3	470.5	82,387.3	229

Table C.1: Execution times in CPU time seconds for the 10,939 effective samples of the GOAL test set.

Construction	Mean	Min.	P25	Median	P75	Max.	Total	\approx hours
Piterman+EQ+RO	3.0	2.2	2.6	2.8	3.0	42.9	21,410.6	59
Slice+P+RO+MADJ+EG	3.7	2.2	2.7	3.2	4.1	36.7	26,398.9	73
Rank+TR+RO	16.0	2.3	2.8	3.7	9.3	443.3	115,563.9	321
Fribourg+M1+R2C	4.0	2.2	2.7	3.1	4.4	410.4	28,970.8	80

Table C.2: Execution times in CPU time seconds for the 7,204 effective samples of the GOAL test set.

Construction	Mean	Min.	P25	Median	P75	Max.	Total	\approx hours
Piterman+EQ+RO	3.6	2.2	2.7	2.9	3.4	365.7	39,663.4	110
Slice+P+RO+MADJ+EG	4.3	2.2	2.9	3.7	5.0	42.4	47,418.2	132
Fribourg+M1+R2C	4.7	2.2	2.8	3.6	5.3	410.4	52,149.0	145

Table C.3: Execution times in CPU time seconds for the 10,998 effective samples of the GOAL test set without the Rank construction.

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Fribourg	2.3	4.0	88.8	100,976.0	$(1.14n)^n$	0.64%
Fribourg+R2C	2.3	3.4	27.4	27,938.3	$(0.92n)^n$	0.64%
Fribourg+M1	2.2	3.6	17.9	6,508.4	$(0.72n)^n$	0.63%
Fribourg+M1+M2	2.3	3.5	13.8	2,707.4	$(0.62n)^n$	0.62%
Fribourg+M1+M2+R2C	2.5	3.5	10.8	2,332.6	$(0.61n)^n$	0.62%
Fribourg+R	2.4	3.7	86.0	101,809.6	$(1.14n)^n$	0.64%

Table C.4: Execution times in CPU time seconds for the four Michel automata.

Construction	Michel 1	Michel 2	Michel 3	Michel 4	Fitted curve	Std. error
Piterman+EQ+RO	2.5	3.8	42.6	75,917.4	$(1.08n)^n$	0.64%
Slice+P+RO+MADJ+EG	2.3	3.6	11.4	159.5	$(0.39n)^n$	0.38%
Rank+TR+RO	2.2	3.0	6.4	30.0	$(0.29n)^n$	0.18%
Fribourg+M1+M2+R2C	2.5	3.5	10.8	2,332.6	$(0.61n)^n$	0.62%

Table C.5: Execution times in CPU time seconds for the four Michel automata.

Bibliography

- [1] C. Althoff, W. Thomas, N. Wallmeier. Observations on Determinization of Büchi Automata. In J. Farré, I. Litovsky, S. Schmitz, eds., *Implementation and Application of Automata*. vol. 3845 of *Lecture Notes in Computer Science*. pp. 262–272. Springer Berlin Heidelberg. 2006.
- [2] C. Göttel. Implementation of an Algorithm for Büchi Complementation. BSc Thesis, University of Fribourg, Switzerland. November 2013.
- [3] D. Kähler, T. Wilke. Complementation, Disambiguation, and Determinization of Büchi Automata Unified. In L. Aceto, I. Damgård, L. Goldberg, et al, eds., *Automata, Languages and Programming*. vol. 5125 of *Lecture Notes in Computer Science*. pp. 724–735. Springer Berlin Heidelberg. 2008.
- [4] O. Kupferman, M. Y. Vardi. Weak Alternating Automata Are Not that Weak. *ACM Trans. Comput. Logic*. 2(3):pp. 408–429. Jul. 2001.
- [5] M. Michel. Complementation is more difficult with automata on infinite words. *CNET, Paris*. 15. 1988.
- [6] D. E. Muller, P. E. Schupp. Simulating Alternating Tree Automata by Nondeterministic Automata: New Results and New Proofs of the Theorems of Rabin, McNaughton and Safra. *Theoretical Computer Science*. 141(1–2):pp. 69 – 107. 1995.
- [7] N. Piterman. From Nondeterministic Buchi and Streett Automata to Deterministic Parity Automata. *Logical Methods in Computer Science*. 3(5):pp. 1–21. 2007.
- [8] S. Safra. On the Complexity of Omega-Automata. In *Foundations of Computer Science, 1988., 29th Annual Symposium on*. pp. 319–327. Oct 1988.
- [9] S. Schewe. Büchi Complementation Made Tight. In *26th International Symposium on Theoretical Aspects of Computer Science-STACS 2009*. pp. 661–672. 2009.
- [10] A. P. Sistla, M. Y. Vardi, P. Wolper. The Complementation Problem for Büchi Automata with Applications to Temporal Logic. *Theoretical Computer Science*. 49(2–3):pp. 217 – 237. 1987.
- [11] W. Thomas. Languages, Automata, and Logic. In G. Rozenberg, A. Salomaa, eds., *Handbook of Formal Languages*. pp. 389–455. Springer Berlin Heidelberg. 1997.
- [12] W. Thomas. Complementation of Büchi Automata Revisited. In J. Karhumäki, H. Maurer, G. Păun, et al, eds., *Jewels are Forever*. pp. 109–120. Springer Berlin Heidelberg. 1999.
- [13] M.-H. Tsai, S. Fogarty, M. Vardi, et al. State of Büchi Complementation. In M. Domaratzki, K. Salomaa, eds., *Implementation and Application of Automata*. vol. 6482 of *Lecture Notes in Computer Science*. pp. 261–271. Springer Berlin Heidelberg. 2011.
- [14] M.-H. Tsai, Y.-K. Tsay, Y.-S. Hwang. GOAL for Games, Omega-Automata, and Logics. In N. Sharygina, H. Veith, eds., *Computer Aided Verification*. vol. 8044 of *Lecture Notes in Computer Science*. pp. 883–889. Springer Berlin Heidelberg. 2013.
- [15] Y.-K. Tsay, Y.-F. Chen, M.-H. Tsai, et al. Goal: A Graphical Tool for Manipulating Büchi Automata and Temporal Formulae. In O. Grumberg, M. Huth, eds., *Tools and Algorithms for the Construction and Analysis of Systems*. vol. 4424 of *Lecture Notes in Computer Science*. pp. 466–471. Springer Berlin Heidelberg. 2007.

- [16] Y.-K. Tsay, Y.-F. Chen, M.-H. Tsai, et al. Goal Extended: Towards a Research Tool for Omega Automata and Temporal Logic. In C. Ramakrishnan, J. Rehof, eds., *Tools and Algorithms for the Construction and Analysis of Systems*. vol. 4963 of *Lecture Notes in Computer Science*. pp. 346–350. Springer Berlin Heidelberg. 2008.
- [17] Y.-K. Tsay, Y.-F. Chen, M.-H. Tsai, et al. Tool support for learning Büchi automata and linear temporal logic. *Formal Aspects of Computing*. 21(3):pp. 259–275. 2009.
- [18] M. Y. Vardi, T. Wilke. Automata: From Logics to Algorithms. In J. Flum, E. Grädel, T. Wilke, eds., *Logic and Automata: History and Perspectives*. vol. 2 of *Texts in Logic and Games*. pp. 629–736. Amsterdam University Press. 2007.